

mybibli

User Manual

The mybibli authors
<https://github.com/guycorbaz/mybibli>

Version 1.8.4 — June 3, 2026

Licensed under AGPL-3.0-or-later.

This manual covers mybibli version 1.3.1 and later (the v1.1.0 floor is still enforced for fresh installs — see the Installation chapter for the seed-gate background).

Contents

Preface	1
1 Installation	3
1.1 System requirements	3
1.2 Pre-installation checklist	3
1.3 Method 1 — Bundled database	4
1.4 Method 2 — External database (Synology NAS example)	6
1.5 Environment variable reference	7
1.6 First boot — the setup wizard	9
1.7 Verifying the install	9
2 Configuration	10
2.1 Storage locations and hierarchy	10
2.2 Labels: V-codes and L-codes	11
2.3 Hardware: barcode scanner	11
2.4 Reference data	12
2.5 System settings	13
2.6 Users	14
3 Everyday use	15
3.1 The home page	15
3.2 Cataloging a new title via barcode	15
3.3 Searching the collection	16
3.4 Moving a volume to another location	17
3.5 Loans	17
3.6 Library audit	18
3.7 Shelving workflow	18
3.8 Wish list	18
3.8.1 Adding an entry	18
3.8.2 Browsing at the bookstore	19
3.8.3 Marking a book as bought	19
3.8.4 Auto-detect on catalog scan	19
4 Multiple media types	20
4.1 Books	20
4.2 Comics and BD	20
4.3 Audio releases	20
4.4 Films and series	21
4.5 Picking the media type at scan time	21
5 Metadata providers	22

5.1	The provider chain	22
5.2	API keys	23
5.3	Where the keys live	24
5.4	Provider health	24
5.5	Rate limits and back-pressure	24
5.6	Troubleshooting empty metadata	25
6	Roles and permissions	26
6.1	Who can do what	26
6.2	Sessions and the inactivity timeout	26
6.3	The last-active-admin guard	27
6.4	Password reset	27
6.5	Anonymous sessions	27
7	Backup and restore	28
7.1	What to back up	28
7.2	Backup procedure	28
7.3	Recommended cadence	29
7.4	Restore procedure	29
7.5	Migration to a new host	30
8	Upgrade and migration	31
8.1	Routine upgrades (1.y to 1.z)	31
8.2	Skipping intermediate versions	31
8.3	The 1.0.x to 1.1.0 jump	32
8.4	Reading the release notes	33
8.5	What's new in 1.8.4	33
8.6	What's new in 1.8.3	33
8.7	What's new in 1.8.2	34
8.8	What's new in 1.8.1	35
8.9	What's new in 1.8.0	35
8.10	What's new in 1.7.11	36
8.11	What's new in 1.7.10	37
8.12	What's new in 1.7.9	37
8.13	What's new in 1.7.8	39
8.14	What's new in 1.7.7	39
8.15	What's new in 1.7.6	40
8.16	What's new in 1.7.5	41
8.17	Rollback	42
9	Best practices	43
9.1	Storage hierarchy	43
9.2	V-code rolls	44
9.3	Genre versus Dewey	44
9.4	Audit cadence	44
9.5	User accounts	44
9.6	Loans	45
9.7	Metadata providers	45
9.8	Upgrades	45
10	Troubleshooting	46
10.1	Reading the logs	46

10.2 App does not start	46
10.3 Cannot log in	47
10.4 Stray users in the <code>users</code> table after the wizard	47
10.5 Metadata never resolves	48
10.6 Search returns too little	48
10.7 Cover images are missing	49
10.8 Where to file bugs	49
11 API & integrations	50
11.1 Why an HTTP API?	50
11.2 Creating an API key	50
11.3 Authenticating	51
11.4 Endpoints	51
11.4.1 GET <code>/api/v1/titles</code>	51
11.4.2 GET <code>/api/v1/titles/{id}</code>	52
11.4.3 GET <code>/api/v1/genres, /locations, /series</code>	52
11.4.4 GET <code>/api/v1/dewey/{prefix}</code>	52
11.4.5 PATCH <code>/api/v1/titles/{id}</code>	52
11.5 Examples	52
11.5.1 curl	52
11.5.2 Python	53
11.5.3 An LLM classifier loop	53
11.6 Revoking a key	54
11.7 What the API does <i>not</i> expose	54
12 Operations & debugging	55
12.1 Where the logs live	55
12.2 Tailing logs from the host	55
12.3 Log levels	56
12.4 Reading the structured JSON	57
12.5 Sharing a log excerpt for a bug report	57
12.6 Retention & rotation	57
13 Glossary	59
Index	62

Preface

Who this manual is for

This manual is written for the person who installs, configures and runs a mybibli instance for their household, a small association, or a community library. It assumes you can:

- open a terminal and run `docker compose` commands;
- edit a plain-text configuration file with the editor of your choice;
- read a URL printed in a log line and open it in a browser.

No Rust, MariaDB or HTMX knowledge is required. The manual is deliberately scoped to the *operator* role — building mybibli from source or modifying its code is covered by `CLAUDE.md` and the planning artifacts in the repository.

Conventions

Commands. Shell snippets are shown in a fixed-width box:

```
docker compose up -d
```

Configuration files. File excerpts use the same style with the file name called out in the surrounding text. Comments inside the snippet use the file's native comment marker.

Callouts. Three sidebars highlight important text:

Note

A *note* provides background or context that helps you understand why something works the way it does. Skipping it is safe.

Tip

A *tip* suggests a practice that makes day-to-day use smoother. Try it once you are comfortable with the basics.

Warning

A *warning* flags a step where a mistake can lose data, lock you out of the admin panel, or expose the instance to unauthorized access. Read these in full.

Cross-references. Chapters and sections are linked in the PDF — click any blue title or page number to jump there. The [index](#) at the end maps every important term to its page of first use.

Getting help

The canonical source of truth for mybibli is the GitHub repository:

<https://github.com/guycorbaz/mybibli>

Open an issue against the `guycorbaz/mybibli` repository when you hit a bug, a documentation gap, or a missing feature. The issue templates ask for the version (visible at the bottom of any page in the admin panel), your install method, and the steps to reproduce.

Manual coverage

This manual covers mybibli **1.1.0 and later**. Older versions shipped a security-relevant bug where the first-launch setup wizard was silently bypassed — see chapter 8 for the upgrade path from 1.0.x.

Happy cataloging.

Chapter 1

Installation

This chapter walks you through installing mybibli on your own hardware. The reference deployment target is a home server, NAS, or any always-on Linux box that can run Docker.

1.1 System requirements

- **Operating system:** any Linux distribution, macOS, or Windows host capable of running Docker. The host kernel must support cgroups v2 (any distribution from 2021 onward).
- **Docker:** version 24 or later, with the `compose` plugin. Test with `docker compose version`.
- **Disk:** at least 1 GB free for the application image plus the database. Cover images add roughly 50 KB per title.
- **Memory:** 512 MB is enough for a household-sized library (a few thousand titles); 1 GB gives generous headroom for the metadata-fetch pipeline.
- **Network:** outbound HTTPS to the metadata providers you choose to use (see chapter 5). LAN-only deployments work without any inbound port forwarding.

Warning

Do not install mybibli versions below 1.1.0. Every published image with a tag below 1.1.0 contains a seed migration that creates the credentials `admin/admin` and `librarian/librarian` on *every* fresh install, including production. The first-launch setup wizard is silently bypassed. Use 1.1.0 or later; the pre-1.1.0 images will be removed from Docker Hub to prevent accidental use.

1.2 Pre-installation checklist

Before you start:

1. Pick a host directory for the mybibli files (e.g. `/srv/mybibli` on Linux, `/volume1/docker/mybibli` on Synology DSM).
2. Decide how cover images are stored. The shipped `docker-compose.yml` declares a Docker *named volume* (`mybibli_covers`) — zero configuration, Docker manages persistence and backup via `docker volume`. If you prefer a host-visible directory (Synology DSM file

browser, manual rsync backups), pick a host path now (e.g. `/volume1/docker/mybibli/covers`) and you will switch the compose file to a bind mount in step 3 below.

3. Decide how logs are stored (v1.7.0+). The shipped `docker-compose.yml` declares a third named volume (`mybibli_logs`) for the daily-rotating log files. Like for covers, you can keep the Docker-managed volume (zero-config) or switch to a host-visible bind mount for shipping logs to DSM Log Center, journald, or an external log collector. See section 1.3.
4. Choose a strong password for the bundled MariaDB root and application users — or, for an external database, create a dedicated database and user up front (see section 1.4).

1.3 Method 1 — Bundled database

This is the simplest path. The shipped `docker-compose.yml` file declares two services: the mybibli application and a MariaDB instance pinned to a known version. Both are managed as one unit.

1. Get the compose file and `.env`

Download the `docker-compose.yml` and `.env.example` from the release you intend to install (or check out the repository at the matching tag):

```
mkdir -p /srv/mybibli && cd /srv/mybibli
curl -O https://raw.githubusercontent.com/guycorbaz/mybibli/v1.7.0/docker-compose.yml
curl -O https://raw.githubusercontent.com/guycorbaz/mybibli/v1.7.0/.env.example
cp .env.example .env
```

1b. (Optional) Switch covers to a bind mount

By default, the shipped `docker-compose.yml` stores downloaded cover images in the Docker-managed named volume `mybibli_covers`. This works zero-config and persists across container upgrades (issue #213). Most operators should leave it as-is.

If you want covers visible in your host file manager (Synology DSM File Station, manual rsync backups, etc.), edit the compose file:

- In the `mybibli` service's volumes: block, comment out the `- mybibli_covers:/app/covers` line and uncomment the `- ${COVERS_HOST_PATH:-./covers-host}:/app/covers` line right below it.
- In `.env`, set `COVERS_HOST_PATH` to the host path you picked in the pre-installation checklist (e.g. `COVERS_HOST_PATH=/volume1/docker/mybibli/covers`).
- Create the directory before starting the stack (`mkdir -p`) so Docker doesn't auto-create it owned by root.

1c. (Optional) Switch logs to a bind mount

New in **v1.7.0**. By default, mybibli writes its daily-rotating log files to the Docker-managed named volume `mybibli_logs` mounted at `/var/log/mybibli` inside the container. Files are named `mybibli.log.YYYY-MM-DD` and the in-process appender purges files older than 30 days automatically.

If you want logs visible directly on the host (Synology DSM Log Center, journald shipping, external log aggregator, manual `tail`), edit the compose file:

- In the `mybibli` service's volumes: block, comment out the `- mybibli_logs:/var/log/mybibli` line and uncomment the `- ${LOGS_HOST_PATH:-./logs-host}:/var/log/mybibli` line right below it.
- In `.env`, set `LOGS_HOST_PATH` to the host path you want (e.g. `LOGS_HOST_PATH=/volume1/docker/mybibli/log`).
- Create the directory before starting the stack (`mkdir -p`) so Docker doesn't auto-create it owned by root.

The logs volume is forensic-only — you can wipe it at any time without affecting catalog data. See chapter 12 for log-tailing and log-level controls.

2. Edit `.env`

Open `.env` in your editor. Every variable carries an explanatory comment; the values that matter for a bundled-database install are:

```
DATABASE_URL=mysql://mybibli:mybibli_password@db:3306/mybibli?charset=utf8mb4
MYSQL_HOST=db
MYSQL_DATABASE=mybibli
MYSQL_USER=mybibli
MYSQL_PASSWORD=<choose-a-strong-password>
MYSQL_ROOT_PASSWORD=<choose-a-strong-password>

HOST_PORT=8080          # change if 8080 is already taken on the host
APP_LANGUAGE=en         # or 'fr' for the French UI by default

MYBIBLI_COOKIE_SECURE=false # set to true only behind HTTPS
```

The `DATABASE_URL` *must* reference the same user, password and database name as the `MYSQL_*` values — the compose file rebuilds the URL from those parts but the Rust binary reads `DATABASE_URL` directly, so the two must agree.

Warning

The default passwords in `.env.example` (`mybibli_password`, `root_password`) are placeholders, not credentials to keep. Change them before the first `docker compose up`. After the database has been created, rotating these passwords requires also rewriting the MariaDB user entries — easier to get it right the first time.

3. Start the stack

```
docker compose up -d
docker compose logs -f mybibli
```

The first launch:

1. pulls the `gcorbaz/mybibli:1.1.0` (or later) image and the MariaDB image;
2. starts the database and waits for it to be healthy;
3. runs the `mybibli` database migrations (one-time schema setup);
4. starts the HTTP listener on the port you configured.

You can stop tailing the logs with `Ctrl-C`; the containers keep running detached.

4. Open the setup wizard

Visit <http://<host>:8080/setup> in a browser. The first-launch setup wizard appears (covered in section 1.6).

1.4 Method 2 — External database (Synology NAS example)

If you already run MariaDB on your NAS — for example via the Synology **MariaDB 10** package — you can point mybibli at it instead of running a second database in Docker. The mybibli image is unchanged; only the compose file and `.env` differ.

1. Create the database and user

Connect to the external MariaDB as `root` and run:

```
CREATE DATABASE mybibli CHARACTER SET utf8mb4 COLLATE utf8mb4_unicode_ci;
CREATE USER 'mybibli'@'%' IDENTIFIED BY 'choose-a-strong-password';
GRANT ALL PRIVILEGES ON mybibli.* TO 'mybibli'@'%';
FLUSH PRIVILEGES;
```

Note

On Synology DSM, MariaDB 10 listens on port 3307 by default (not 3306 — that port is reserved for the older MySQL package). Open **MariaDB 10** → **Settings** and tick *Enable TCP/IP connection* so containers on the same Docker bridge can reach the server.

2. Trim the compose file

Edit `docker-compose.yml` and remove the bundled database:

- delete (or comment out) the entire `db:` service;
- delete (or comment out) the `depends_on:` block on the `mybibli` service;
- in the bottom `volumes:` block, remove `mybibli_db_data:` only — **keep** `mybibli_covers:`. Cover images are local to the app and must survive container upgrades regardless of where the database lives (issue #213).

3. Point `.env` at the external database

```
DATABASE_URL=mysql://mybibli:strong-password@host.docker.internal:3307/mybibli?charset=utf8mb4
```

Substitute the host name and port that match your network. From a mybibli container running on the same Docker bridge as the NAS host, `host.docker.internal` resolves to the host on macOS and recent Linux Docker installs; on older Docker engines, use the LAN IP of the NAS instead (e.g. 192.168.1.10).

The `MYSQL_*` variables are now unused — leave them blank or delete them.

4. Start the application

```
docker compose up -d
docker compose logs -f mybibli
```

The first launch runs the schema migrations against the external database. The Synology MariaDB package will hold this schema for as long as the package exists; backing it up is covered by the NAS-level backup tools (see chapter 7).

Tip

The external-database setup is the recommended path when you already maintain a database server, because it gives you one backup pipeline to monitor instead of two. The bundled-database setup is faster to get going if you have no other workloads on the host.

1.5 Environment variable reference

The full reference lives in the file `.env.example` at the repository root — every variable carries an inline comment. The table below summarises which variable applies where.

All boolean variables follow a **strict accept-set**: only the strings `1`, `true` and `TRUE` count as “on”. Anything else (including the empty string, `0`, and `false`) is treated as “off”. This avoids the classic footgun where a stale shell value like `0` reads as “set” and silently flips an opt-out.

Variable	Default	Purpose
<i>Database</i>		
DATABASE_URL	— (required)	MariaDB connection string (DSN).
MYSQL_HOST	db	Hostname for the bundled DB service.
MYSQL_PORT	3306	Port for the bundled DB service.
MYSQL_DATABASE	mybibli	DB name for the bundled service.
MYSQL_USER	mybibli	DB user for the bundled service.
MYSQL_PASSWORD	—	DB user password for the bundled service.
MYSQL_ROOT_PASSWORD	—	Root password for the bundled service.
<i>HTTP server</i>		
HOST	0.0.0.0	Bind address inside the container.
PORT	8080	Bind port inside the container.
HOST_PORT	8080	Port published by Docker on the host.
<i>Application</i>		
MYBIBLI_LOG_LEVEL	info	v1.7.0+. Runtime log level. Plain level (trace/debug/info/warn/error) or an EnvFilter directive list. Takes precedence over RUST_LOG . Also flippable from /admin > System .
MYBIBLI_LOG_DIR	/var/log/mybibli	v1.7.0+. Directory for daily-rotating log files inside the container. Mapped to the mybibli_logs named volume.
LOGS_HOST_PATH	./logs-host	v1.7.0+. Optional bind-mount host path; used only when you uncomment the bind-mount line in docker-compose.yml .
RUST_LOG	mybibli=info	Legacy tracing filter. Honored when MYBIBLI_LOG_LEVEL is unset.
APP_LANGUAGE	en	Default UI language for anonymous visitors. v1.7.0 adds de + it to en / fr .
COVERS_DIR	./covers	Filesystem path for cover-image storage.
COVERS_HOST_PATH	./covers-host	Optional bind-mount host path; used only when you uncomment the bind-mount line in docker-compose.yml .
<i>Hardening</i>		
MYBIBLI_COOKIE_SECURE	false	Set to true only behind HTTPS.
CSP_REPORT_ONLY	false	Switch CSP header to report-only mode.
<i>Metadata providers (migrated to DB on first boot)</i>		
GOOGLE_BOOKS_API_KEY	—	Optional key for Google Books.
OMDB_API_KEY	—	Optional key for OMDb (films/series).
TMDB_API_KEY	—	Optional key for TMDb (films/series).
<i>Optional dev / test overrides</i>		
MYBIBLI_SKIP_SETUP	false	Bypass the first-launch wizard. <i>Never on prod.</i>
MYBIBLI_SKIP_STARTUP_PURGE	false	Skip the boot-time auto-purge.
MYBIBLI_SEED_DEV_USERS	false	Seed dev admin/admin + librarian/librarian . <i>Never on prod.</i>

Provider base-URL overrides (**BNF_API_BASE_URL**, **GOOGLE_BOOKS_API_BASE_URL**, ...) exist solely to point the metadata pipeline at the in-tree mock server during automated tests. Leave them unset in production.

1.6 First boot — the setup wizard

On a clean install, every request is redirected to `/setup` until the wizard completes. The wizard has four steps:

1. **Admin account.** Choose a username and a strong password for the first administrator. This step is the whole reason the wizard exists — once the account is created, the wizard's gate predicate flips and subsequent visitors land on the public catalog instead.
2. **Metadata providers.** Paste in any API keys you have on hand (Google Books, OMDb, TMDb). All keyed providers are optional — you can skip this step entirely and add keys later from `/admin > System > Metadata Providers`.
3. **Preferences.** Pick the default UI language for anonymous visitors and the loan overdue threshold (in days).
4. **Done.** A summary screen. After confirming, the wizard writes `settings.setup_completed_at` and disables itself for the lifetime of the installation.

The wizard is idempotent: if you close the browser between steps, reopen `/setup` and the wizard resumes server-side at the right step (it queries the database on every load — there is no client state to lose).

1.7 Verifying the install

After the wizard finishes:

- Log in at <http://<host>:8080/login> with the admin credentials you just created.
- Visit `/admin > Health`. The page should show non-zero entity counts (zero on a clean install — that is expected) and a green check next to each metadata provider that has been reachable in the last 5 minutes.
- Scan an ISBN (the search field at the top of the home page) to confirm the metadata pipeline resolves a real title. The first scan takes a couple of seconds; the result appears in the timeline of recent additions on the home page once the background fetch completes.

If anything goes wrong, jump to chapter [10](#).

Chapter 2

Configuration

Once the setup wizard is behind you, the day-one configuration work consists of three blocks: defining your *storage locations*, populating the *reference data* (genres, contributor roles, volume states, location types), and adjusting the system-wide *settings* (overdue threshold, auto-purge cadence, provider keys).

Everything in this chapter is configured from the `/admin` panel — there is no separate configuration file. The panel has five tabs:

- **Health** — read-only status dashboard.
- **Users** — create / deactivate users, change roles.
- **Reference data** — the four taxonomy tables.
- **Trash** — view and restore soft-deleted items.
- **System** — global settings.

Only the `Admin` role sees the panel; `Librarian` and `Anonymous` users get a 403 from every `/admin/*` route.

2.1 Storage locations and hierarchy

Storage locations are the physical spots where your volumes live: a room, a shelf, a row, a numbered box, a specific cell in a cabinet. mybibli organises them as a **tree** — every location has at most one parent. The hierarchy is arbitrary: you decide both the levels and the labels.

A workable household hierarchy might be three levels deep:

Living room	(room)
Bookshelf A	(shelf)
Row 1	(row)
Row 2	(row)
Bookshelf B	(shelf)
Row 1	(row)
Bedroom	(room)
Nightstand	(shelf)

Three levels is enough to find any volume in seconds without forcing the hierarchy to mirror every nook of your living space. A flat list (one big “shelf” per physical shelf) works equally well

for collections under a few hundred volumes.

Editing the tree

Open `/locations`. The tree is rendered as a nested list with inline `create` / `rename` / `move` / `delete` affordances. Drag-and-drop is *not* supported (mybibli ships with a CSP-strict no-JS policy); moves are done by editing a node and picking a new parent from a dropdown.

Location node types

Each location node carries a **node type** — “room”, “shelf”, “row”, “box”, anything you like. Node types are managed under `/admin > Reference data > Location node types`. They are pure labels; mybibli does not enforce that a “shelf” must sit under a “room”.

Tip

Renaming a node type cascades to every location that uses it (the type is stored by name in the locations table). Deleting a node type is refused as long as one location still references it.

2.2 Labels: V-codes and L-codes

mybibli prints two kinds of barcode labels:

- **V-codes** — one per *volume* (the physical book on a shelf). Format: `v` followed by a zero-padded sequence (e.g. `V0042`).
- **L-codes** — one per *storage location*. Format: `L` followed by a zero-padded sequence (e.g. `L0007`).

V-codes and L-codes are scanned with the same barcode scanner you use for ISBNs at cataloging time. The scanner’s prefix (`v` or `L`) tells mybibli which lookup to perform.

Printing labels

Open `/locations` or the relevant volume page and click **Print label**. mybibli renders a printable HTML page with the Code-128 barcode and human-readable text. Print on plain paper or sticker stock — the encoding is forgiving enough that a low-cost inkjet works for V-codes and L-codes alike.

Tip

Print V-code stickers *before* you start cataloging in bulk. A roll of pre-printed V0001 through V0500 means you can grab a sticker, stick it on the back cover, and scan in one motion — no context-switch between the cataloging UI and the printer.

2.3 Hardware: barcode scanner

mybibli is designed around a barcode scanner but does not strictly require one — every ISBN, V-code and L-code field accepts manual keyboard input as a fallback. That said, for any collection beyond a handful of titles, a cheap USB scanner pays for itself within the first dozen scans.

What to buy

Any *HID-keyboard-mode* USB scanner in the 15–30 € range works. The HID mode is what makes the scanner driver-free: the OS sees it as a regular keyboard and the scanned digits land directly in whatever input field has focus. Almost every general-purpose USB scanner on the market ships in this mode by default; look for the words “HID”, “keyboard emulation” or “plug-and-play” on the product page.

The scanner must support the two symbologies mybibli uses:

- **EAN-13** — the 13-digit code on the back cover of books, CDs, DVDs, and most consumer products. Universally supported.
- **Code-128** — the symbology mybibli uses to encode V-codes and L-codes on the labels it prints (see section 2.2). Universally supported.

You do *not* need a 2D imager (the kind that reads QR codes) — a basic 1D laser or LED scanner is enough. mybibli does not use QR codes.

Configuration

Three settings on the scanner matter. Most scanners ship them correctly out of the box; if your scans behave strangely, check the scanner’s manual or the configuration sheet of barcodes that comes in the box:

- **Terminator / suffix:** Enter (also called CR or Carriage Return). mybibli’s cataloging form auto-submits on Enter, so this turns each scan into one round-trip. Anything else (Tab, no terminator) will not work without manual key presses.
- **Prefix:** none. Some scanners can prepend a fixed string before the payload — leave it disabled.
- **Keyboard layout:** *must* match the layout of the host the browser runs on. If your laptop uses an AZERTY layout but the scanner is configured for QWERTY, the digits come out as letters (e.g. 1234567890 arrives as &é" (-è_çà) and no ISBN ever resolves. The scanner’s manual will have a code chart to switch.

Tip

USB scanners on a Synology NAS host: it does not matter that the scanner is plugged into the laptop you use to access the mybibli web UI rather than the NAS itself. The scanner is a keyboard for the browser, not for the server. The same setup works with a phone-side scanner app, an industrial Bluetooth scanner paired to the laptop, or any USB stick scanner.

What to do when the scanner is unavailable

Every scan field also accepts hand-typed input. Click into the field labelled “Scan ISBN” (or “Scan V-code” / “Scan L-code” elsewhere), type the digits, press Enter. The form behaves identically.

2.4 Reference data

Four taxonomy tables drive the dropdowns you see while cataloging. They live under `/admin > Reference data:`

Genres The genres assigned to titles. mybibli ships with a baseline FR/EN seed (Fiction, Non-fiction, Biography, ...) — edit or extend to match your collection.

Contributor roles Author, Illustrator, Translator, Editor, Narrator, ...Used by title detail pages to attribute the right people to the right slots.

Volume states The physical condition of an exemplaire — New, Good, Worn, Damaged, ...Useful for loan returns (mark a volume Worn when the borrower damaged it).

Location node types See section 2.1 above.

CRUD semantics

All four tables share the same CRUD pattern:

- **Create** adds a row. If the name collides with a soft-deleted row of the same table, the existing row is reactivated and the result reads *Reactivated* instead of *Created*.
- **Rename** updates the name. For location node types, the new name cascades to every location that uses the type.
- **Delete** performs a soft-delete *only when the row is unused*. If any title still references the genre / role / state / type, the deletion is refused with a count of dependents.

Reference data is not translated

Reference data values are stored verbatim — the FR seed creates a genre called “Roman” and that is the literal text English-speaking users will see if they switch the UI to English. Only the UI chrome (button labels, navigation, error messages) is translated.

2.5 System settings

The `/admin > System` tab groups every global setting. The panel is divided into three forms, each with its own save button:

Loans

- **Overdue threshold (days)**. Defines how many days a loan can run before it surfaces as overdue on the dashboard. Default 30.
- **Session inactivity timeout (seconds)**. Idle time after which an authenticated session is wiped. Default 1800 s (30 minutes).
- **Auto-purge interval (seconds)**. How often the background scheduler hard-deletes items older than 30 days in the trash. Default 86400 s (24 hours).

Metadata providers

One row per keyed provider (Google Books, OMDb, TMDb). Each row has a key input and a *Clear* button. Saving a key takes effect on the next metadata fetch — no restart, no re-deploy.

Note

If the same provider key is set in `.env` and cleared from the UI, the next reboot re-migrates the value from `.env`. The operator’s deployment-time intent wins on boot. To make a

Clear durable across reboots, also remove the variable from `.env` (or `docker-compose.yml`).

Language

Sets the default UI language for *anonymous* visitors. Authenticated users override this with the language toggle in the navigation bar — the choice is stored on their user row.

2.6 Users

Open `/admin > Users` to manage user accounts. Each row shows the username, role, active flag, and last-seen timestamp. The available actions are:

- **Create user** (top of the page) — username, password, role.
- **Edit** — change role or rename. Cannot change another user's password from this view; password resets go through the user's own **Account** page.
- **Deactivate** — soft-delete the user. Their sessions are wiped immediately; they cannot log back in.
- **Reactivate** — undo a deactivation (from the **Trash** tab).

Two safety guards apply:

- **Self-deactivate is blocked.** You cannot deactivate the account you are logged in as.
- **Last-active-admin guard.** The system always keeps at least one active admin. Deactivating the last one is refused with a 409 conflict and a clear error.

The three roles (Anonymous, Librarian, Admin) are detailed in chapter 6.

Chapter 3

Everyday use

This chapter walks through the workflows you will run dozens of times once the library is live: cataloging a new volume, searching your collection, moving a volume, lending it out, taking it back, and auditing what is on the shelves.

3.1 The home page

The home page is the dashboard. From top to bottom:

- **Scan field.** A focused input that accepts an ISBN, a V-code or an L-code. The scanner-guard layer routes the input to the right page depending on its prefix.
- **Indicators row.** Recent additions, overdue loans, series with gaps, unshelved volumes — each click jumps to the filtered list.
- **Recent activity.** The last seven days of cataloging and returns.

3.2 Cataloging a new title via barcode

1. Focus the home-page scan field (it is selected by default on page load). Scan the ISBN / EAN-13 of the back cover.
2. mybibli redirects to the catalog page, which displays a skeleton title with the barcode you scanned. A background task fires off requests against the metadata provider chain.
3. Within a few seconds, the skeleton fills in: title, contributors, publisher, year, cover image. A small indicator confirms which provider resolved the title.
4. Pick a **location** from the dropdown (or scan the L-code of the destination shelf) and pick a **volume state** (typically *New*). Click **Add volume**.
5. Print and stick a V-code label on the volume.

If the metadata pipeline fails to resolve the ISBN, the title appears with an *Unresolved* badge. You can edit the title fields manually and the volume is still created.




Figure 3.1: The home page on a live install: search bar, genre filter chips, “À traiter” counters, “Aperçu de la collection” totals, and a recent-additions strip with cover thumbnails fetched through the metadata-provider chain.

Tip

Cataloging in batches works best when the V-code sticker roll, the scanner and the printer are all within arm’s reach. Most of the time is spent scanning, not waiting for mybibli — the metadata fetch is asynchronous, so you do not have to wait for it to finish before moving on to the next volume.

3.3 Searching the collection

The `/search` page accepts:

- full or partial **title**;
- **contributor name** (matches on author by default; other contributor roles are scored

lower);

- **ISBN / EAN-13** — exact match;
- **V-code** or **L-code** — exact match;
- **Dewey code** (for libraries that use Dewey).

Results are sorted by relevance. Click a result to open the title page; click a contributor to filter the catalog by their work; click a location label to see every volume currently in that location.

Tip

You can “search” a single volume by scanning its V-code from any page that has the home-page scan field — that includes the home page, `/search` and `/loans`. Faster than typing.

3.4 Moving a volume to another location

1. Open the volume page (search by V-code or click through from a title page).
2. Click **Move**.
3. Scan the L-code of the destination location, or pick it from the dropdown.
4. Confirm.

mybibli records the move in the volume’s location history so you can trace where a book has lived over time.

3.5 Loans

Lending a volume

1. Open `/loans` and click **New loan**.
2. Scan or pick the **borrower**. If the borrower does not exist yet, click *Create borrower* inline.
3. Scan or pick the **volume** (V-code).
4. Optionally set a custom **due date**; otherwise mybibli derives it from the overdue threshold (chapter 2).
5. Confirm.

The volume’s current location is recorded on the loan row so a later return can restore it to the right shelf.

Returning a volume

1. Open `/loans` and find the active loan (filter on *Active* or scan the V-code).
2. Click **Return**.
3. Pick the post-return state (*Good*, *Worn*, ...).

4. Confirm. The volume snaps back to its pre-loan location automatically.

Overdue loans

Loans that exceed the overdue threshold turn red on the `/loans` page and on the home-page *Overdue* indicator. You can run `/loans?filter=overdue` for a focused list.

3.6 Library audit

An *audit* is the routine “walk the shelves and confirm everything is where the database says it is”. mybibli does not ship a dedicated audit mode — the equivalent workflow is:

1. Filter `/search` by the location you intend to audit, sorted by V-code.
2. Walk the shelf, scanning each V-code.
3. For any volume scanned that is *not* in the filtered list, mybibli surfaces a *Volume not in this location* message. Move it (section 3.4) and move on.
4. Any volume *in* the filtered list that you did *not* scan after walking the shelf is suspect — either misplaced or missing. Mark it for follow-up.

Tip

A quarterly audit catches misplaced volumes early — the longer a book has been on the wrong shelf, the harder it is to remember where it should be.

3.7 Shelving workflow

After a returns batch, the volumes pile up near the scanner station. mybibli’s *shelving* flow speeds up putting them back:

1. Pick up a volume and scan its V-code.
2. mybibli shows the volume page with the original location prominently displayed.
3. Walk the volume to that shelf. Scan the L-code on the shelf to confirm.
4. Repeat.

Confirming with the L-code is optional but recommended — it leaves a fresh entry in the volume’s location history that the next audit can rely on.

3.8 Wish list

The *wish list* is a separate list of books you’d like to buy, kept distinct from the owned catalog. It lives under `/wishlist` (see the navigation bar) and is available to Librarians and Admins.

3.8.1 Adding an entry

The *Add to wish list* form has two modes:

- **By ISBN.** Scan or type a 13-digit ISBN. mybibli looks the book up through the same metadata provider chain the catalog uses, then shows a preview card (title, author, publisher, year, cover). Click *Add to wish list* to save. Cover images come from the resolved provider's URL directly (e.g. Open Library Covers or Google Books); downloading the cover locally is on the roadmap.
- **Free-form.** Use this when you don't have an ISBN handy (a self-published title, a recommendation jotted down at lunch). Only *Title* is required; author, publisher and notes are optional.

3.8.2 Browsing at the bookstore

The list view collapses into a mobile-card stack on narrow viewports — designed to scan quickly with one thumb at a bookstore. Each card shows the title, the author (if known), the ISBN (if any), the publisher, and the date you added it.

To print a paper copy or save a PDF for offline use, click *Print* at the top of the list. mybibli renders a clean print-friendly page; use the browser's *Print to PDF* (File → Print → Save as PDF, or Ctrl/Cmd-P) to generate a PDF you can take to the bookstore.

3.8.3 Marking a book as bought

When you buy a book, open its wish list entry (or tap *Bought* directly from the list row), confirm in the modal, and the entry soft-deletes. The action is reversible from the Admin Trash panel for 30 days.

3.8.4 Auto-detect on catalog scan

When you catalog a freshly bought book through the standard scan flow (*/catalog*), mybibli checks whether its ISBN matches an active wish list entry. If it does, an emerald banner appears in the feedback area: *This ISBN is on your wish list — Remove it?* One click on the banner's *Remove from wish list* button opens the same *Mark as bought* modal, and confirming removes the entry. No silent deletions.

Chapter 4

Multiple media types

mybibli is not just for books. Four media types are first-class citizens:

- **Book** — the default. Single-volume novels, essays, reference works, children’s picture books.
- **Comic / BD** — single-issue comics or multi-position omnibus albums (e.g. a “Tintin anthology” that physically gathers volumes 1–5).
- **Audio** — CDs, vinyl records, audiobooks on physical media.
- **Film / series** — DVDs, Blu-ray boxsets, TV-series collections.

The media type drives three things:

1. **Which metadata provider is consulted first.** ISBNs route to BnF + Google Books + Open Library; films route to TMDb + OMDb; audio routes to MusicBrainz.
2. **Which fields are exposed on the title page.** Audio titles surface a track list; films surface director, runtime, age rating.
3. **The icon shown next to the title** in search results and the catalog list.

4.1 Books

The default media type. Catalog as described in section 3.2. Multi-volume series (encyclopaedias, trilogies, ...) are modeled with a *Series* entity in chapter 2.

4.2 Comics and BD

A single **omnibus** volume can contain several positions in a series. mybibli supports this with the *multi-position* feature: at volume creation time, you can set the volume’s *position range* (for example, “contains volumes 5–7”). Series-gap detection (chapter 3) understands these ranges and will not flag positions 5, 6 or 7 as missing once the omnibus is on the shelf.

4.3 Audio releases

Audio titles use MusicBrainz as the primary metadata source. The barcode on the back of a CD (EAN-13) resolves the same way as an ISBN.

- **Tracks.** MusicBrainz returns the track list; mybibli renders it on the title page.
- **Format.** CD vs vinyl vs cassette is captured on the volume row (not on the title) — two identical pressings on different media are two volumes of the same title.

4.4 Films and series

Films and TV series use OMDb and TMDb as the primary metadata sources. A series boxset is modeled as a *title* with a multi-position volume covering the disc range.

Tip

OMDb is free for low volumes (under 1000 requests per day); TMDb is free with attribution. Either covers a household-sized DVD / Blu-ray shelf. See chapter 5 for setup.

4.5 Picking the media type at scan time

After you scan an unfamiliar barcode, mybibli guesses the media type from the EAN prefix (the GS1 country prefix and the resolved provider's response). The guess is shown on the catalog page; you can override it before clicking **Add volume** if it is wrong. Re-typing a title later is a few clicks (`/title/<id>`) and the provider chain re-runs against the new type if you trigger *Re-fetch metadata*.

Chapter 5

Metadata providers

When you scan an ISBN or an EAN-13, mybibli does not look the title up in a local catalog — it queries one or more external *metadata providers* until one returns a clean answer. This chapter explains the chain, where to get API keys, and how to rotate them safely.

5.1 The provider chain

mybibli walks the providers in the order best-suited to the guessed media type:

Provider	Media types	Notes	Key?
BnF	Books	French Bibliothèque nationale catalogue; best for French-language titles.	no
Open Library	Books	Volunteer catalogue; broad English / multilingual coverage.	no
Google Books	Books	Commercial; good for trade non-fiction.	optional
MusicBrainz	Audio	Volunteer music metadata; very deep on CDs / LPs.	no
OMDb	Films, series	Free tier; rate-limited.	required
TMDb	Films, series	Free with attribution.	required
BDGest	Comics, BD	FR-language comic catalogue (web-scraped).	no

The first provider to return a *clean* record wins. “Clean” means: title present, at least one contributor, and a year. Partial matches are recorded but mybibli keeps walking the chain until it finds a clean one or runs out of providers.

Cover-image fallback

When the resolving provider returns metadata but no cover URL (BnF never carries image URLs in its UNIMARC payload; Google Books occasionally lacks `imageLinks` for self-published titles) mybibli falls back to the *Open Library Covers API* (<https://covers.openlibrary.org/>) keyed by ISBN. That covers index is broad and independent of Open Library’s catalogue completeness, so most books still get a cover even when their text metadata comes from a different provider. The fallback runs only when (a) the chain succeeded *and* (b) the resolving provider gave no

`cover_url`, so a title for which no provider has a cover anywhere will land cover-less without retry storms.

When no provider has a cover: upload your own

Some titles will never receive a cover from the automatic chain. In practice this is dominated by French, Swiss, and German trade and academic publishers (Dunod, Eyrolles, Slatkine, Mardaga, Kana, Casterman, and the like): Google Books returns nothing for their ISBNs and the Open Library Covers index does not hold them either. This is a structural gap in the free data sources, not a bug — no amount of re-fetching will fill it.

For those titles, **manual upload is the realistic path**. On a title's detail page a Librarian (or Admin) sees an *Upload cover* button under the cover area; titles with no cover get a more prominent call-to-action and a one-line reminder of this gap. Accepted formats are JPG, PNG, WebP and GIF, up to 10 MiB. The image is resized and stored locally, and a manually uploaded cover is protected from being overwritten by a later automatic re-fetch.

To find the titles that still need a cover, use the *Titles without cover* filter on the home page (Librarian and above), or follow the *Review titles without a cover* link shown on the Admin page's Health tab after a bulk cover re-fetch.

5.2 API keys

Three providers require an API key:

Google Books

Optional but raises the daily request budget significantly. Without a key the Google Books provider stays disabled — the chain still works (BnF → Open Library are key-less), but English-language trade non-fiction that BnF misses will land with empty metadata.

The Google Cloud Console flow trips up many first-timers because the API-key creation step is buried two menus deep. Walk through this exactly:

1. Open <https://console.cloud.google.com/> and sign in.
2. Top-bar project picker (next to the *Google Cloud* logo) → *New project* → name it anything (e.g. *mybibli-metadata*) → *Create*. Wait a few seconds for the project to be selected.
3. Open the hamburger menu (top-left, three horizontal lines) → *APIs & Services* → *Library*. Search for **Books API**, click it, then *Enable*.
4. Still under *APIs & Services*, open *Credentials*. Click *+ Create credentials* → *API key*. A modal pops up with a freshly generated key — copy it now.
5. Recommended: click *Restrict key* → under *API restrictions* pick *Restrict key* → tick only *Books API*. Leave *Application restrictions* on *None* (mybibli does not send a Referer header). Save.
6. In mybibli: `/admin > System > Metadata Providers` → paste into the **Google Books API key** field → *Save*. The next scan picks it up — no restart needed.

To verify the key is active, scan an English-only ISBN that BnF will not have (*Effective Java*, 9780134685991, is a good canary). The metadata should populate within a few seconds. If it stays empty, see *Troubleshooting* at the end of the chapter.

OMDb

1. Visit <https://www.omdbapi.com/apikey.aspx>.
2. Pick the free tier (1000 requests/day).
3. Confirm via the activation email you receive.
4. Paste the key into mybibli.

TMDb

1. Sign up at <https://www.themoviedb.org/>.
2. Open *Settings* → *API* and request a key for *Personal / non-commercial use*.
3. Paste the key into mybibli.

5.3 Where the keys live

Keys travel through two layers:

1. **Environment variables** (`GOOGLE_BOOKS_API_KEY`, `OMDB_API_KEY`, `TMDB_API_KEY`). These are migrated *once* on boot into the `settings` table — they are essentially a deployment-time bootstrap.
2. **The settings table**, exposed through `/admin > System > Metadata Providers`. This is the live source of truth at request time. Rotating a key from the UI takes effect on the very next metadata fetch.

The implication is asymmetrical:

- If you set a key only in the environment, the first boot copies it into the database and afterwards the database wins.
- If you set a key only in the UI, the environment is never consulted again.
- If you set both and then *Clear* the key from the UI, the database value goes empty — but on the next boot, the environment value is migrated back in. To make the UI clear durable, also remove the environment variable from `.env` / `docker-compose.yml`.

5.4 Provider health

`/admin > Health` pings every registered provider every five minutes and reports their status as a green check or a red cross. A red cross does not always mean the provider is down — it can also mean the API key is missing or invalid. The Health tab includes the last-checked timestamp and the HTTP status code for diagnostics.

5.5 Rate limits and back-pressure

mybibli applies a small per-provider rate limit (a couple of requests per second). Bulk-cataloging a hundred volumes in a row will not trigger remote 429 responses. If a provider does return 429, the request is logged and the next provider in the chain is consulted instead — the catalog flow never blocks waiting for a rate-limited provider.

Warning

Never commit a real API key to git. The `.env` file is in `.gitignore` for this reason; `.env.example` carries only placeholders. If a key leaks, rotate it at the provider console immediately — pasting a new key into mybibli takes effect within seconds.

5.6 Troubleshooting empty metadata

If a freshly scanned ISBN ends up with empty title / author / description fields, walk this list before suspecting a bug:

1. **Provider status in** `/admin > Health`. If Google Books or Open Library show a red cross, those providers will be skipped on every fetch. Re-check the last-checked HTTP code — 403 or 401 usually means a missing or invalid API key, not a remote outage.
2. **Language of the title.** BnF is the primary provider and is FR-leaning. For English-only books, the chain falls through to Google Books — which means the key (see above) is the unblocker.
3. **ISBN format.** mybibli normalises EAN-13 vs ISBN-10 internally, but some providers index only one of the two. A second scan a few minutes later sometimes hits a different fallback path.
4. **Manually edit and save.** You can always edit title, genre, contributors etc. by hand from the title's detail page — saving works even when no provider returned anything. The genre `Non classé` is the default for any title that came in without classification data.

If the fields are still empty after configuring the right keys and waiting a few seconds, open an issue on the project repo with the failing ISBN — the maintainers add provider mocks for reproducible cases.

Chapter 6

Roles and permissions

mybibli has three user roles ordered by privilege:

1. **Anonymous** — anyone who lands on the site without a session cookie. Read-only access to the catalog. Cannot loan, edit or delete anything.
2. **Librarian** — authenticated user whose role is `librarian`. Day-to-day cataloging, loans, returns, location moves. Cannot manage users or change system settings.
3. **Admin** — authenticated user whose role is `admin`. Everything a Librarian can do, plus the `/admin` panel: users, reference data, trash, system settings.

6.1 Who can do what

Action	Anon	Lib	Admin
Browse / search the catalog	yes	yes	yes
View title and volume pages	yes	yes	yes
View location tree	yes	yes	yes
Scan / catalog a new title	no	yes	yes
Add or remove a volume	no	yes	yes
Move a volume between locations	no	yes	yes
Edit a title's metadata	no	yes	yes
Create / edit borrowers	no	yes	yes
Register / return loans	no	yes	yes
Edit storage location tree	no	yes	yes
Open <code>/admin</code>	no	no	yes
Create / deactivate users	no	no	yes
Edit reference data	no	no	yes
Restore from trash	no	no	yes
Permanently delete from trash	no	no	yes
Edit system settings	no	no	yes

6.2 Sessions and the inactivity timeout

After login, mybibli sets a session cookie that travels with every subsequent request. The cookie is:

- `HttpOnly` — not exposed to JavaScript;
- `SameSite=Lax` — sent on top-level same-site navigations but not on cross-site requests;
- `Secure` (only when `MYBIBLI_COOKIE_SECURE` is set to `true` — see chapter 1).

The session has an **inactivity timeout** (default 30 minutes, configurable from `/admin > System > Loans`). Every authenticated request resets the countdown; if no request lands within the timeout window, the session is wiped server-side and the next request redirects to the login page.

A small *keep-alive toast* appears in the corner of the screen roughly two minutes before expiry — click it to extend the session without leaving your current page.

6.3 The last-active-admin guard

mybibli refuses to leave the installation without any active admin. The guard applies to two actions:

- **Deactivating a user.** If the user is the only active admin, the request returns a 409 conflict and the deactivation is refused.
- **Permanently deleting a user from the trash.** Same rule.

Self-deactivation is also blocked unconditionally (you cannot deactivate the account you are currently logged in as).

Warning

The guard counts only *active* admins. If you mass-deactivate admins and the guard fires, undo by reactivating someone from the trash tab. There is no recovery path through the wizard or the database without manual SQL surgery.

6.4 Password reset

mybibli does not implement password-reset emails — household / small-community deployments rarely benefit from the email plumbing. An admin can reset another user's password by:

1. Open `/admin > Users`.
2. Pick the user, click **Deactivate**.
3. Click **Reactivate** from the trash tab.
4. Re-create the user with a fresh password.

A user who knows their own password can change it from the `/account` page.

6.5 Anonymous sessions

Anonymous visitors also get a session row (used for CSRF protection and pending HTMX updates). Anonymous sessions expire after seven days of inactivity and are garbage-collected by a daily background task — they do not affect logged-in users.

Chapter 7

Backup and restore

A self-hosted library has no cloud safety net. This chapter covers what to back up, how often, and how to restore.

7.1 What to back up

Three artifacts together form a complete mybibli backup:

1. **The MariaDB database.** Holds every title, volume, location, contributor, loan, user, audit log and setting. Lose this and you start from zero.
2. **Cover images.** Downloaded during cataloging and persisted in the Docker volume `mybibli_covers` (or in your host bind-mount path if you switched per section 1.3). Lose this and the catalog still works, but every title page shows a generic placeholder until you re-fetch metadata title-by-title (issue #213 — pre-1.1.4 installs without this volume had covers disappear on every container upgrade).
3. **The `.env` file.** Holds passwords (database, API keys not yet migrated). Lose this and a re-deploy needs you to re-enter every credential.

The Docker images themselves are *not* backup-critical — you can always pull them again from Docker Hub.

7.2 Backup procedure

Database — bundled MariaDB

From the host running the compose stack:

```
docker compose exec -T db mariadb-dump \  
  -u root -p"$MYSQL_ROOT_PASSWORD" \  
  --single-transaction --routines --triggers \  
  mybibli \  
  > backups/mybibli-$(date +%Y%m%d-%H%M%S).sql
```

The `--single-transaction` flag is important: it dumps a consistent snapshot without locking the table for the duration of the dump. `--routines --triggers` preserves any stored routines (mybibli does not currently use any, but the flag is cheap and future-proof).

Compress the dump if disk space is a concern:

```
gzip backups/mybibli-*.sql
```

Database — external MariaDB

Run `mariadb-dump` directly against the external server. Synology DSM users can also schedule a *Hyper Backup* task that targets the MariaDB package — it produces a binary backup that restores faster than a SQL dump.

Covers and `.env`

For a host-mounted covers directory (bind-mount installs), a plain `tar` archive is enough:

```
tar czf backups/mybibli-files-$(date +%Y%m%d).tar.gz \
.env "$COVERS_HOST_PATH"
```

For the default install (named Docker volume), back up the volume through Docker — the host path of a named volume varies by Docker version and is not stable, so don't try to `tar` it directly:

```
docker run --rm \
-v mybibli_covers:/source:ro \
-v "$(pwd)/backups:/backup" \
alpine \
tar czf "/backup/mybibli-covers-$(date +%Y%m%d).tar.gz" -C /source .
tar czf "backups/mybibli-env-$(date +%Y%m%d).tar.gz" .env
```

The throwaway `alpine` container mounts the named volume read-only, the host `backups/` directory writable, and writes the tarball. Same shape works for restoring — see the Restore section below.

7.3 Recommended cadence

- **Daily** — automated database dump. Even on a seldom-changed catalog, a daily dump means worst-case data loss is bounded to one day.
- **Weekly** — covers and `.env`. These change only when you catalog new titles or rotate keys; weekly is plenty.
- **Off-site copy** — at least monthly. Rsync the `backups/` directory to a remote NAS, a USB drive you swap in / out, or a cloud-storage bucket. A fire that consumes the home server should not also consume your backups.

Tip

A backup that has never been restored is not a backup. Once a quarter, restore your latest dump into a throwaway docker environment and verify the title and loan counts match what `/admin > Health` reports on production.

7.4 Restore procedure

Database

For the bundled-database setup:

```
docker compose stop mybibli
docker compose exec -T db mariadb -u root -p"$MYSQL_ROOT_PASSWORD" \
    -e "DROP DATABASE mybibli; CREATE DATABASE mybibli CHARACTER SET utf8mb4;"
zcat backups/mybibli-20260513-1200.sql.gz | \
    docker compose exec -T db mariadb -u root -p"$MYSQL_ROOT_PASSWORD" mybibli
docker compose start mybibli
```

For external MariaDB, the same `mariadb` client commands work against the external server.

Covers and `.env`

For a host-bind-mount install:

```
tar xzf backups/mybibli-files-20260513.tar.gz
```

For the default named-volume install, restore through a throwaway container in the symmetric shape to the backup:

```
docker compose stop mybibli
docker run --rm \
    -v mybibli_covers:/target \
    -v "$(pwd)/backups:/backup:ro" \
    alpine \
    sh -c "rm -rf /target/* && tar xzf /backup/mybibli-covers-20260513.tar.gz -C /target"
tar xzf backups/mybibli-env-20260513.tar.gz
docker compose start mybibli
```

Restart the stack so the `mybibli` container picks up the restored `.env`.

7.5 Migration to a new host

The backup procedure is also the migration procedure:

1. On the old host: produce a fresh database dump and tar archive.
2. Transfer both to the new host.
3. On the new host: install Docker, pull the compose file from the release matching the dump's version, place `.env` and `covers/` alongside, restore the database, start the stack.
4. Verify by counting titles on `/admin > Health` and spot-checking a few title pages.

Warning

The dump is tied to the *schema* version. If the source host ran 1.1.0 and the new host pulls 1.5.0, the new host's migrations will upgrade the restored database in place — that is fine. The reverse (restoring a 1.5.0 dump onto a 1.1.0 install) is *not* supported and will fail at boot. Always restore on a host running the same or newer `mybibli` release.

Chapter 8

Upgrade and migration

Routine upgrades are designed to be safe. The exception is the 1.0.x to 1.1.0 jump, which is a one-time breaking change documented below.

8.1 Routine upgrades (1.y to 1.z)

1. Read the release notes for every version between your current and the target version.
2. Take a fresh **database backup** (chapter 7).
3. Update the image tag in `docker-compose.yml` (or repull `:latest` if you track that tag).
4. Run `docker compose up -d`. The new container boots, runs any new migrations against the existing database, and starts serving requests.
5. Watch the logs for migration errors: `docker compose logs mybibli | grep -i migrate`.

Migrations are append-only — schema changes are forward compatible. A 1.4 install upgraded to 1.5 keeps every row from 1.4 and applies new tables / columns in place.

Tip

For homelab installs, the `:latest` Docker tag is a reasonable default. For production-style setups (a community library, an association), pin to an explicit version tag so an unattended pull does not bring in a release you have not read about.

8.2 Skipping intermediate versions

Going from a much older install (say, 1.3.x) directly to the current release is supported and safe for the database. The migration runner applies every pending migration in timestamp order at boot, regardless of how many you skip. Schema migrations are purely additive — no `DROP COLUMN`, no `DROP TABLE` in any release — and the four data backfills that exist (e.g. the 1.1.0 CSRF token seed, the 1.1.8 `manually_edited_fields` cleanup) are idempotent.

What still applies on a multi-version jump:

- **Back up before upgrading.** There is no automatic rollback. A migration that fails mid-execution leaves the database in an inconsistent state; the recovery path is the pre-upgrade backup (chapter 7).

- **Pre-1.1.4 cover JPGs are gone.** Until 1.1.4 the `mybibli_covers` volume was not declared, so cover images written before that version were lost on every container upgrade. Skipping straight to 1.7 does not recover them — re-fetch via the admin **Bulk re-fetch missing covers** action (added in 1.2) or the per-title **Re-fetch metadata** button.
- **The first boot may take a few minutes.** A jump that runs 20+ migrations against a large catalog is silent at the UX level — wait for `/health` to respond before worrying.

The 1.0.x to 1.1.0 jump is the one exception that needs the dedicated procedure below.

8.3 The 1.0.x to 1.1.0 jump

Warning

This is a one-time breaking change. 1.0.x installs are subject to the issue documented in [#173](#): the seed migrations create the credentials `admin/admin` and `librarian/librarian` on every fresh install, including production, and silently bypass the first-launch setup wizard.

The recommended path depends on whether your 1.0.x install holds real data:

If your 1.0.x install is empty (no titles cataloged)

The simplest path: wipe the database and reinstall on 1.1.0+.

1. Stop the stack: `docker compose down`.
2. Remove the database volume: `docker volume rm mybibli_mybibli_db_data` (or the equivalent volume name on your host — `docker volume ls` lists them).
3. Pull the 1.1.0 image: `docker compose pull`.
4. Update the image tag in `docker-compose.yml` if you pinned a previous version.
5. Start: `docker compose up -d`.
6. Open `/setup` and walk through the wizard.

If your 1.0.x install holds real data

1. Take a database backup (chapter [7](#)).
2. Take a covers + `.env` backup.
3. Log in as the seeded admin and rotate the password *before* you upgrade. The 1.1.0 upgrade flow itself is safe, but rotation closes the only window where the default credentials could leak.
4. Update the image tag in `docker-compose.yml` and run `docker compose up -d`.
5. The 1.1.0 boot detects the existing admin and does *not* re-enable the wizard — the rotated admin remains the active admin.

The new MYBIBLI_SEED_DEV_USERS variable

Starting with 1.1.0, the seed migration is gated by MYBIBLI_SEED_DEV_USERS. The default is 0 (no seed). Only set it to 1 for local-development workflows or automated tests — never in production. The first-launch wizard becomes the canonical onboarding path.

8.4 Reading the release notes

Every mybibli release publishes a GitHub Release page with:

- a one-paragraph summary of the change set;
- the migration impact (none / additive / breaking);
- a link to the diff against the previous tag;
- attached PDFs of this manual (English + French) once the release-time PDF build is wired into CI.

Read the notes before upgrading. Breaking changes are tagged with a prominent *Breaking* label.

8.5 What’s new in 1.8.4

Version 1.8.4 is a three-issue hardening and quality closure on v1.8.3. *No operator action required*; the one bookkeeping migration runs automatically. Routine `docker compose pull && docker compose up -d`.

- **#21 — database-layer hardening.** The connection pool now sets explicit acquire and idle timeouts, so a request waiting on a saturated pool fails fast instead of hanging and a NAS that idles overnight reclaims stale connections. The `/health` endpoint was promoted from a bare “process is alive” stub to a real readiness probe: it runs a `SELECT 1` and returns 503 when the database is unreachable, so an external monitor (or your container healthcheck) can tell a live process from one that has lost its database.
- **#55 — admin users-panel hardening.** The user list’s page parameter is now bounds-checked (a hostile or accidental `?page=` value can no longer overflow the offset calculation); attempting to deactivate your own account or the last remaining admin now shows a properly translated message instead of an internal key; and the create/edit forms validate username length (3–255 characters) up front rather than surfacing a database error.
- **#22 — internal test-suite quality.** A consolidation pass on the end-to-end test suite (shared login and lookup helpers, deterministic record resolution, and parallel-safe unique test data). No user-visible change — listed for completeness.

8.6 What’s new in 1.8.3

Version 1.8.3 is a second consecutive open-issue closure sweep on v1.8.2, retiring five GitHub issues. *No schema migration that needs operator action*; routine `docker compose pull && docker compose up -d`.

- **#41 — session resolver consistency.** When an admin soft-deletes a user while that user is browsing, or when an authenticated session is past its inactivity timeout, the underlying database row was previously left intact and the Session in memory ended up

in a half-state (an authenticated identity that nonetheless denied authentication). The resolver now soft-deletes the orphan row in both cases and mints a fresh anonymous session, so the next request always sees a consistent state.

- **#20 — shorter pagination on the home page.** Previously every page number from 1 to N was rendered. On a large catalog (10 000 titles = 400 pages) this filled the page with 400 anchors and made the tab order useless. The nav now shows at most eleven links: first, last, the current page ± 3 , and a literal `...` where a gap was collapsed.
- **#37 — long-lived tabs no longer 403 after a background session rotation.** When the background anonymous-session purge (or the new #41 cleanup) runs, the tab's stored CSRF token used to drift from the server's stored token, and the very next form submission failed with a *session expired* feedback. The server now signals the rotation on its next response; the JavaScript layer rewrites the page's CSRF token in place, and the next mutation succeeds.
- **#23 — metadata provider robustness.** A configuration of `session_inactivity_timeout_hours = 0` (silently accepted before) would have logged everyone out on every request; the loader now refuses zero. Provider rate-limit detection (HTTP 429) uses a structured signal instead of pattern-matching the error message, so future provider message changes can't slip past unnoticed. And a sloppy upstream response that returns an empty string for a missing field no longer overwrites a real value in the database.
- **#29 — security and i18n hygiene.** Replaced the unmaintained YAML parser dependency (`serde_yaml` 0.9, security advisory RUSTSEC-2024-0320) with its actively maintained fork (`serde_yaml_ng` 0.10). New continuous-integration check rejects any future translation that accidentally introduces raw HTML markup into the locale files, closing a latent cross-site-scripting surface.

8.7 What's new in 1.8.2

Version 1.8.2 is an open-issue closure sweep on v1.8.1 with *zero user-visible behaviour change*. The driving objective is shrinking the open-issue list: seven GitHub issues retired in one patch (#151, #365, #370, #374 directly, plus parents #26, #30 and #18 via the carrier issues). No schema migration, no settings change; routine `docker compose pull && docker compose up -d`.

- **#151 — spec / impl alignment for the nav bar.** Three planning sources disagreed on which links Anonymous users see. Story 9-18 had already shipped Option A; this patch aligns the UX specification and the Epic 9.18 acceptance criteria so future readers don't reopen the debate. Doc-only.
- **#365 — session and CSRF token hygiene.** The function that mints session tokens was defined in three different places, with subtly drifting comments. They are now one canonical implementation. Same move for the CSRF token generator. A latent footgun in the test-only `Session` extractor path that returned an empty CSRF token is closed.
- **#370 — typed error variants.** When two users edit the same record at once, the resulting 409 message now reads “Another user just modified this <entity>” (or “This <entity> was just deleted” if the row was soft-deleted between read and write) instead of a generic “This record was modified by another user”. Internal cleanup also moves the `FeedbackEntry` HTML builder out of a routes module so the error layer no longer depends on it.
- **#374 — CI guards + accessibility leftover.** Two new automated tests guard the

release pipeline: one reproduces the Python parser that extracts the Cargo.toml version from the release tag, the other replays every database migration on a fresh DB and asserts the seed data still satisfies the schema’s NOT NULL / unique / foreign-key constraints. On the user side, eight hardcoded English `aria-label` strings on the home-page filter chips now ship in all four UI languages, and the catalog form pages re-enable the WCAG colour-contrast audit after a global placeholder colour bump.

8.8 What’s new in 1.8.1

Version 1.8.1 is a focused polish patch on v1.8.0. It closes one production-surfaced bug and four code-quality cleanups carried from the issue queue. *No schema migration that needs operator action*; the migration to add the new `is_primary` column on `contributor_roles` runs automatically at boot. Routine `docker compose pull && docker compose up -d`.

- **#300 — phantom-volume guard.** A subtle bug since v1.0: the catalog session keeps the last-scanned ISBN’s title in a sticky *current title* slot that was never cleared. If you scanned a V-code for a different physical book without first re-scanning its ISBN, the new volume was silently attached to the previous title — so a title could end up with two or three volumes you didn’t actually own. The fix surfaces a confirmation modal the moment you scan a V-code for a title that already has volumes: “Add another copy of <title>?”. Click *Add another copy* for the legitimate multi-exemplaire flow; *Cancel* to discard the scan and re-scan the new book’s ISBN first.
- **#216 — readable error pages from direct navigation.** Until now, navigating directly to a URL that errors (e.g. a stale bookmark to a deleted title) rendered the error feedback fragment as a styleless block on a blank page. A new middleware wraps these responses in a minimal HTML shell so direct-nav errors are now legible and offer a Home link to backtrack. HTMX requests (in-app interactions) are unaffected.
- **#19 — rename-resilient author role.** The author-first sort on contributor lists and the metadata-fetch author lookup used to match on the literal name “Auteur”, so renaming the role from the admin reference-data panel silently broke them. A new `is_primary` marker on the seeded role replaces the string check — rename freely.
- **#87 — per-target OOB swap mode.** Forms that update multiple things in one save (e.g. the provider-keys form updating Google Books, OMDb and TMDb together) can now stack their per-sub-action feedbacks instead of overwriting them. Default behaviour is unchanged.
- **#18 — sortable column headers announce sort state.** The five sortable columns on the home browse table now carry the `aria-sort` attribute, so screen-reader users hear which column is sorted and in which direction.

8.9 What’s new in 1.8.0

Version 1.8.0 is the seventh themed minor, “Covers, honestly”. It makes the manual cover-upload path discoverable for the titles that no free metadata provider will ever cover, and completes the internal `base_context()` refactor begun in 1.7.11. *No schema migration, no settings change*; routine `docker compose pull && docker compose up -d`.

- **#350 — the manual cover-upload affordance is now visible.** A title’s detail page shows a prominent “Upload a cover” button plus a one-line explainer when it has no cover: many French, Swiss and German publishers (Dunod, Eyrolles, Slatkine, Mardaga, Kana,

Casterman) are indexed by neither Google Books nor Open Library, so manual upload is the realistic path. A manually uploaded cover is protected from being overwritten by a later automatic re-fetch. See chapter 5 (*Metadata providers*) for the full background on this publisher gap.

- **#355 — “Titles without cover” filter on the home page** (Librarian and above). Lists every title with no cover so you can walk them one by one and upload manually. It deliberately includes titles with no ISBN/ISSN/UPC — those can never be auto-fetched, so they need manual upload most.
- **Admin Health follow-up.** After a bulk cover re-fetch, the Health tab shows a “Review titles without a cover” link straight to that filtered list.
- **#108 / #354 — internal quality.** The title-card markup is now a single shared partial (no more three-site duplication), and every full-page handler builds its shared header/navigation fields through one helper — an invisible-to-you change that keeps the four UI languages and the navigation bar consistent across every page.

8.10 What’s new in 1.7.11

Version 1.7.11 bundles one user-visible WCAG AA contrast fix and three code-review-finding clean-ups, plus a companion contrast fix surfaced during the v1.7.11 local Rule 13 gate. *No schema migration, no settings change*; routine `docker compose pull && docker compose up -d`.

User-visible (accessibility):

- **#345 — trash count admin badge contrast bumped to WCAG AA.** The badge in `templates/components/admin_tabs.html` paired `bg-red-500 (#fb2c36)` with `text-white`, yielding only 3.8:1 — below the WCAG 2.2 AA 4.5:1 threshold for normal text. Switched to `bg-red-700 (#b91c1c, 6.61:1)` — comfortably above AA and above the AAA 7.0 threshold. Badge stays visually a “warning red”. New `templates_audit::templates_forbid_bg_red_500_with_text_w` regression guard panics on any future template pairing the two — the combination is intrinsically inaccessible, so a static audit is the right enforcement level.
- **#352 (companion) — locations tree L-code + node-type spans contrast.** Same shape as #345, different element. The L-code span on `/locations` (e.g. L1004) and the V-code span on `/location/:id` used naked `text-stone-400 (#a8a29e)` on `bg-stone-50 (#fafaf9)` light mode = 2.41:1. Switched to `text-stone-600 dark:text-stone-400` for 7.5:1 on light and 8:1 on dark.

Code-review-finding clean-ups (developer-visible):

- **#30 — CI/tooling.** `.github/workflows/release.yml` now extracts the Cargo.toml version via Python’s `stdlib tomlib` (3.11+, available on ubuntu-latest) instead of `grep`, making it robust against UTF-8 BOM, CRLF, single-quoted values, and a future `[workspace.package]` section. `tests/e2e/tsconfig.json` adds `noUncheckedIndexedAccess + exactOptionalPropertyTypes`; the 12 surfaced narrow non-null assertions on regex-match indexing are backfilled. The CI typecheck (e2e job) now catches any future violation. Sub-items 2 (tag-mismatch smoke test) and 3 (seed-migration drift audit) deferred to their own focused issues.
- **#340 — TEST_MODE seed-overdue-loan endpoint.** New `POST /debug/seed-overdue-loan` (TEST_MODE-gated, Admin role check, mirror of `debug_set_session_timeout`) that back-dates `loans.loaned_at` via direct SQL — UI seeding cannot travel time. Closes the overdue

branch of CR #119: `home.spec.ts` can finally seed a real overdue loan and assert the indicator end-to-end, removing the silent `test.skip` that previously green'd a missing-indicator regression.

- **#35 slice — `base_context()` DRY helper.** New `BaseContextFields` struct in `src/utils.rs` carries the 20 fields every full-page template struct repeats (`lang`, `role`, `csrf_token`, 10 nav strings, ...). Built by `base_context(&session, locale, current_page, &uri, session_timeout_secs)` and applied to `loans_page` / `borrowers_page` / `series_index` handlers as proof-of-pattern. Reduces i18n key duplication: a typo on `nav.catalog` now surfaces in one place, not 17. Remaining ~14 full-page handlers + the deferred sub-items (`generate_session_token` dedup, `Session::anonymous_with_token` fallback removal) are tracked in a follow-up issue for incremental migration.

8.11 What's new in 1.7.10

Version 1.7.10 is a single-fix patch closing a silent-data-loss vulnerability. *No schema migration, no settings change*; routine `docker compose pull && docker compose up -d`.

User-visible:

- **#347 — manually uploaded covers are no longer overwritten by the metadata-fetch chain.** The async background fetch after the initial scan, and the bulk-refetch admin action in a race with a fresh upload, could silently overwrite a manually uploaded cover (the v1.7.6 #335 safety net) because `update_cover_image_url` did an unconditional UPDATE that ignored `manually_edited_fields`. The failure path was equally vulnerable, blanking the manual cover to NULL when the provider download failed. The fix is a 22-line guard pattern identical to `do_update`'s with 4 integration tests covering success / download-failure / regression / specificity.

Companion finding (no code change):

The broader cover-miss gap on the production NAS (73/147 titles, ~50%) is structural and not closable from the provider chain. Google Books returns `totalItems=0` for the dominant FR/CH/DE tech publishers (Dunod, Eyrolles, Slatkine, Mardaga), and Open Library Covers does not index their ISBNs either. The v1.7.9 piste-1 fix (highest-resolution `imageLinks` variant) was prod-validated as ineffective at moving that count. Manual upload (#335) remains the only realistic safety net for those titles. A v1.8 UX track to make the upload affordance more visible is being scoped.

8.12 What's new in 1.7.9

Version 1.7.9 is a cover-quality + admin-UX + 4-hardening bundle. One user-visible cover-quality fix, one admin-visible config surface, and four code-review-finding hardenings. *Two new K/V settings rows (seeded at the prior hardcoded defaults), no breaking changes*; routine `docker compose pull && docker compose up -d`.

User-visible:

- **#333 piste 1 — Google Books cover URL prefers the highest-resolution `imageLinks` variant.** The provider used to read only `thumbnail` from the Google Books response, even when the payload also carried `small` / `medium` / `large` / `extraLarge`. The downstream 400 px Lanczos resize in `CoverService::process_and_save_bytes` then worked from a lossy source. The new fallback chain (`extraLarge -> large -> medium -> small -> thumbnail`

-> `smallThumbnail`) picks the best available variant per title, so existing titles attach a sharper cover after re-fetch and the rare title that only ships `medium` upward but no `thumbnail` finally attaches a cover at all. Closes the (1) audit piste of the 50% prod miss-rate investigation; piste 3 manual upload already shipped in v1.7.6 #335; piste 2 BDGest cover-only fetch by ISBN remains open as a separate cycle.

Admin-visible:

- **#334 — metadata-chain + provider-health timeouts in /admin > System.** Two timeouts that were previously hardcoded (per-provider chain timeout) or env-var-only (provider-health probe timeout) now appear inside the Metadata Providers block of the System tab. Both are bounded 1..=60 seconds, persisted in the K/V `settings` table (migration 20260526075659), and hot-reload through `Arc<RwLock<AppSettings>>` so the next chain run / next ping round picks up the new value with no restart. The legacy environment variables `MYBIBLI_METADATA_CHAIN_PROVIDER_TIMEOUT_SECS` and `MYBIBLI_PROVIDER_HEALTH_TIMEOUT_SECS` are honored as one-shot bootstrap via `config::migrate_legacy_env_vars` (mirrors the v1.7.1 #308 log-level pattern). Useful for deployments with slow DNS+TLS (rural fiber, VPN, satellite).

Hardening (no user-visible behaviour change):

- **#46 — anonymous-session purge survives crash-loops.** The daily purge of anonymous-session rows older than 7 days previously slept a fixed 24h after every boot. On a host that crash-looped or rolling-deployed faster than 24h the purge never fired and the `sessions` table grew unbounded. The next-run delay is now computed from a persisted `last_anonymous_session_purge_at` K/V row (migration 20260526075700): if 24h have already elapsed since the previous run, the purge fires immediately (catch-up); otherwise it sleeps the remainder. Spec §Task 4.3 contract preserved for fresh installs.
- **#28 — observability trio.** (a) `LoanService::register_loan` retry trail now distinguishes a non-transient business-rule failure that follows a transient retry (e.g. concurrent soft-delete between attempts that flips the second attempt into a `BadRequest`) — it used to read in the log as “warn!(“retrying”)” immediately followed by the business-rule error, implying the retry caused it. (b) The 4 catalog “fire-and-forget” `fetch_metadata_chain` spawn sites now go through a new `tasks::metadata_fetch::spawn_fetch` wrapper that captures panics from the inner future and logs them via `tracing::error!` instead of dropping the `JoinHandle` silently. (c) `CLAUDE.md` single-tenant retry note now explicit that `sqlx::Error::PoolTimedOut` is intentionally *not* retried.
- **#24 — templates_audit hardening pass.** Five edge cases closed: (a) `strip_html_comments` rewritten as a state machine — an unterminated `<!--` no longer consumes the rest of the source file (previously any inline-markup violation past an accidentally-truncated comment was invisible to the audit). (b) The same function now slices the source `&str` instead of casting UTF-8 bytes to `char`, so accented characters in failure-report snippets render correctly. (c) New `strip_svg_inner_blocks` masks SVG inner content so an inline `<svg><style>...</style></svg>` can no longer false-positive the bare-`<style>` check. (d) The `<script>` and `<style>` attribute regex now uses a quote-aware grammar `(?:[>'"|"["]*'|'["']*')*` so a literal `>` inside a quoted attribute value cannot terminate the match early. (e) The `src=/type=` script-bypass filter behaviour was acknowledged as imprecise but not exploitable; no code change needed.
- **#196 — home-search typing-race flake.** The long-standing flake at `tests/e2e/specs/journeys/home-` (6 epic retrospectives in a row) finally closed. `simulateTyping` switched from `locator.pressSequentially`

to `keyboard.type` after an explicit `focus()` — the former re-focuses the element between key actions and could drop a keystroke under default-worker parallelism. The URL-state poll on the test was bumped from 2s to 5s to give the debounced search + HTMX swap + `hx-push-url` chain headroom under load.

8.13 What's new in 1.7.8

Version 1.7.8 is a hardening patch with one user-visible UX fix and three code-review-finding hardenings. *No schema migration, no breaking changes*; routine `docker compose pull && docker compose up -d`.

User-visible:

- **#136 — concurrent-delete modals surface feedback instead of silently no-op'ing.** Two librarians with the same `/borrower/:id`, `/contributor/:id`, or return-loan modal open used to see the loser's Delete click vanish silently: the modal handler returned 404, the default `AppError::NotFound` retargeted to `#feedback-list`, and these pages don't declare that slot — HTMX swap silently dropped. The three modal-trigger GET handlers now return 200 + inline feedback fragment + `HX-Retarget` to the page's existing slot (`#borrower-feedback`, `#contributor-feedback`, or the `?target=` resolution for the return-loan flow). The second librarian sees a localized "already deleted by another session" toast in all four supported locales.

Hardening (no user-visible behaviour change):

- **#39 — atomic login session swap.** `services::auth::authenticate_session` previously ran the INSERT new session row and the UPDATE prior anonymous session row as two independent queries. A partial failure between them left an orphan anonymous session row alongside the new authenticated one (both carried valid CSRF tokens for the same browser). Both queries now run inside a single `sqlx::Transaction` — atomic swap or full rollback.
- **#47 — CSRF test suite coverage gaps.** Four high-value Rust integration tests added: HTMX-branch language POST, form-field-token logout (the nav-bar shape), stale-token replay against a rotated session, and a 403 body assertion proving the rejection envelope carries the `FeedbackEntry` markup. The remaining 5 E2E items (parallel-pollution teardown + selector tightening) are deferred to a follow-up E2E pass.
- **#217 — modal Confirm scope tightening.** `modal.js::originatesFromConfirm` used to return true for ANY `<form>` descendant of the open dialog. A future modal with a nested form would have tagged the nested form's submissions with `X-Modal-Confirm: true` and tripped the server-side `ModalConfirmRetargetGuard`. The predicate now matches only the dialog's primary form (`dialog.querySelector("form")`) — the macro's Confirm form by construction. An E2E regression-guard injects a secondary form into an open dialog and asserts the header is NOT sent.

8.14 What's new in 1.7.7

Version 1.7.7 is a test-rigor patch bundling one user-visible enhancement and seven code-review-finding clean-ups closing silent-no-op E2E patterns and locking concurrency / locale-resolution contracts. *No schema migration, no breaking changes*; routine `docker compose pull && docker compose up -d`.

User-visible:

- **#336 — series detail card layout.** The `/series/:id` page used to list each position as a compact coloured cell that only carried the volume number. Now every filled position renders as a card with the cover thumbnail at top (or the media-type icon fallback when no cover has been downloaded), a "Vol. N" prefix in the accent color, the title (truncated), and the primary contributor. Missing positions render as dashed-border gap cards that still carry the "Vol. N" so collectors see at a glance which slots are still gaps. The grid is responsive (2 columns on mobile scaling to 6 columns on extra-wide screens) and preserves the ARIA `role="grid" / role="gridcell"` contract so screen-reader navigation is unchanged.

Test rigor (no user-visible behaviour change):

- **#43 — DB-snapshot helper on contributor POST rejection.** A new `captureEntityCounts()` Playwright helper proves that an anonymous-POST 403/303 rejection left the DB untouched (status-code assertion alone could miss a future regression where the middleware rejects with the right status while the mutation also fires).
- **#62 — real assertions in the admin permanent- delete spec.** The previous spec wrapped every assertion in `if (tableExists) { ...}`, silently no-op'ing on a clean CI DB. The rewrite seeds soft-deleted entities per-test and exercises the full type-to-confirm friction.
- **#88, #89 — admin system settings concurrency + locale resolution.** Two missing test branches around the partitioned optimistic-locking design (different settings rows don't collide) and the locale-resolution chain (Accept-Language and authenticated user preference both win over the global default).
- **#96 — concurrent first-launch race (integration test).** Two `tokio::spawn` Step 1 POSTs race the first-admin guard and assert exactly one admin row afterwards — the AC13 invariant that was unit-tested by shape but never exercised end-to-end.
- **#119 — dashboard indicator E2E seeds its own unshelved volume.** The "What needs attention" smoke used to early-return on a clean DB. The unshelved indicator is now properly seeded via `scanTitleAndVolume()`; the overdue indicator requires backdating a loan timestamp and is tracked separately as a follow-up.
- **#120 — single-active-filter contract.** Two new render unit tests lock the contract that `?filter=overdue&q=X&sort=Y` short-circuits the search-fragment branch (the indicator-list owns the response; `#browse-results` stays in the DOM as the HTMX target but carries no search content).

8.15 What's new in 1.7.6

Version 1.7.6 is a focused patch with two user-visible feature fixes and three code-review-finding hardening commits. *No schema migration, no breaking changes*; routine `docker compose pull && docker compose up -d`.

User-visible:

- **#331 — media type now editable on /title/:id.** A BD or comic catalogued via its ISBN used to lock as a "book" forever — the metadata edit form displayed `media_type` read-only. The edit form now renders a dropdown with the six variants (book, bd, cd, dvd,

magazine, report); changing it triggers an `admin_audit` entry (`title_media_type_edit`) with the before/after pair. Useful for fixing a BD that the BnF resolved as a book at scan time, or for any misclassification.

- **#335 — manual cover upload.** The provider-chain fallback (Open Library Covers API) misses most French BDs, mangas, and small-publisher technical books because that index is English-publishing-heavy. Click *Upload cover* under the cover thumbnail on `/title/:id` (Librarian+) to drop a JPG / PNG / WebP / GIF from disk (10 MiB). The image goes through the exact same decode/resize/JPEG-encode pipeline as provider-downloaded covers, so the on-disk artifact stays byte-compatible. `cover_image_url` is added to `manually_edited_fields`, which makes the bulk-refetch admin action *skip* the manually-curated title on subsequent runs — your hand-picked cover is preserved.

Hardening (no user-visible behaviour change):

- **#109 — recent-additions card test coverage.** The home-page render tests previously left `publication_date` and `cover_image_url` as `None`, never exercising the corresponding `{% if let Some(d) = ...%}` branches. A wrong format string or stray tag in those branches would have slipped past unit tests. A new factory `fake_search_result_full` and a new render test close the gap.
- **#40 — CSRF exempt-route normalisation.** The exempt-list comparison was a literal string compare; `POST /login/` or `POST //login` (variants Axum's router dispatches identically to `POST /login`) would have hit CSRF validation instead of bypassing. New `normalize_exempt_path` helper strips trailing slashes and collapses repeated slashes — percent-encoded forms are intentionally *not* decoded (a `/login%2F..%2Fadmin` bypass attempt stays rejected).
- **#36 — skip session resolver on static assets.** `session_resolve_middleware` previously walked the resolver on every request, including `/static/*`, `/covers/*`, and `/logo/*`. A single home-page load fanned out dozens of asset fetches, each triggering a session-row read (and an anonymous-row INSERT on a fresh tab). The middleware now short-circuits on those prefixes alongside the existing `/health` liveness-probe bypass.

8.16 What's new in 1.7.5

Version 1.7.5 is a focused patch bundling one user-visible bug fix and seven code-review-finding hardening commits. *No schema migration, no breaking changes*; routine `docker compose pull` && `docker compose up -d`.

User-visible:

- **#325 — add / remove contributor silently failed on /title/:id.** Clicking *+ Add contributor* on a title detail page opened the form, accepted name and role, but *Save* produced no observable effect. Root cause: the contributor form hardcodes `hx-target="#feedback-list"`, but the slot was only declared on `/catalog` and `/admin` — not on `/title/:id`. The same defect silently affected the contributor *Remove* button. Regression of the v1.7.2 fix that added the title- detail flow without porting the slot. Added a templates-audit guard that panics if any page references `hx-target="#feedback-list"` without declaring `id="feedback-list"`.

Hardening (no user-visible behaviour change):

- **#91 — form re-render on validation error.** Admin → System → Loans/Language saves return 400 with the form body re-rendered (preserves the typed value) + an inline

error feedback entry, instead of an empty bare 400 that HTMX 2.0 silently dropped.

- **#42 — strict first-child CSRF token.** The `forms_include_csrf_token` audit now requires `_csrf_token` as the very first `<input>` of every POST form, not just within the first five.
- **#48 — audit Rust HTML literals.** New sibling test `rust_emitted_post_forms_include_csrf_token` walks `src/**/*.rs` and refuses any `<form method="POST">` string literal that does not declare its CSRF input nearby. Locks the existing `src/routes/locations.rs:301` site; any new Rust- emitted form must also declare its hidden token.
- **#220 — JSON-object HX-Trigger.** `modal_confirm_retarguard` now tolerates the HTMX JSON-object HX-Trigger shape in addition to the comma-separated shape. Defensive: the current emitter uses comma-separated.
- **#31 — scanner-guard pipeline respect.** The scanner-guard JavaScript that forwards USB-scanner bursts to a focused modal text input now honors `maxLength` and skips while an IME composition is active. Sub-items 3/4/5 were already closed as N/A to the current modal shape.
- **#90 — DRY admin_system.** The duplicated optimistic-lock UPDATE SQL between `save_setting` and `run_provider_update` is consolidated; the latter now delegates and `map_errs` the conflict message to interpolate the provider label.
- **#152 — anonymous /series test.** Adds the missing role-gate test that locks FR95's anonymous-readable contract for `/series` (mirrors `/catalog` and `/locations`).

8.17 Rollback

Rollback is *not* a first-class flow. Schema migrations move forward only — there is no `down.sql`. The supported recovery strategy is:

1. Restore the pre-upgrade database backup (chapter 7).
2. Pin the image tag back to the prior version.
3. Start the stack.

This works because the backup is consistent with the prior schema. Trying to point an older image at a newly-migrated database will fail to boot.

Chapter 9

Best practices

Concentrated lessons from real-world deployments. Every section is optional — pick what fits your scale.

9.1 Storage hierarchy




Figure 9.1: The Locations page on a live install: a household hierarchy with a study room, an in-room bookcase, ten shelves, and two flat-level rooms — each node carrying its L-code and a live volume count.

- **Three levels are usually enough.** Room → shelf → row covers most households. A fourth level (e.g. “box inside row”) becomes paperwork you have to maintain.

- **Name locations the way you talk about them.** “Living-room shelf, top” beats “LR-S01-R01”. The names are visible on barcode labels; treat them as user-facing.
- **Print L-codes once.** Re-printing because you rebalanced shelves is wasted effort — L-codes are stable identifiers; the human-readable name on the label can drift and that is fine.

9.2 V-code rolls

- Print V-codes in rolls of 50 to 100. Pre-printed stickers in a holder beside the scanner station make cataloging *fast*.
- **Stick the V-code before scanning.** It removes a future failure mode where you mislabel a book and have to scan-find it again.
- **Place the sticker consistently.** Inside back cover, upper-right corner, parallel to the spine — pick a rule and stick to it. Future-you doing an audit will thank present-you.

9.3 Genre versus Dewey

mybibli supports both broad genres and Dewey codes. Pick one as your primary navigation aid:

- **Use genres** when the audience browses by mood (“I want a thriller”) and the collection is small enough that “Thriller” is granular enough.
- **Use Dewey** when the audience browses by topic (“books on 18th-century French history”) or when the collection has reached the size where “Non-fiction” is no longer a useful filter.
- Using both is fine; mybibli does not force a choice.

9.4 Audit cadence

- **Quarterly** for collections under a few hundred volumes.
- **Monthly** for collections above that, especially if more than one person catalogs.
- Always run an audit *after* a major shelf rebalance — moving books physically without telling mybibli is the single biggest source of “where is this volume?” bugs.

9.5 User accounts

- **One human, one account.** Sharing a single admin account between two people loses the audit trail and makes the last-active-admin guard less useful.
- **Most editors should be Librarians, not Admins.** The admin panel handles installation-time concerns and rare cleanups; you do not need admin rights to catalog or loan.
- **Rotate the seeded admin password on first login.** On 1.1.0+ the setup wizard forces this; on legacy 1.0.x installs do it manually before exposing the instance to a network.

9.6 Loans

- **Keep borrowers identifiable.** “Marie L.” beats “Marie” the moment the second Marie shows up. Email or phone in the notes field helps for chase-ups.
- **Run the overdue list weekly.** The dashboard indicator nags you for a reason — overdue loans compounded over a year become books you no longer own.
- **Mark damage at return time.** Switching a returned volume to *Worn* or *Damaged* captures the information while the borrower is in front of you; chasing the cost later from memory rarely works.

9.7 Metadata providers

- **Keep keys minimal.** You do not need all three keyed providers from day one — BnF + Open Library cover books for free. Add OMDb + TMDb the day you catalog your first DVD.
- **Treat keys as secrets.** `.env` is in `.gitignore`; keep backups of `.env` encrypted (chapter 7).

9.8 Upgrades

- **Read release notes before pulling :latest.** See chapter 8.
- **Back up immediately before any upgrade.** Even “additive” migrations have non-zero failure surface; a same-minute backup turns a worry into a chore.
- **Pin a version tag for production.** Auto-updates for a household instance are fine; for a community library, an unattended major bump can interrupt service.

Chapter 10

Troubleshooting

When things go wrong, start with the logs and the `/admin > Health` dashboard. This chapter catalogues the problems users actually hit, in rough order of frequency.

10.1 Reading the logs

```
docker compose logs --since 10m mybibli
docker compose logs -f mybibli
```

Logs are emitted as structured JSON. The most useful fields are `level`, `target` (which module logged this), `message`, and any contextual key=value pairs that the log line carries (`user_id`, `volume_id`, ...).

To filter:

```
docker compose logs --since 10m mybibli | grep -i error
```

10.2 App does not start

Symptom: container exits immediately

Run `docker compose logs mybibli` and look for one of:

- missing required environment variable: `DATABASE_URL` — the `.env` file is missing or the variable is unset. See section [1.5](#).
- Failed to create database pool — the database hostname / port / credentials in `DATABASE_URL` are wrong, or the DB is not reachable. From the host, try `docker compose exec mybibli sh -c 'apk add curl; curl <db-host>:<db-port>'` to confirm reachability.
- Failed to run database migrations — usually a permissions issue. Confirm the DB user has `ALL PRIVILEGES` on the database (see section [1.4](#)).
- Address already in use — the published host port is taken. Change `HOST_PORT` in `.env`.

Symptom: every page returns 404

The container is up but routes are missing. Most common cause: you pulled an image tag that does not exist (typo) and Docker fell back to a stale older tag. `docker compose pull` and restart.

10.3 Cannot log in

Symptom: “Invalid credentials” on a fresh 1.0.x install

The seed credentials are `admin/admin`. If they do not work either:

- the database migration that creates them failed silently — look for migration errors in the logs;
- the previous admin already rotated the password (check with whoever installed mybibli);
- you are looking at a 1.1.0+ install, where the seed credentials are planted by the migration but immediately soft-deleted by the issue-#173 seed gate at startup — use the wizard at `/setup` instead.

10.4 Stray users in the users table after the wizard

Symptom: four users in the database, only one was created in the wizard

If you inspect the `users` table directly (via phpMyAdmin or the MariaDB CLI) right after completing `/setup`, you will see *four* rows even though you only entered one username and password. This is by design. Three of the rows cannot be used to log in.

username	role	active	deleted_at
admin	admin	1	set, $\approx$ first-boot time
librarian	librarian	1	set, $\approx$ first-boot time
SYSTEM	system	0	NULL
<your-username>	admin	1	NULL

- `admin` and `librarian` are planted by the seed migrations (which run on every fresh install for the dev / E2E workflow) and immediately *soft-deleted* by the issue-#173 seed gate at startup. Their `deleted_at` is set, so the login query (which filters `deleted_at IS NULL`) refuses them.
- `SYSTEM` is a non-loginable utility account used to attribute the audit-log entries produced by the 30-day auto-purge background task. Its role is `system` (outside the login enum `admin/librarian`), its `password_hash` is the literal string `NO_LOGIN_SYSTEM_USER_HASH_NOT_A_VALID_ARGON2` and its `active` flag is 0.

Tip

To confirm the soft-deleted rows are inert, try logging in as `admin/admin` — it returns an authentication failure. The `/admin > Users` panel filters out both the `SYSTEM` row and any `deleted_at`-set row, so the admin UI shows the single account you created in the wizard.

The two soft-deleted rows are physically removed from the table by the 30-day auto-purge scheduler — expect them to disappear roughly one month after the boot date stamped in `deleted_at`.

Symptom: login redirects back to `/login`

The session cookie is being rejected by the browser. Causes:

- `MYBIBLI_COOKIE_SECURE=true` on a plain-HTTP deployment — the browser refuses to accept a Secure cookie over HTTP. Set the variable to `false` or front the instance with HTTPS.
- Cookies blocked by browser policy (private window with strict third-party blocking, etc.). Test in a default profile.

Symptom: locked out — last admin deactivated

The last-active-admin guard should prevent this, but if a previous admin was deactivated then permanently deleted, recovery requires direct SQL access:

```
docker compose exec db mariadb -u root -p"$MYSQL_ROOT_PASSWORD" mybibli \  
-e "UPDATE users SET deleted_at=NULL, active=1 WHERE username='your-admin';"
```

Then log in normally.

10.5 Metadata never resolves

Symptom: skeleton title stays unresolved

Open `/admin > Health`. Look at each provider’s status. A red cross usually means one of:

- the provider is genuinely down (rare for BnF / Open Library; more common for the smaller services);
- the API key is missing or invalid — open `/admin > System > Metadata Providers` and re-enter the key;
- the host cannot reach the upstream. Confirm with `docker compose exec mybibli sh -c "wget -q https://www.googleapis.com/books/v1 -O -"` or equivalent.

Symptom: only French ISBNs resolve

You disabled (or never set) the Google Books / Open Library keys. BnF is excellent for French-published titles but limited for foreign-language titles. Add at least Open Library (no key required).

10.6 Search returns too little

Symptom: known titles do not surface

- Confirm the title is not soft-deleted (check `/admin > Trash`). Search hides soft-deleted items.
- Confirm the title’s contributors are spelled the way you searched. mybibli’s search is fuzzy but not magical; “Camus” will find “Albert Camus” but “Camu” might not.

Symptom: search is slow

For collections above 10,000 titles, queries that combine contributor + free-text can take a second or two. This is a known limitation; planned indexes will close the gap.

10.7 Cover images are missing

Symptom: every title shows the placeholder

The cover directory is empty or unreadable. Confirm:

- `docker-compose.yml` declares the `mybibli_covers` named volume (or a bind mount on `COVERS_HOST_PATH`) mapped to `/app/covers`. Pre-1.1.4 installs did NOT declare this volume — see the next subsection;
- the metadata fetch actually returned a cover URL (check the title page — if the URL is missing, the provider did not supply one).

Symptom: covers vanished after upgrading the image

You upgraded from a pre-1.1.4 install and all previously cataloged covers are gone. This is issue #213 — the older compose template did not persist `/app/covers`, so the writable layer holding those JPGs was destroyed when the new container replaced the old one.

Fix forward: pull the 1.1.4 `docker-compose.yml`, run `docker compose down && docker compose up -d`. Subsequent covers are persisted in the `mybibli_covers` volume.

Recovery of lost covers: open each affected title's detail page and click *Re-fetch metadata* to repopulate from the provider chain. A bulk action is tracked at issue #214.

10.8 Where to file bugs

Issues you cannot resolve, or behaviour that contradicts this manual: open a ticket at <https://github.com/guycorbaz/mybibli/issues>. The issue template asks for the version (visible in the page footer), the install method, and the steps to reproduce — please fill them in.

Chapter 11

API & integrations

mybibli ships a small JSON HTTP API at `/api/v1/*` that lets external scripts and AI tools read your catalog and update the classification of titles. The API runs on the same web server as the human-facing UI — there is no separate daemon, no extra port to open.

This chapter walks through:

- creating an API key in the admin panel,
- the two scopes (*read only* vs. *read + write*),
- the endpoints exposed under `/api/v1/*`,
- three end-to-end examples (`curl`, Python, an AI classifier loop).

11.1 Why an HTTP API?

The everyday use case is offloading classification work to an external tool — for instance, asking an LLM to suggest a Dewey code, a genre, or a clean subtitle for every newly-added English book. The LLM walks the catalog over `GET /api/v1/titles`, fetches the detail of each via `GET /api/v1/titles/{id}`, computes a suggestion, and writes it back with `PATCH /api/v1/titles/{id}`. Every `PATCH` is audited so you can roll back changes if a model produced garbage.

The API is *not* meant for high-throughput batch ingest — for that, use the SQL database directly during a maintenance window. It *is* meant for the kind of slow, careful, opinionated write that fits inside an LLM agent loop.

11.2 Creating an API key

API keys are managed under **Administration** → **API keys**. Only Admin users see this tab.

1. Click **Administration** in the navigation bar.
2. Open the **API keys** tab.
3. Fill in a **label** — a name only you see, e.g. “*AI classifier*” or “*Backup script*”. The label is for the audit log; the key plaintext itself is random and unrelated to it.
4. Choose a **scope**:

- **Read only** — GET endpoints only. Safe for read-only dashboards, mirroring tools, or curiosity-driven scripts.
 - **Read + write** — additionally allows PATCH `/api/v1/titles/{id}` for a narrow allow-list of fields (genre, dewey, description, subtitle). Use this for LLM agents or curation tools.
5. Click **Create key**. A modal appears showing the full plaintext exactly once — copy it now. Once you close the modal there is no recovery path; if you lose the key, revoke it and create a new one.

The plaintext shape is `mybibli_<scope>_<random40>`, where `<scope>` is `ro` or `rw`. Only the `argon2` hash and the first 12 characters of the plaintext (the prefix) are persisted; the prefix is what `mybibli` uses to narrow candidates during authentication. The hash itself cannot be reversed.

11.3 Authenticating

Every `/api/v1/*` request must carry one of:

- Authorization: Bearer `<plaintext>` — preferred.
- X-API-Key: `<plaintext>` — convenience fallback.

A missing or invalid key returns 401 `Unauthorized` with a JSON envelope:

```
{ "error": "unauthorized",
  "reason": "missing_api_key" }
```

A read-only key on a write endpoint returns 403 `Forbidden` with the missing scope named:

```
{ "error": "forbidden",
  "reason": "insufficient_scope",
  "scope_required": "write" }
```

CSRF protection is intentionally bypassed for `/api/*` because bearer-token authentication does not ride on cookies and the cross-site replay risk that CSRF defends against does not apply.

11.4 Endpoints

11.4.1 GET `/api/v1/titles`

Paginated list. Supports `?page=N`, `?q=<search>`, `?genre_id=N`, and `?dewey_prefix=8`.

Returns:

```
{ "items": [ { "id": 42, "title": "...",
               "subtitle": null,
               "language": "en",
               "media_type": "book",
               "isbn": "9780...",
               "publisher": "...",
               "publication_year": 2020,
               "genre_id": 7,
               "genre_name": "Fiction",
               "dewey_code": "813.54" },
  ... ],
```



```
"page": 1, "total_pages": 12,
"total_items": 287 }
```

11.4.2 GET /api/v1/titles/{id}

Full detail, including contributors, volumes (with their location), and series assignments. The `version` field is the optimistic-locking handle — keep it next to the data your script will edit.

11.4.3 GET /api/v1/genres, /locations, /series

Reference data. Use these when an LLM needs to know which genres exist before it suggests a `genre_id`.

11.4.4 GET /api/v1/dewey/{prefix}

Bucket lookup. Pass a partial Dewey prefix (e.g. 8 for literature) and get back every title whose `dewey_code` starts with it. Useful for “find all books that already have a Dewey code in the 800s so I can compare”.

11.4.5 PATCH /api/v1/titles/{id}

Write-scope only. Updates the four whitelisted fields:

- `subtitle` — nullable.
- `description` — nullable.
- `dewey_code` — nullable; digits, dots, slashes only.
- `genre_id` — must reference an existing, non-deleted genre.

The JSON body uses partial-update semantics:

- Field **omitted** → no change.
- Field set to `null` → clear to `NULL` (for the three nullable fields above; `null` on `genre_id` returns 400).
- Field set to a value → update.

`version` is required. On mismatch the response is 409 `Conflict` carrying both the expected and the supplied version — refetch and retry.

A successful PATCH writes one row into `admin_audit` with action `api_patch_title`, attribution via the issuing admin (`user_id`) and the API key (`details.via_api_key_id + details.via_api_key_label`), and a per-field `old/new` diff. A no-op PATCH (no field actually changed) returns 200 but does NOT bump `version` and does NOT write an audit row.

11.5 Examples

11.5.1 curl

```
KEY="mybibli_ro..."
```

```
# List the first page of titles.
curl -H "Authorization: Bearer $KEY" \
```

```

https://library.example.org/api/v1/titles

# Show one title in full.
curl -H "Authorization: Bearer $KEY" \
      https://library.example.org/api/v1/titles/42

# PATCH a Dewey code (needs a write-scope key).
KEY_WRITE="mybibli_rw_..."
curl -X PATCH \
      -H "Authorization: Bearer $KEY_WRITE" \
      -H "Content-Type: application/json" \
      -d '{"version": 3, "dewey_code": "813.54"}' \
      https://library.example.org/api/v1/titles/42

```

11.5.2 Python

```

import os, requests

BASE = "https://library.example.org/api/v1"
KEY  = os.environ["MYBIBLI_API_KEY"]
H    = {"Authorization": f"Bearer {KEY}"}

# Walk every page.
page = 1
while True:
    r = requests.get(f"{BASE}/titles?page={page}",
                     headers=H, timeout=30)
    r.raise_for_status()
    data = r.json()
    for t in data["items"]:
        print(t["id"], t["title"], t.get("dewey_code"))
    if page >= data["total_pages"]:
        break
    page += 1

```

11.5.3 An LLM classifier loop

The intended write-flow looks like this — pseudocode:

```

for t in walk_titles(KEY_READ):
    if t["dewey_code"]:
        continue                # already classified
    detail = get(f"/titles/{t['id']}")
    suggestion = ask_llm(detail) # returns {"dewey_code": ...}
    if not suggestion.get("dewey_code"):
        continue
    patch(
        f"/titles/{t['id']}",
        json = {
            "version":    detail["version"],
            "dewey_code": suggestion["dewey_code"],
        },
        key  = KEY_WRITE,

```

)

If two classifiers race on the same title, the loser gets a 409 and re-fetches; no silent clobber. The audit log records who (key) suggested what — handy for rolling back a model that has a bad day.

11.6 Revoking a key

Revocation is one click in the admin panel. A revoked key *cannot* be restored; the audit history of past usage is preserved alongside the now-defunct row. If you suspect a key has leaked (committed to a public repo, e-mailed by mistake, found in a log file), revoke immediately and create a new one.

11.7 What the API does *not* expose

By design, the v1 API is read-mostly. It does not:

- create or delete titles, volumes, contributors, borrowers, loans, or series;
- update any field outside the four-field PATCH allow-list;
- return passwords, password hashes, session tokens, or any other credentials;
- accept multipart uploads (no cover-image PUT).

These are intentional v1 limits — the use case that justified the API is classification, and that fits inside the narrow allow-list. Wider write-surface would change the security posture (CSRF considerations, mass-mutation rate limits) and is deferred until a concrete need surfaces.

Chapter 12

Operations & debugging

This chapter documents where logs live on disk, how to control their verbosity, and the commands you typically need when a user reports odd behavior and you want to see what the app actually did.

12.1 Where the logs live

Starting with **v1.7.0**, mybibli writes its logs to two places simultaneously:

- **stdout** (unchanged) — visible via `docker compose logs mybibli`. Best for live-tailing during an active session.
- **A daily-rotating file on disk**, by default under `/var/log/mybibli/` inside the container, mounted on a named volume (`mybibli_logs`) on the host so files survive `docker compose pull && up -d`. Best for *post-mortem* debugging.

Files are named `mybibli.log.YYYY-MM-DD` (one per UTC day). A background task purges files older than 30 days automatically.

Changing the in-container log path

The default in-container log directory `/var/log/mybibli/` is controlled by the `MYBIBLI_LOG_DIR` environment variable. Change it only if you remount the volume elsewhere — the shipped `docker-compose.yml` maps the `mybibli_logs` named volume (or a bind mount via `LOGS_HOST_PATH`) onto `/var/log/mybibli`, so changing one without the other lands logs in the writable layer where they get wiped on every `docker compose up -d` after a pull.

See section [1.3](#) for the bind-mount setup (host-visible logs for DSM Log Center / journald shipping).

12.2 Tailing logs from the host

```
# Tail the current day's file in real time
docker compose exec mybibli tail -f /var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d)

# Last 100 lines of the current day's file
docker compose exec mybibli tail -n 100 /var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d)

# Show all available days
docker compose exec mybibli ls /var/log/mybibli/
```

```
# Search across all retained days for a keyword
docker compose exec mybibli sh -c "grep ERROR /var/log/mybibli/mybibli.log.*"
```

If you mounted the volume as a bind mount (`LOGS_HOST_PATH` in `.env`, see section 1.3), the files are visible directly on the host without needing `docker compose exec`.

12.3 Log levels

Two environment variables control verbosity:

- `MYBIBLI_LOG_LEVEL` (preferred, v1.7.0+) — accepts either a plain level (`trace`, `debug`, `info`, `warn`, `error`) or a tracing-subscriber `EnvFilter` directive list, e.g. `mybibli=debug,sqlx::query=warn`.
- `RUST_LOG` — legacy fallback, same format. Honored when `MYBIBLI_LOG_LEVEL` is unset.

Default is `info` — production-safe (`info` + `warn` + `error`, no per-query SQL noise, no per-request trace floods).

Recommended levels

info Production default. Catches authentication events, metadata fetches, scheduled-task runs, user-facing errors. Roughly 10–100 lines per day on a single-household install.

mybibli=debug Active investigation of an mybibli feature. Adds request-routing, cache hits/misses, scan dispatch decisions. Roughly 10× the `info` volume.

mybibli=trace Deepest one-shot investigation. Includes per-call tracing within mybibli’s own code. Very noisy; use briefly and bump back to `info`.

mybibli=debug,sqlx::query=info Active investigation AND see every SQL statement (useful when a query is suspect).

How to change the level

Two paths, depending on whether you want the change to survive container restarts or just for one investigation.

From the admin UI (v1.7.1+, recommended)

Log in as Admin, go to `/admin?tab=system`, scroll to the **Logging** section. Enter the new directive in the `Log level` field (anything that `tracing-subscriber::EnvFilter` accepts — plain level or full directive list). Click **Save**. The next log line uses the new filter — no container restart, no `docker compose up -d`. The value is persisted to the `settings` table, so it survives container restarts.

From `docker-compose.yml` / `.env`

Edit the file, set `MYBIBLI_LOG_LEVEL`, then:

```
docker compose up -d
```

The container restarts and the new level takes effect from the next log line. **Caveat:** the admin-saved value in the `settings` table wins on the next boot, so if you’ve ever saved a level via the

admin UI, that’s what will load — the env var only matters on a fresh install before the first save.

12.4 Reading the structured JSON

Logs are emitted as line-delimited JSON. Each line carries:

- `timestamp` — ISO 8601 UTC.
- `level` — TRACE / DEBUG / INFO / WARN / ERROR.
- `target` — the Rust module that emitted the line (e.g. `mybibli::routes::catalog`). Filter on this to narrow a search to a subsystem.
- `fields.message` — the human-readable summary.
- `additional fields.*` — contextual keys like `user_id`, `volume_id`, `isbn`, `title_id`.

To pretty-print:

```
docker compose exec mybibli tail -n 50 /var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d) \
| jq -r ' [.timestamp, .level, .fields.message] | @tsv'
```

12.5 Sharing a log excerpt for a bug report

When opening an issue at <https://github.com/guycorbaz/mybibli/issues>:

1. Note the wall-clock time when the bug occurred (e.g. “around 14:30 UTC today”).
2. Extract the relevant 5-minute window:

```
# Replace 14:25 / 14:35 with your window.
docker compose exec mybibli sh -c \
"awk '/\"timestamp\": \"$(date -u +%Y-%m-%d)T14:2[5-9]/, /\"timestamp\": \"$(date -u +%Y-%m-%d)T14:3[0-5]/' \
/var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d)"
```

3. Paste the output into the issue body, fenced as ````json`.

Redact anything you consider sensitive — log lines may contain borrower names, ISBNs of titles you own, or session-token prefixes (but never full session tokens or password hashes; those are kept out of log emission by design).

12.6 Retention & rotation

- Rotation happens **automatically** at UTC midnight, handled in-process by `tracing-appender::rolling`.
- Files older than **30 days** are purged once per day by a background task. The retention window is hardcoded in v1.7.0; a future release will expose it as an admin setting.
- The current day’s file is never deleted (the cutoff is strict).

If you want to keep logs longer than 30 days, copy the rotated files out of the volume to long-term storage:

```
docker compose cp mybibli:/var/log/mybibli/. ./logs-archive/
```

Chapter 13

Glossary

Admin A user account with role `admin`. The only role that can open the `/admin` panel and edit users, reference data, trash or system settings.

Anonymous A visitor without a session cookie. Read-only access to the public catalog. No cataloging, no loans, no edits.

Audit The routine “walk the shelves and confirm everything is where the database says it is” workflow described in section 3.6.

BDGest A French-language comic / BD catalogue used as a metadata source for that media type.

BnF *Bibliothèque nationale de France*. The French national library catalogue, used as the primary metadata source for French-language books.

Borrower A person who loans volumes from the library. Modeled as its own entity (name, optional contact information, loan history).

Cataloging The act of registering a new title and / or volume in the database. Typically barcode-driven: scan, confirm, stick label, repeat.

Cover image The visual front of a book / disc / record, downloaded from a metadata provider and cached under `COVERS_DIR` on disk.

CSP Content Security Policy. A browser-enforced instructions header that restricts which scripts and styles the page may execute. mybibli ships a strict policy that forbids inline scripts and styles.

CSRF token A per-session secret embedded into every form. mybibli rejects POST / PUT / DELETE requests whose token does not match — a defense against cross-site request forgery.

Dewey code The Dewey Decimal Classification number, used for topic-based catalog navigation. Optional per title.

EAN-13 The 13-digit barcode standard. ISBNs are encoded as EAN-13 since 2007; CDs, DVDs and other media carry EAN-13 codes too.

Exemplaire Synonym for *volume* (the physical copy). Used by some French libraries; mybibli prefers *volume*.

- Genre** A free-form classification of a title's subject (Fiction, Non-fiction, Mystery, ...). Defined in the genres reference table.
- Hard-delete** Deleting a row outright. mybibli only hard-deletes from the `/admin > Trash` panel and from the auto-purge background task; everything else is a soft-delete.
- HTMX** The JavaScript library mybibli uses to swap fragments of the page in place without full reloads. Operators do not need to understand HTMX to use mybibli.
- ISBN** International Standard Book Number. Identifier printed on the back cover of books, encoded as an EAN-13 barcode.
- Last-active-admin guard** The rule that prevents deactivating or permanently deleting the last user with the `admin` role. See section 2.6.
- L-code** A storage-location barcode label. Format: L followed by a zero-padded number.
- Librarian** A user account with role `librarian`. Can catalog, loan, return and edit most data but cannot administer users or settings.
- Loan** A record that a specific volume is currently out with a specific borrower. Tracks the loan date, due date, return date, and the volume's pre-loan location.
- Location** The physical place where a volume lives. Modeled as a tree (room → shelf → row), identified by an L-code.
- MariaDB** The database engine mybibli stores its data in. A fork of MySQL.
- Media type** The kind of object a title represents: Book, Comic / BD, Audio, Film / series. Drives provider selection and field display.
- Multi-position volume** A single physical volume that contains several positions in a series — typically an omnibus comic album.
- MusicBrainz** A volunteer-maintained music metadata catalogue, used as the primary source for audio media.
- Node type** A label applied to storage locations to indicate what kind of container they are (room, shelf, row, box). Pure label — no semantic enforcement.
- OMDb** Open Movie Database. Metadata source for films and TV series.
- Open Library** An open volunteer-driven catalogue of books, used as a metadata source.
- Optimistic locking** The technique mybibli uses to detect concurrent edits. Two users editing the same title get a clear error on the second save instead of silently overwriting each other.
- Reactivation** Restoring a soft-deleted row. For users, clears `deleted_at` and lets them log in again. For reference data, the term also covers re-creating a row with the same name as a soft-deleted one, which transparently undoes the delete.
- Reference data** The four taxonomy tables: genres, contributor roles, volume states, location node types. Editable from `/admin > Reference data`.
- Series** A named ordering of related titles (e.g. a trilogy or an open-ended comic series). Each title's position in the series is a small integer or a range (for omnibuses).

Setup wizard The four-step first-launch flow at `/setup` that an empty instance presents until an admin account exists.

Soft-delete Setting a row's `deleted_at` column to “now” instead of physically removing the row. Soft-deleted rows are invisible to normal queries but surface in `/admin > Trash` for 30 days, after which the auto-purge hard-deletes them.

Title The bibliographic entity (a book, a film, an album). A title can have multiple volumes — for example, two copies of the same novel on different shelves.

TMDb The Movie Database. A community-curated metadata source for films and TV series.

Trash The `/admin > Trash` tab. Lists soft-deleted entities of every type and lets an admin restore them or permanently delete them.

V-code A per-volume barcode label. Format: `V` followed by a zero-padded number.

Volume The physical object on a shelf — one copy of a title. A title can have many volumes; each carries its own V-code, location and state.

Volume state A reference-data row describing the condition of a volume: New, Good, Worn, Damaged,

Index

1.1.0 upgrade, 32

admin panel, 10
admin role, 26, 59
anonymous, 26, 59
audit, 18, 59

backup
 .env, 28
 covers, 28
 database, 28
BDGest, 22, 59
BnF, 22, 59
borrower, 17, 59

Code-128, 12
comic / BD, 20
contributor roles, 13
covers, 59
 directory, 3
 persistence, 6
CSRF, 59

Dewey, 44, 59

EAN-13, 12
env vars
 .env.example, 7
 provider keys, 24

film, 20

genres, 13, 60
Google Books, 22

HID keyboard mode, 12

i18n, 13
inactivity timeout, 27
installation, 3

L-code, 11, 60
last-active-admin guard, 60

librarian, 26, 60
loan, 60
location, 15, 60
location tree, 10
logs
 directory, 4

media types, 20, 60
metadata providers, 3, 22
multi-position, 20, 60
MusicBrainz, 20, 22, 60
MYIBILI_SEED_DEV_USERS, 33

node type, 11, 13, 60

OMDb, 21, 22, 60
omnibus, 20
Open Library, 22, 60

roles, 14, 26

scan field, 15
security
 1.1.0 floor, 3
series, 20, 60
session cookie, 26
setup wizard, 6, 9, 61
shelving, 18
soft-delete, 61
storage locations, 10
SYSTEM user, 47

TMDb, 21, 22, 61
troubleshooting
 DATABASE_URL, 46

V-code, 11, 61
volume, 10, 11, 61
volume state, 13, 15, 61

wish list, 18
/wishlist, 18