

mybibli

User Manual

The mybibli authors
<https://github.com/guycorbaz/mybibli>

Version 1.7.1 — May 21, 2026

Licensed under AGPL-3.0-or-later.

This manual covers mybibli version 1.3.1 and later (the v1.1.0 floor is still enforced for fresh installs — see the Installation chapter for the seed-gate background).

Contents

Preface	1
1 Installation	3
1.1 System requirements	3
1.2 Pre-installation checklist	3
1.3 Method 1 — Bundled database	4
1.4 Method 2 — External database (Synology NAS example)	6
1.5 Environment variable reference	7
1.6 First boot — the setup wizard	9
1.7 Verifying the install	9
2 Configuration	10
2.1 Storage locations and hierarchy	10
2.2 Labels: V-codes and L-codes	11
2.3 Hardware: barcode scanner	11
2.4 Reference data	12
2.5 System settings	13
2.6 Users	14
3 Everyday use	15
3.1 The home page	15
3.2 Cataloging a new title via barcode	16
3.3 Searching the collection	16
3.4 Moving a volume to another location	16
3.5 Loans	17
3.6 Library audit	17
3.7 Shelving workflow	18
3.8 Wish list	18
3.8.1 Adding an entry	18
3.8.2 Browsing at the bookstore	18
3.8.3 Marking a book as bought	18
3.8.4 Auto-detect on catalog scan	19
4 Multiple media types	20
4.1 Books	20
4.2 Comics and BD	20
4.3 Audio releases	20
4.4 Films and series	21
4.5 Picking the media type at scan time	21
5 Metadata providers	22

5.1	The provider chain	22
5.2	API keys	23
5.3	Where the keys live	24
5.4	Provider health	24
5.5	Rate limits and back-pressure	24
5.6	Troubleshooting empty metadata	24
6	Roles and permissions	26
6.1	Who can do what	26
6.2	Sessions and the inactivity timeout	26
6.3	The last-active-admin guard	27
6.4	Password reset	27
6.5	Anonymous sessions	27
7	Backup and restore	28
7.1	What to back up	28
7.2	Backup procedure	28
7.3	Recommended cadence	29
7.4	Restore procedure	29
7.5	Migration to a new host	30
8	Upgrade and migration	31
8.1	Routine upgrades (1.y to 1.z)	31
8.2	The 1.0.x to 1.1.0 jump	31
8.3	Reading the release notes	32
8.4	What's new in 1.7.1	32
8.5	What's new in 1.7.0	33
8.6	What's new in 1.6.2	34
8.7	What's new in 1.6.1	35
8.8	What's new in 1.6.0	35
8.9	What's new in 1.5.2	36
8.10	What's new in 1.5.1	36
8.11	What's new in 1.5.0	37
8.12	What's new in 1.4.0	38
8.13	What's new in 1.3.1	38
8.14	What's new in 1.3.0	39
8.15	What's new in 1.2.2	39
8.16	What's new in 1.2.1	40
8.17	What's new in 1.2.0	41
8.18	What's new in 1.1.9	41
8.19	What's new in 1.1.8	42
8.20	What's new in 1.1.7	42
8.21	What's new in 1.1.6	43
8.22	What's new in 1.1.5	43
8.23	What's new in 1.1.4	44
8.24	What's new in 1.1.3	44
8.25	What's new in 1.1.2	44
8.26	Rollback	45
9	Best practices	46
9.1	Storage hierarchy	46
9.2	V-code rolls	47

9.3	Genre versus Dewey	47
9.4	Audit cadence	47
9.5	User accounts	47
9.6	Loans	47
9.7	Metadata providers	48
9.8	Upgrades	48
10	Troubleshooting	49
10.1	Reading the logs	49
10.2	App does not start	49
10.3	Cannot log in	50
10.4	Stray users in the <code>users</code> table after the wizard	50
10.5	Metadata never resolves	51
10.6	Search returns too little	51
10.7	Cover images are missing	52
10.8	Where to file bugs	52
11	API & integrations	53
11.1	Why an HTTP API?	53
11.2	Creating an API key	53
11.3	Authenticating	54
11.4	Endpoints	54
11.4.1	GET <code>/api/v1/titles</code>	54
11.4.2	GET <code>/api/v1/titles/{id}</code>	55
11.4.3	GET <code>/api/v1/genres, /locations, /series</code>	55
11.4.4	GET <code>/api/v1/dewey/{prefix}</code>	55
11.4.5	PATCH <code>/api/v1/titles/{id}</code>	55
11.5	Examples	55
11.5.1	<code>curl</code>	55
11.5.2	Python	56
11.5.3	An LLM classifier loop	56
11.6	Revoking a key	57
11.7	What the API does <i>not</i> expose	57
12	Operations & debugging	58
12.1	Where the logs live	58
12.2	Tailing logs from the host	58
12.3	Log levels	59
12.4	Reading the structured JSON	60
12.5	Sharing a log excerpt for a bug report	60
12.6	Retention & rotation	60
13	Glossary	62
	Index	65

Preface

Who this manual is for

This manual is written for the person who installs, configures and runs a mybibli instance for their household, a small association, or a community library. It assumes you can:

- open a terminal and run `docker compose` commands;
- edit a plain-text configuration file with the editor of your choice;
- read a URL printed in a log line and open it in a browser.

No Rust, MariaDB or HTMX knowledge is required. The manual is deliberately scoped to the *operator* role — building mybibli from source or modifying its code is covered by `CLAUDE.md` and the planning artifacts in the repository.

Conventions

Commands. Shell snippets are shown in a fixed-width box:

```
docker compose up -d
```

Configuration files. File excerpts use the same style with the file name called out in the surrounding text. Comments inside the snippet use the file's native comment marker.

Callouts. Three sidebars highlight important text:

Note

A *note* provides background or context that helps you understand why something works the way it does. Skipping it is safe.

Tip

A *tip* suggests a practice that makes day-to-day use smoother. Try it once you are comfortable with the basics.

Warning

A *warning* flags a step where a mistake can lose data, lock you out of the admin panel, or expose the instance to unauthorized access. Read these in full.

Cross-references. Chapters and sections are linked in the PDF — click any blue title or page number to jump there. The [index](#) at the end maps every important term to its page of first use.

Getting help

The canonical source of truth for mybibli is the GitHub repository:

<https://github.com/guycorbaz/mybibli>

Open an issue against the `guycorbaz/mybibli` repository when you hit a bug, a documentation gap, or a missing feature. The issue templates ask for the version (visible at the bottom of any page in the admin panel), your install method, and the steps to reproduce.

Manual coverage

This manual covers mybibli **1.1.0 and later**. Older versions shipped a security-relevant bug where the first-launch setup wizard was silently bypassed — see chapter 8 for the upgrade path from 1.0.x.

Happy cataloging.

Chapter 1

Installation

This chapter walks you through installing mybibli on your own hardware. The reference deployment target is a home server, NAS, or any always-on Linux box that can run Docker.

1.1 System requirements

- **Operating system:** any Linux distribution, macOS, or Windows host capable of running Docker. The host kernel must support cgroups v2 (any distribution from 2021 onward).
- **Docker:** version 24 or later, with the `compose` plugin. Test with `docker compose version`.
- **Disk:** at least 1 GB free for the application image plus the database. Cover images add roughly 50 KB per title.
- **Memory:** 512 MB is enough for a household-sized library (a few thousand titles); 1 GB gives generous headroom for the metadata-fetch pipeline.
- **Network:** outbound HTTPS to the metadata providers you choose to use (see chapter 5). LAN-only deployments work without any inbound port forwarding.

Warning

Do not install mybibli versions below 1.1.0. Every published image with a tag below 1.1.0 contains a seed migration that creates the credentials `admin/admin` and `librarian/librarian` on *every* fresh install, including production. The first-launch setup wizard is silently bypassed. Use 1.1.0 or later; the pre-1.1.0 images will be removed from Docker Hub to prevent accidental use.

1.2 Pre-installation checklist

Before you start:

1. Pick a host directory for the mybibli files (e.g. `/srv/mybibli` on Linux, `/volume1/docker/mybibli` on Synology DSM).
2. Decide how cover images are stored. The shipped `docker-compose.yml` declares a Docker *named volume* (`mybibli_covers`) — zero configuration, Docker manages persistence and backup via `docker volume`. If you prefer a host-visible directory (Synology DSM file

browser, manual rsync backups), pick a host path now (e.g. `/volume1/docker/mybibli/covers`) and you will switch the compose file to a bind mount in step 3 below.

3. Decide how logs are stored (v1.7.0+). The shipped `docker-compose.yml` declares a third named volume (`mybibli_logs`) for the daily-rotating log files. Like for covers, you can keep the Docker-managed volume (zero-config) or switch to a host-visible bind mount for shipping logs to DSM Log Center, journald, or an external log collector. See section 1.3.
4. Choose a strong password for the bundled MariaDB root and application users — or, for an external database, create a dedicated database and user up front (see section 1.4).

1.3 Method 1 — Bundled database

This is the simplest path. The shipped `docker-compose.yml` file declares two services: the mybibli application and a MariaDB instance pinned to a known version. Both are managed as one unit.

1. Get the compose file and `.env`

Download the `docker-compose.yml` and `.env.example` from the release you intend to install (or check out the repository at the matching tag):

```
mkdir -p /srv/mybibli && cd /srv/mybibli
curl -O https://raw.githubusercontent.com/guycorbaz/mybibli/v1.7.0/docker-compose.yml
curl -O https://raw.githubusercontent.com/guycorbaz/mybibli/v1.7.0/.env.example
cp .env.example .env
```

1b. (Optional) Switch covers to a bind mount

By default, the shipped `docker-compose.yml` stores downloaded cover images in the Docker-managed named volume `mybibli_covers`. This works zero-config and persists across container upgrades (issue #213). Most operators should leave it as-is.

If you want covers visible in your host file manager (Synology DSM File Station, manual rsync backups, etc.), edit the compose file:

- In the `mybibli` service's volumes: block, comment out the `- mybibli_covers:/app/covers` line and uncomment the `- ${COVERS_HOST_PATH:-./covers-host}:/app/covers` line right below it.
- In `.env`, set `COVERS_HOST_PATH` to the host path you picked in the pre-installation checklist (e.g. `COVERS_HOST_PATH=/volume1/docker/mybibli/covers`).
- Create the directory before starting the stack (`mkdir -p`) so Docker doesn't auto-create it owned by root.

1c. (Optional) Switch logs to a bind mount

New in **v1.7.0**. By default, mybibli writes its daily-rotating log files to the Docker-managed named volume `mybibli_logs` mounted at `/var/log/mybibli` inside the container. Files are named `mybibli.log.YYYY-MM-DD` and the in-process appender purges files older than 30 days automatically.

If you want logs visible directly on the host (Synology DSM Log Center, journald shipping, external log aggregator, manual `tail`), edit the compose file:

- In the `mybibli` service's volumes: block, comment out the `- mybibli_logs:/var/log/mybibli` line and uncomment the `- ${LOGS_HOST_PATH:-./logs-host}:/var/log/mybibli` line right below it.
- In `.env`, set `LOGS_HOST_PATH` to the host path you want (e.g. `LOGS_HOST_PATH=/volume1/docker/mybibli/log`).
- Create the directory before starting the stack (`mkdir -p`) so Docker doesn't auto-create it owned by root.

The logs volume is forensic-only — you can wipe it at any time without affecting catalog data. See chapter 12 for log-tailing and log-level controls.

2. Edit `.env`

Open `.env` in your editor. Every variable carries an explanatory comment; the values that matter for a bundled-database install are:

```
DATABASE_URL=mysql://mybibli:mybibli_password@db:3306/mybibli?charset=utf8mb4
MYSQL_HOST=db
MYSQL_DATABASE=mybibli
MYSQL_USER=mybibli
MYSQL_PASSWORD=<choose-a-strong-password>
MYSQL_ROOT_PASSWORD=<choose-a-strong-password>

HOST_PORT=8080          # change if 8080 is already taken on the host
APP_LANGUAGE=en         # or 'fr' for the French UI by default

MYBIBLI_COOKIE_SECURE=false # set to true only behind HTTPS
```

The `DATABASE_URL` *must* reference the same user, password and database name as the `MYSQL_*` values — the compose file rebuilds the URL from those parts but the Rust binary reads `DATABASE_URL` directly, so the two must agree.

Warning

The default passwords in `.env.example` (`mybibli_password`, `root_password`) are placeholders, not credentials to keep. Change them before the first `docker compose up`. After the database has been created, rotating these passwords requires also rewriting the MariaDB user entries — easier to get it right the first time.

3. Start the stack

```
docker compose up -d
docker compose logs -f mybibli
```

The first launch:

1. pulls the `gcorbaz/mybibli:1.1.0` (or later) image and the MariaDB image;
2. starts the database and waits for it to be healthy;
3. runs the `mybibli` database migrations (one-time schema setup);
4. starts the HTTP listener on the port you configured.

You can stop tailing the logs with `Ctrl-C`; the containers keep running detached.

4. Open the setup wizard

Visit <http://<host>:8080/setup> in a browser. The first-launch setup wizard appears (covered in section 1.6).

1.4 Method 2 — External database (Synology NAS example)

If you already run MariaDB on your NAS — for example via the Synology **MariaDB 10** package — you can point mybibli at it instead of running a second database in Docker. The mybibli image is unchanged; only the compose file and `.env` differ.

1. Create the database and user

Connect to the external MariaDB as `root` and run:

```
CREATE DATABASE mybibli CHARACTER SET utf8mb4 COLLATE utf8mb4_unicode_ci;
CREATE USER 'mybibli'@'%' IDENTIFIED BY 'choose-a-strong-password';
GRANT ALL PRIVILEGES ON mybibli.* TO 'mybibli'@'%';
FLUSH PRIVILEGES;
```

Note

On Synology DSM, MariaDB 10 listens on port 3307 by default (not 3306 — that port is reserved for the older MySQL package). Open **MariaDB 10** → **Settings** and tick *Enable TCP/IP connection* so containers on the same Docker bridge can reach the server.

2. Trim the compose file

Edit `docker-compose.yml` and remove the bundled database:

- delete (or comment out) the entire `db:` service;
- delete (or comment out) the `depends_on:` block on the `mybibli` service;
- in the bottom `volumes:` block, remove `mybibli_db_data:` only — **keep** `mybibli_covers:`. Cover images are local to the app and must survive container upgrades regardless of where the database lives (issue #213).

3. Point `.env` at the external database

```
DATABASE_URL=mysql://mybibli:strong-password@host.docker.internal:3307/mybibli?charset=utf8mb4
```

Substitute the host name and port that match your network. From a mybibli container running on the same Docker bridge as the NAS host, `host.docker.internal` resolves to the host on macOS and recent Linux Docker installs; on older Docker engines, use the LAN IP of the NAS instead (e.g. 192.168.1.10).

The `MYSQL_*` variables are now unused — leave them blank or delete them.

4. Start the application

```
docker compose up -d
docker compose logs -f mybibli
```

The first launch runs the schema migrations against the external database. The Synology MariaDB package will hold this schema for as long as the package exists; backing it up is covered by the NAS-level backup tools (see chapter 7).

Tip

The external-database setup is the recommended path when you already maintain a database server, because it gives you one backup pipeline to monitor instead of two. The bundled-database setup is faster to get going if you have no other workloads on the host.

1.5 Environment variable reference

The full reference lives in the file `.env.example` at the repository root — every variable carries an inline comment. The table below summarises which variable applies where.

All boolean variables follow a **strict accept-set**: only the strings `1`, `true` and `TRUE` count as “on”. Anything else (including the empty string, `0`, and `false`) is treated as “off”. This avoids the classic footgun where a stale shell value like `0` reads as “set” and silently flips an opt-out.

Variable	Default	Purpose
<i>Database</i>		
DATABASE_URL	— (required)	MariaDB connection string (DSN).
MYSQL_HOST	db	Hostname for the bundled DB service.
MYSQL_PORT	3306	Port for the bundled DB service.
MYSQL_DATABASE	mybibli	DB name for the bundled service.
MYSQL_USER	mybibli	DB user for the bundled service.
MYSQL_PASSWORD	—	DB user password for the bundled service.
MYSQL_ROOT_PASSWORD	—	Root password for the bundled service.
<i>HTTP server</i>		
HOST	0.0.0.0	Bind address inside the container.
PORT	8080	Bind port inside the container.
HOST_PORT	8080	Port published by Docker on the host.
<i>Application</i>		
MYBIBLI_LOG_LEVEL	info	v1.7.0+. Runtime log level. Plain level (trace/debug/info/warn/error) or an EnvFilter directive list. Takes precedence over RUST_LOG . Also flippable from /admin > System .
MYBIBLI_LOG_DIR	/var/log/mybibli	v1.7.0+. Directory for daily-rotating log files inside the container. Mapped to the mybibli_logs named volume.
LOGS_HOST_PATH	./logs-host	v1.7.0+. Optional bind-mount host path; used only when you uncomment the bind-mount line in docker-compose.yml .
RUST_LOG	mybibli=info	Legacy tracing filter. Honored when MYBIBLI_LOG_LEVEL is unset.
APP_LANGUAGE	en	Default UI language for anonymous visitors. v1.7.0 adds de + it to en / fr .
COVERS_DIR	./covers	Filesystem path for cover-image storage.
COVERS_HOST_PATH	./covers-host	Optional bind-mount host path; used only when you uncomment the bind-mount line in docker-compose.yml .
<i>Hardening</i>		
MYBIBLI_COOKIE_SECURE	false	Set to true only behind HTTPS.
CSP_REPORT_ONLY	false	Switch CSP header to report-only mode.
<i>Metadata providers (migrated to DB on first boot)</i>		
GOOGLE_BOOKS_API_KEY	—	Optional key for Google Books.
OMDB_API_KEY	—	Optional key for OMDb (films/series).
TMDB_API_KEY	—	Optional key for TMDb (films/series).
<i>Optional dev / test overrides</i>		
MYBIBLI_SKIP_SETUP	false	Bypass the first-launch wizard. <i>Never on prod.</i>
MYBIBLI_SKIP_STARTUP_PURGE	false	Skip the boot-time auto-purge.
MYBIBLI_SEED_DEV_USERS	false	Seed dev admin/admin + librarian/librarian . <i>Never on prod.</i>

Provider base-URL overrides (**BNF_API_BASE_URL**, **GOOGLE_BOOKS_API_BASE_URL**, ...) exist solely to point the metadata pipeline at the in-tree mock server during automated tests. Leave them unset in production.

1.6 First boot — the setup wizard

On a clean install, every request is redirected to `/setup` until the wizard completes. The wizard has four steps:

1. **Admin account.** Choose a username and a strong password for the first administrator. This step is the whole reason the wizard exists — once the account is created, the wizard's gate predicate flips and subsequent visitors land on the public catalog instead.
2. **Metadata providers.** Paste in any API keys you have on hand (Google Books, OMDb, TMDb). All keyed providers are optional — you can skip this step entirely and add keys later from `/admin > System > Metadata Providers`.
3. **Preferences.** Pick the default UI language for anonymous visitors and the loan overdue threshold (in days).
4. **Done.** A summary screen. After confirming, the wizard writes `settings.setup_completed_at` and disables itself for the lifetime of the installation.

The wizard is idempotent: if you close the browser between steps, reopen `/setup` and the wizard resumes server-side at the right step (it queries the database on every load — there is no client state to lose).

1.7 Verifying the install

After the wizard finishes:

- Log in at <http://<host>:8080/login> with the admin credentials you just created.
- Visit `/admin > Health`. The page should show non-zero entity counts (zero on a clean install — that is expected) and a green check next to each metadata provider that has been reachable in the last 5 minutes.
- Scan an ISBN (the search field at the top of the home page) to confirm the metadata pipeline resolves a real title. The first scan takes a couple of seconds; the result appears in the timeline of recent additions on the home page once the background fetch completes.

If anything goes wrong, jump to chapter [10](#).

Chapter 2

Configuration

Once the setup wizard is behind you, the day-one configuration work consists of three blocks: defining your *storage locations*, populating the *reference data* (genres, contributor roles, volume states, location types), and adjusting the system-wide *settings* (overdue threshold, auto-purge cadence, provider keys).

Everything in this chapter is configured from the `/admin` panel — there is no separate configuration file. The panel has five tabs:

- **Health** — read-only status dashboard.
- **Users** — create / deactivate users, change roles.
- **Reference data** — the four taxonomy tables.
- **Trash** — view and restore soft-deleted items.
- **System** — global settings.

Only the `Admin` role sees the panel; `Librarian` and `Anonymous` users get a 403 from every `/admin/*` route.

2.1 Storage locations and hierarchy

Storage locations are the physical spots where your volumes live: a room, a shelf, a row, a numbered box, a specific cell in a cabinet. mybibli organises them as a **tree** — every location has at most one parent. The hierarchy is arbitrary: you decide both the levels and the labels.

A workable household hierarchy might be three levels deep:

Living room	(room)
Bookshelf A	(shelf)
Row 1	(row)
Row 2	(row)
Bookshelf B	(shelf)
Row 1	(row)
Bedroom	(room)
Nightstand	(shelf)

Three levels is enough to find any volume in seconds without forcing the hierarchy to mirror every nook of your living space. A flat list (one big “shelf” per physical shelf) works equally well

for collections under a few hundred volumes.

Editing the tree

Open `/locations`. The tree is rendered as a nested list with inline `create` / `rename` / `move` / `delete` affordances. Drag-and-drop is *not* supported (mybibli ships with a CSP-strict no-JS policy); moves are done by editing a node and picking a new parent from a dropdown.

Location node types

Each location node carries a **node type** — “room”, “shelf”, “row”, “box”, anything you like. Node types are managed under `/admin > Reference data > Location node types`. They are pure labels; mybibli does not enforce that a “shelf” must sit under a “room”.

Tip

Renaming a node type cascades to every location that uses it (the type is stored by name in the locations table). Deleting a node type is refused as long as one location still references it.

2.2 Labels: V-codes and L-codes

mybibli prints two kinds of barcode labels:

- **V-codes** — one per *volume* (the physical book on a shelf). Format: `v` followed by a zero-padded sequence (e.g. `V0042`).
- **L-codes** — one per *storage location*. Format: `L` followed by a zero-padded sequence (e.g. `L0007`).

V-codes and L-codes are scanned with the same barcode scanner you use for ISBNs at cataloging time. The scanner’s prefix (`v` or `L`) tells mybibli which lookup to perform.

Printing labels

Open `/locations` or the relevant volume page and click **Print label**. mybibli renders a printable HTML page with the Code-128 barcode and human-readable text. Print on plain paper or sticker stock — the encoding is forgiving enough that a low-cost inkjet works for V-codes and L-codes alike.

Tip

Print V-code stickers *before* you start cataloging in bulk. A roll of pre-printed V0001 through V0500 means you can grab a sticker, stick it on the back cover, and scan in one motion — no context-switch between the cataloging UI and the printer.

2.3 Hardware: barcode scanner

mybibli is designed around a barcode scanner but does not strictly require one — every ISBN, V-code and L-code field accepts manual keyboard input as a fallback. That said, for any collection beyond a handful of titles, a cheap USB scanner pays for itself within the first dozen scans.

What to buy

Any *HID-keyboard-mode* USB scanner in the 15–30 € range works. The HID mode is what makes the scanner driver-free: the OS sees it as a regular keyboard and the scanned digits land directly in whatever input field has focus. Almost every general-purpose USB scanner on the market ships in this mode by default; look for the words “HID”, “keyboard emulation” or “plug-and-play” on the product page.

The scanner must support the two symbologies mybibli uses:

- **EAN-13** — the 13-digit code on the back cover of books, CDs, DVDs, and most consumer products. Universally supported.
- **Code-128** — the symbology mybibli uses to encode V-codes and L-codes on the labels it prints (see section 2.2). Universally supported.

You do *not* need a 2D imager (the kind that reads QR codes) — a basic 1D laser or LED scanner is enough. mybibli does not use QR codes.

Configuration

Three settings on the scanner matter. Most scanners ship them correctly out of the box; if your scans behave strangely, check the scanner’s manual or the configuration sheet of barcodes that comes in the box:

- **Terminator / suffix:** Enter (also called CR or Carriage Return). mybibli’s cataloging form auto-submits on Enter, so this turns each scan into one round-trip. Anything else (Tab, no terminator) will not work without manual key presses.
- **Prefix:** none. Some scanners can prepend a fixed string before the payload — leave it disabled.
- **Keyboard layout:** *must* match the layout of the host the browser runs on. If your laptop uses an AZERTY layout but the scanner is configured for QWERTY, the digits come out as letters (e.g. 1234567890 arrives as &é" (-è_çà) and no ISBN ever resolves. The scanner’s manual will have a code chart to switch.

Tip

USB scanners on a Synology NAS host: it does not matter that the scanner is plugged into the laptop you use to access the mybibli web UI rather than the NAS itself. The scanner is a keyboard for the browser, not for the server. The same setup works with a phone-side scanner app, an industrial Bluetooth scanner paired to the laptop, or any USB stick scanner.

What to do when the scanner is unavailable

Every scan field also accepts hand-typed input. Click into the field labelled “Scan ISBN” (or “Scan V-code” / “Scan L-code” elsewhere), type the digits, press Enter. The form behaves identically.

2.4 Reference data

Four taxonomy tables drive the dropdowns you see while cataloging. They live under `/admin > Reference data:`

Genres The genres assigned to titles. mybibli ships with a baseline FR/EN seed (Fiction, Non-fiction, Biography, ...) — edit or extend to match your collection.

Contributor roles Author, Illustrator, Translator, Editor, Narrator, ...Used by title detail pages to attribute the right people to the right slots.

Volume states The physical condition of an exemplaire — New, Good, Worn, Damaged, ...Useful for loan returns (mark a volume Worn when the borrower damaged it).

Location node types See section 2.1 above.

CRUD semantics

All four tables share the same CRUD pattern:

- **Create** adds a row. If the name collides with a soft-deleted row of the same table, the existing row is reactivated and the result reads *Reactivated* instead of *Created*.
- **Rename** updates the name. For location node types, the new name cascades to every location that uses the type.
- **Delete** performs a soft-delete *only when the row is unused*. If any title still references the genre / role / state / type, the deletion is refused with a count of dependents.

Reference data is not translated

Reference data values are stored verbatim — the FR seed creates a genre called “Roman” and that is the literal text English-speaking users will see if they switch the UI to English. Only the UI chrome (button labels, navigation, error messages) is translated.

2.5 System settings

The `/admin > System` tab groups every global setting. The panel is divided into three forms, each with its own save button:

Loans

- **Overdue threshold (days)**. Defines how many days a loan can run before it surfaces as overdue on the dashboard. Default 30.
- **Session inactivity timeout (seconds)**. Idle time after which an authenticated session is wiped. Default 1800 s (30 minutes).
- **Auto-purge interval (seconds)**. How often the background scheduler hard-deletes items older than 30 days in the trash. Default 86400 s (24 hours).

Metadata providers

One row per keyed provider (Google Books, OMDb, TMDb). Each row has a key input and a *Clear* button. Saving a key takes effect on the next metadata fetch — no restart, no re-deploy.

Note

If the same provider key is set in `.env` and cleared from the UI, the next reboot re-migrates the value from `.env`. The operator’s deployment-time intent wins on boot. To make a

Clear durable across reboots, also remove the variable from `.env` (or `docker-compose.yml`).

Language

Sets the default UI language for *anonymous* visitors. Authenticated users override this with the language toggle in the navigation bar — the choice is stored on their user row.

2.6 Users

Open `/admin > Users` to manage user accounts. Each row shows the username, role, active flag, and last-seen timestamp. The available actions are:

- **Create user** (top of the page) — username, password, role.
- **Edit** — change role or rename. Cannot change another user's password from this view; password resets go through the user's own **Account** page.
- **Deactivate** — soft-delete the user. Their sessions are wiped immediately; they cannot log back in.
- **Reactivate** — undo a deactivation (from the **Trash** tab).

Two safety guards apply:

- **Self-deactivate is blocked.** You cannot deactivate the account you are logged in as.
- **Last-active-admin guard.** The system always keeps at least one active admin. Deactivating the last one is refused with a 409 conflict and a clear error.

The three roles (Anonymous, Librarian, Admin) are detailed in chapter 6.

Chapter 3

Everyday use

This chapter walks through the workflows you will run dozens of times once the library is live: cataloging a new volume, searching your collection, moving a volume, lending it out, taking it back, and auditing what is on the shelves.

3.1 The home page




Figure 3.1: The home page on a live install: search bar, genre filter chips, “À traiter” counters, “Aperçu de la collection” totals, and a recent-additions strip with cover thumbnails fetched through the metadata-provider chain.

The home page is the dashboard. From top to bottom:

- **Scan field.** A focused input that accepts an ISBN, a V-code or an L-code. The scanner-guard layer routes the input to the right page depending on its prefix.
- **Indicators row.** Recent additions, overdue loans, series with gaps, unshelved volumes — each click jumps to the filtered list.
- **Recent activity.** The last seven days of cataloging and returns.

3.2 Cataloging a new title via barcode

1. Focus the home-page scan field (it is selected by default on page load). Scan the ISBN / EAN-13 of the back cover.
2. mybibli redirects to the catalog page, which displays a skeleton title with the barcode you scanned. A background task fires off requests against the metadata provider chain.
3. Within a few seconds, the skeleton fills in: title, contributors, publisher, year, cover image. A small indicator confirms which provider resolved the title.
4. Pick a **location** from the dropdown (or scan the L-code of the destination shelf) and pick a **volume state** (typically *New*). Click **Add volume**.
5. Print and stick a V-code label on the volume.

If the metadata pipeline fails to resolve the ISBN, the title appears with an *Unresolved* badge. You can edit the title fields manually and the volume is still created.

Tip

Cataloging in batches works best when the V-code sticker roll, the scanner and the printer are all within arm's reach. Most of the time is spent scanning, not waiting for mybibli — the metadata fetch is asynchronous, so you do not have to wait for it to finish before moving on to the next volume.

3.3 Searching the collection

The `/search` page accepts:

- full or partial **title**;
- **contributor name** (matches on author by default; other contributor roles are scored lower);
- **ISBN / EAN-13** — exact match;
- **V-code** or **L-code** — exact match;
- **Dewey code** (for libraries that use Dewey).

Results are sorted by relevance. Click a result to open the title page; click a contributor to filter the catalog by their work; click a location label to see every volume currently in that location.

Tip

You can “search” a single volume by scanning its V-code from any page that has the home-page scan field — that includes the home page, `/search` and `/loans`. Faster than typing.

3.4 Moving a volume to another location

1. Open the volume page (search by V-code or click through from a title page).
2. Click **Move**.

3. Scan the L-code of the destination location, or pick it from the dropdown.
4. Confirm.

mybibli records the move in the volume's location history so you can trace where a book has lived over time.

3.5 Loans

Lending a volume

1. Open `/loans` and click **New loan**.
2. Scan or pick the **borrower**. If the borrower does not exist yet, click *Create borrower* inline.
3. Scan or pick the **volume** (V-code).
4. Optionally set a custom **due date**; otherwise mybibli derives it from the overdue threshold (chapter 2).
5. Confirm.

The volume's current location is recorded on the loan row so a later return can restore it to the right shelf.

Returning a volume

1. Open `/loans` and find the active loan (filter on *Active* or scan the V-code).
2. Click **Return**.
3. Pick the post-return state (*Good*, *Worn*, ...).
4. Confirm. The volume snaps back to its pre-loan location automatically.

Overdue loans

Loans that exceed the overdue threshold turn red on the `/loans` page and on the home-page *Overdue* indicator. You can run `/loans?filter=overdue` for a focused list.

3.6 Library audit

An *audit* is the routine “walk the shelves and confirm everything is where the database says it is”. mybibli does not ship a dedicated audit mode — the equivalent workflow is:

1. Filter `/search` by the location you intend to audit, sorted by V-code.
2. Walk the shelf, scanning each V-code.
3. For any volume scanned that is *not* in the filtered list, mybibli surfaces a *Volume not in this location* message. Move it (section 3.4) and move on.
4. Any volume *in* the filtered list that you did *not* scan after walking the shelf is suspect — either misplaced or missing. Mark it for follow-up.

Tip

A quarterly audit catches misplaced volumes early — the longer a book has been on the wrong shelf, the harder it is to remember where it should be.

3.7 Shelving workflow

After a returns batch, the volumes pile up near the scanner station. mybibli's *shelving* flow speeds up putting them back:

1. Pick up a volume and scan its V-code.
2. mybibli shows the volume page with the original location prominently displayed.
3. Walk the volume to that shelf. Scan the L-code on the shelf to confirm.
4. Repeat.

Confirming with the L-code is optional but recommended — it leaves a fresh entry in the volume's location history that the next audit can rely on.

3.8 Wish list

The *wish list* is a separate list of books you'd like to buy, kept distinct from the owned catalog. It lives under `/wishlist` (see the navigation bar) and is available to Librarians and Admins.

3.8.1 Adding an entry

The *Add to wish list* form has two modes:

- **By ISBN.** Scan or type a 13-digit ISBN. mybibli looks the book up through the same metadata provider chain the catalog uses, then shows a preview card (title, author, publisher, year, cover). Click *Add to wish list* to save. Cover images come from the resolved provider's URL directly (e.g. Open Library Covers or Google Books); downloading the cover locally is on the roadmap.
- **Free-form.** Use this when you don't have an ISBN handy (a self-published title, a recommendation jotted down at lunch). Only *Title* is required; author, publisher and notes are optional.

3.8.2 Browsing at the bookstore

The list view collapses into a mobile-card stack on narrow viewports — designed to scan quickly with one thumb at a bookstore. Each card shows the title, the author (if known), the ISBN (if any), the publisher, and the date you added it.

To print a paper copy or save a PDF for offline use, click *Print* at the top of the list. mybibli renders a clean print-friendly page; use the browser's *Print to PDF* (File → Print → Save as PDF, or Ctrl/Cmd-P) to generate a PDF you can take to the bookstore.

3.8.3 Marking a book as bought

When you buy a book, open its wish list entry (or tap *Bought* directly from the list row), confirm in the modal, and the entry soft-deletes. The action is reversible from the Admin Trash panel

for 30 days.

3.8.4 Auto-detect on catalog scan

When you catalog a freshly bought book through the standard scan flow (`/catalog`), mybibli checks whether its ISBN matches an active wish list entry. If it does, an emerald banner appears in the feedback area: *This ISBN is on your wish list — Remove it?* One click on the banner's *Remove from wish list* button opens the same *Mark as bought* modal, and confirming removes the entry. No silent deletions.

Chapter 4

Multiple media types

mybibli is not just for books. Four media types are first-class citizens:

- **Book** — the default. Single-volume novels, essays, reference works, children’s picture books.
- **Comic / BD** — single-issue comics or multi-position omnibus albums (e.g. a “Tintin anthology” that physically gathers volumes 1–5).
- **Audio** — CDs, vinyl records, audiobooks on physical media.
- **Film / series** — DVDs, Blu-ray boxsets, TV-series collections.

The media type drives three things:

1. **Which metadata provider is consulted first.** ISBNs route to BnF + Google Books + Open Library; films route to TMDb + OMDb; audio routes to MusicBrainz.
2. **Which fields are exposed on the title page.** Audio titles surface a track list; films surface director, runtime, age rating.
3. **The icon shown next to the title** in search results and the catalog list.

4.1 Books

The default media type. Catalog as described in section 3.2. Multi-volume series (encyclopaedias, trilogies, ...) are modeled with a *Series* entity in chapter 2.

4.2 Comics and BD

A single **omnibus** volume can contain several positions in a series. mybibli supports this with the *multi-position* feature: at volume creation time, you can set the volume’s *position range* (for example, “contains volumes 5–7”). Series-gap detection (chapter 3) understands these ranges and will not flag positions 5, 6 or 7 as missing once the omnibus is on the shelf.

4.3 Audio releases

Audio titles use MusicBrainz as the primary metadata source. The barcode on the back of a CD (EAN-13) resolves the same way as an ISBN.

- **Tracks.** MusicBrainz returns the track list; mybibli renders it on the title page.
- **Format.** CD vs vinyl vs cassette is captured on the volume row (not on the title) — two identical pressings on different media are two volumes of the same title.

4.4 Films and series

Films and TV series use OMDb and TMDb as the primary metadata sources. A series boxset is modeled as a *title* with a multi-position volume covering the disc range.

Tip

OMDb is free for low volumes (under 1000 requests per day); TMDb is free with attribution. Either covers a household-sized DVD / Blu-ray shelf. See chapter 5 for setup.

4.5 Picking the media type at scan time

After you scan an unfamiliar barcode, mybibli guesses the media type from the EAN prefix (the GS1 country prefix and the resolved provider's response). The guess is shown on the catalog page; you can override it before clicking **Add volume** if it is wrong. Re-typing a title later is a few clicks (`/title/<id>`) and the provider chain re-runs against the new type if you trigger *Re-fetch metadata*.

Chapter 5

Metadata providers

When you scan an ISBN or an EAN-13, mybibli does not look the title up in a local catalog — it queries one or more external *metadata providers* until one returns a clean answer. This chapter explains the chain, where to get API keys, and how to rotate them safely.

5.1 The provider chain

mybibli walks the providers in the order best-suited to the guessed media type:

Provider	Media types	Notes	Key?
BnF	Books	French Bibliothèque nationale catalogue; best for French-language titles.	no
Open Library	Books	Volunteer catalogue; broad English / multilingual coverage.	no
Google Books	Books	Commercial; good for trade non-fiction.	optional
MusicBrainz	Audio	Volunteer music metadata; very deep on CDs / LPs.	no
OMDb	Films, series	Free tier; rate-limited.	required
TMDb	Films, series	Free with attribution.	required
BDGest	Comics, BD	FR-language comic catalogue (web-scraped).	no

The first provider to return a *clean* record wins. “Clean” means: title present, at least one contributor, and a year. Partial matches are recorded but mybibli keeps walking the chain until it finds a clean one or runs out of providers.

Cover-image fallback

When the resolving provider returns metadata but no cover URL (BnF never carries image URLs in its UNIMARC payload; Google Books occasionally lacks `imageLinks` for self-published titles) mybibli falls back to the *Open Library Covers API* (<https://covers.openlibrary.org/>) keyed by ISBN. That covers index is broad and independent of Open Library’s catalogue completeness, so most books still get a cover even when their text metadata comes from a different provider. The fallback runs only when (a) the chain succeeded *and* (b) the resolving provider gave no

`cover_url`, so a title for which no provider has a cover anywhere will land cover-less without retry storms.

5.2 API keys

Three providers require an API key:

Google Books

Optional but raises the daily request budget significantly. Without a key the Google Books provider stays disabled — the chain still works (BnF → Open Library are key-less), but English-language trade non-fiction that BnF misses will land with empty metadata.

The Google Cloud Console flow trips up many first-timers because the API-key creation step is buried two menus deep. Walk through this exactly:

1. Open <https://console.cloud.google.com/> and sign in.
2. Top-bar project picker (next to the *Google Cloud* logo) → *New project* → name it anything (e.g. *mybibli-metadata*) → *Create*. Wait a few seconds for the project to be selected.
3. Open the hamburger menu (top-left, three horizontal lines) → *APIs & Services* → *Library*. Search for **Books API**, click it, then *Enable*.
4. Still under *APIs & Services*, open *Credentials*. Click *+ Create credentials* → *API key*. A modal pops up with a freshly generated key — copy it now.
5. Recommended: click *Restrict key* → under *API restrictions* pick *Restrict key* → tick only *Books API*. Leave *Application restrictions* on *None* (mybibli does not send a Referer header). Save.
6. In mybibli: `/admin > System > Metadata Providers` → paste into the **Google Books API key** field → *Save*. The next scan picks it up — no restart needed.

To verify the key is active, scan an English-only ISBN that BnF will not have (*Effective Java*, 9780134685991, is a good canary). The metadata should populate within a few seconds. If it stays empty, see *Troubleshooting* at the end of the chapter.

OMDb

1. Visit <https://www.omdbapi.com/apikey.aspx>.
2. Pick the free tier (1000 requests/day).
3. Confirm via the activation email you receive.
4. Paste the key into mybibli.

TMDb

1. Sign up at <https://www.themoviedb.org/>.
2. Open *Settings* → *API* and request a key for *Personal / non-commercial use*.
3. Paste the key into mybibli.

5.3 Where the keys live

Keys travel through two layers:

1. **Environment variables** (`GOOGLE_BOOKS_API_KEY`, `OMDB_API_KEY`, `TMDB_API_KEY`). These are migrated *once* on boot into the `settings` table — they are essentially a deployment-time bootstrap.
2. **The settings table**, exposed through `/admin > System > Metadata Providers`. This is the live source of truth at request time. Rotating a key from the UI takes effect on the very next metadata fetch.

The implication is asymmetrical:

- If you set a key only in the environment, the first boot copies it into the database and afterwards the database wins.
- If you set a key only in the UI, the environment is never consulted again.
- If you set both and then *Clear* the key from the UI, the database value goes empty — but on the next boot, the environment value is migrated back in. To make the UI clear durable, also remove the environment variable from `.env` / `docker-compose.yml`.

5.4 Provider health

`/admin > Health` pings every registered provider every five minutes and reports their status as a green check or a red cross. A red cross does not always mean the provider is down — it can also mean the API key is missing or invalid. The Health tab includes the last-checked timestamp and the HTTP status code for diagnostics.

5.5 Rate limits and back-pressure

mybibli applies a small per-provider rate limit (a couple of requests per second). Bulk-cataloging a hundred volumes in a row will not trigger remote 429 responses. If a provider does return 429, the request is logged and the next provider in the chain is consulted instead — the catalog flow never blocks waiting for a rate-limited provider.

Warning

Never commit a real API key to git. The `.env` file is in `.gitignore` for this reason; `.env.example` carries only placeholders. If a key leaks, rotate it at the provider console immediately — pasting a new key into mybibli takes effect within seconds.

5.6 Troubleshooting empty metadata

If a freshly scanned ISBN ends up with empty title / author / description fields, walk this list before suspecting a bug:

1. **Provider status in `/admin > Health`**. If Google Books or Open Library show a red cross, those providers will be skipped on every fetch. Re-check the last-checked HTTP code — 403 or 401 usually means a missing or invalid API key, not a remote outage.

2. **Language of the title.** BnF is the primary provider and is FR-leaning. For English-only books, the chain falls through to Google Books — which means the key (see above) is the unblocker.
3. **ISBN format.** mybibli normalises EAN-13 vs ISBN-10 internally, but some providers index only one of the two. A second scan a few minutes later sometimes hits a different fallback path.
4. **Manually edit and save.** You can always edit title, genre, contributors etc. by hand from the title's detail page — saving works even when no provider returned anything. The genre `Non classé` is the default for any title that came in without classification data.

If the fields are still empty after configuring the right keys and waiting a few seconds, open an issue on the project repo with the failing ISBN — the maintainers add provider mocks for reproducible cases.

Chapter 6

Roles and permissions

mybibli has three user roles ordered by privilege:

1. **Anonymous** — anyone who lands on the site without a session cookie. Read-only access to the catalog. Cannot loan, edit or delete anything.
2. **Librarian** — authenticated user whose role is `librarian`. Day-to-day cataloging, loans, returns, location moves. Cannot manage users or change system settings.
3. **Admin** — authenticated user whose role is `admin`. Everything a Librarian can do, plus the `/admin` panel: users, reference data, trash, system settings.

6.1 Who can do what

Action	Anon	Lib	Admin
Browse / search the catalog	yes	yes	yes
View title and volume pages	yes	yes	yes
View location tree	yes	yes	yes
Scan / catalog a new title	no	yes	yes
Add or remove a volume	no	yes	yes
Move a volume between locations	no	yes	yes
Edit a title's metadata	no	yes	yes
Create / edit borrowers	no	yes	yes
Register / return loans	no	yes	yes
Edit storage location tree	no	yes	yes
Open <code>/admin</code>	no	no	yes
Create / deactivate users	no	no	yes
Edit reference data	no	no	yes
Restore from trash	no	no	yes
Permanently delete from trash	no	no	yes
Edit system settings	no	no	yes

6.2 Sessions and the inactivity timeout

After login, mybibli sets a session cookie that travels with every subsequent request. The cookie is:

- `HttpOnly` — not exposed to JavaScript;
- `SameSite=Lax` — sent on top-level same-site navigations but not on cross-site requests;
- `Secure` (only when `MYBIBLI_COOKIE_SECURE` is set to `true` — see chapter 1).

The session has an **inactivity timeout** (default 30 minutes, configurable from `/admin > System > Loans`). Every authenticated request resets the countdown; if no request lands within the timeout window, the session is wiped server-side and the next request redirects to the login page.

A small *keep-alive toast* appears in the corner of the screen roughly two minutes before expiry — click it to extend the session without leaving your current page.

6.3 The last-active-admin guard

mybibli refuses to leave the installation without any active admin. The guard applies to two actions:

- **Deactivating a user.** If the user is the only active admin, the request returns a 409 conflict and the deactivation is refused.
- **Permanently deleting a user from the trash.** Same rule.

Self-deactivation is also blocked unconditionally (you cannot deactivate the account you are currently logged in as).

Warning

The guard counts only *active* admins. If you mass-deactivate admins and the guard fires, undo by reactivating someone from the trash tab. There is no recovery path through the wizard or the database without manual SQL surgery.

6.4 Password reset

mybibli does not implement password-reset emails — household / small-community deployments rarely benefit from the email plumbing. An admin can reset another user's password by:

1. Open `/admin > Users`.
2. Pick the user, click **Deactivate**.
3. Click **Reactivate** from the trash tab.
4. Re-create the user with a fresh password.

A user who knows their own password can change it from the `/account` page.

6.5 Anonymous sessions

Anonymous visitors also get a session row (used for CSRF protection and pending HTMX updates). Anonymous sessions expire after seven days of inactivity and are garbage-collected by a daily background task — they do not affect logged-in users.

Chapter 7

Backup and restore

A self-hosted library has no cloud safety net. This chapter covers what to back up, how often, and how to restore.

7.1 What to back up

Three artifacts together form a complete mybibli backup:

1. **The MariaDB database.** Holds every title, volume, location, contributor, loan, user, audit log and setting. Lose this and you start from zero.
2. **Cover images.** Downloaded during cataloging and persisted in the Docker volume `mybibli_covers` (or in your host bind-mount path if you switched per section 1.3). Lose this and the catalog still works, but every title page shows a generic placeholder until you re-fetch metadata title-by-title (issue #213 — pre-1.1.4 installs without this volume had covers disappear on every container upgrade).
3. **The `.env` file.** Holds passwords (database, API keys not yet migrated). Lose this and a re-deploy needs you to re-enter every credential.

The Docker images themselves are *not* backup-critical — you can always pull them again from Docker Hub.

7.2 Backup procedure

Database — bundled MariaDB

From the host running the compose stack:

```
docker compose exec -T db mariadb-dump \  
  -u root -p"$MYSQL_ROOT_PASSWORD" \  
  --single-transaction --routines --triggers \  
  mybibli \  
  > backups/mybibli-$(date +%Y%m%d-%H%M%S).sql
```

The `--single-transaction` flag is important: it dumps a consistent snapshot without locking the table for the duration of the dump. `--routines --triggers` preserves any stored routines (mybibli does not currently use any, but the flag is cheap and future-proof).

Compress the dump if disk space is a concern:

```
gzip backups/mybibli-*.sql
```

Database — external MariaDB

Run `mariadb-dump` directly against the external server. Synology DSM users can also schedule a *Hyper Backup* task that targets the MariaDB package — it produces a binary backup that restores faster than a SQL dump.

Covers and `.env`

For a host-mounted covers directory (bind-mount installs), a plain `tar` archive is enough:

```
tar czf backups/mybibli-files-$(date +%Y%m%d).tar.gz \
.env "$COVERS_HOST_PATH"
```

For the default install (named Docker volume), back up the volume through Docker — the host path of a named volume varies by Docker version and is not stable, so don't try to `tar` it directly:

```
docker run --rm \
-v mybibli_covers:/source:ro \
-v "$(pwd)/backups:/backup" \
alpine \
tar czf "/backup/mybibli-covers-$(date +%Y%m%d).tar.gz" -C /source .
tar czf "backups/mybibli-env-$(date +%Y%m%d).tar.gz" .env
```

The throwaway `alpine` container mounts the named volume read-only, the host `backups/` directory writable, and writes the tarball. Same shape works for restoring — see the Restore section below.

7.3 Recommended cadence

- **Daily** — automated database dump. Even on a seldom-changed catalog, a daily dump means worst-case data loss is bounded to one day.
- **Weekly** — covers and `.env`. These change only when you catalog new titles or rotate keys; weekly is plenty.
- **Off-site copy** — at least monthly. Rsync the `backups/` directory to a remote NAS, a USB drive you swap in / out, or a cloud-storage bucket. A fire that consumes the home server should not also consume your backups.

Tip

A backup that has never been restored is not a backup. Once a quarter, restore your latest dump into a throwaway docker environment and verify the title and loan counts match what `/admin > Health` reports on production.

7.4 Restore procedure

Database

For the bundled-database setup:

```
docker compose stop mybibli
docker compose exec -T db mariadb -u root -p"$MYSQL_ROOT_PASSWORD" \
    -e "DROP DATABASE mybibli; CREATE DATABASE mybibli CHARACTER SET utf8mb4;"
zcat backups/mybibli-20260513-1200.sql.gz | \
    docker compose exec -T db mariadb -u root -p"$MYSQL_ROOT_PASSWORD" mybibli
docker compose start mybibli
```

For external MariaDB, the same `mariadb` client commands work against the external server.

Covers and `.env`

For a host-bind-mount install:

```
tar xzf backups/mybibli-files-20260513.tar.gz
```

For the default named-volume install, restore through a throwaway container in the symmetric shape to the backup:

```
docker compose stop mybibli
docker run --rm \
    -v mybibli_covers:/target \
    -v "$(pwd)/backups:/backup:ro" \
    alpine \
    sh -c "rm -rf /target/* && tar xzf /backup/mybibli-covers-20260513.tar.gz -C /target"
tar xzf backups/mybibli-env-20260513.tar.gz
docker compose start mybibli
```

Restart the stack so the `mybibli` container picks up the restored `.env`.

7.5 Migration to a new host

The backup procedure is also the migration procedure:

1. On the old host: produce a fresh database dump and tar archive.
2. Transfer both to the new host.
3. On the new host: install Docker, pull the compose file from the release matching the dump's version, place `.env` and `covers/` alongside, restore the database, start the stack.
4. Verify by counting titles on `/admin > Health` and spot-checking a few title pages.

Warning

The dump is tied to the *schema* version. If the source host ran 1.1.0 and the new host pulls 1.5.0, the new host's migrations will upgrade the restored database in place — that is fine. The reverse (restoring a 1.5.0 dump onto a 1.1.0 install) is *not* supported and will fail at boot. Always restore on a host running the same or newer `mybibli` release.

Chapter 8

Upgrade and migration

Routine upgrades are designed to be safe. The exception is the 1.0.x to 1.1.0 jump, which is a one-time breaking change documented below.

8.1 Routine upgrades (1.y to 1.z)

1. Read the release notes for every version between your current and the target version.
2. Take a fresh **database backup** (chapter 7).
3. Update the image tag in `docker-compose.yml` (or repull `:latest` if you track that tag).
4. Run `docker compose up -d`. The new container boots, runs any new migrations against the existing database, and starts serving requests.
5. Watch the logs for migration errors: `docker compose logs mybibli | grep -i migrate`.

Migrations are append-only — schema changes are forward compatible. A 1.4 install upgraded to 1.5 keeps every row from 1.4 and applies new tables / columns in place.

Tip

For homelab installs, the `:latest` Docker tag is a reasonable default. For production-style setups (a community library, an association), pin to an explicit version tag so an unattended pull does not bring in a release you have not read about.

8.2 The 1.0.x to 1.1.0 jump

Warning

This is a one-time breaking change. 1.0.x installs are subject to the issue documented in [#173](#): the seed migrations create the credentials `admin/admin` and `librarian/librarian` on every fresh install, including production, and silently bypass the first-launch setup wizard.

The recommended path depends on whether your 1.0.x install holds real data:

If your 1.0.x install is empty (no titles cataloged)

The simplest path: wipe the database and reinstall on 1.1.0+.

1. Stop the stack: `docker compose down`.
2. Remove the database volume: `docker volume rm mybibli_mybibli_db_data` (or the equivalent volume name on your host — `docker volume ls` lists them).
3. Pull the 1.1.0 image: `docker compose pull`.
4. Update the image tag in `docker-compose.yml` if you pinned a previous version.
5. Start: `docker compose up -d`.
6. Open `/setup` and walk through the wizard.

If your 1.0.x install holds real data

1. Take a database backup (chapter 7).
2. Take a covers + `.env` backup.
3. Log in as the seeded admin and rotate the password *before* you upgrade. The 1.1.0 upgrade flow itself is safe, but rotation closes the only window where the default credentials could leak.
4. Update the image tag in `docker-compose.yml` and run `docker compose up -d`.
5. The 1.1.0 boot detects the existing admin and does *not* re-enable the wizard — the rotated admin remains the active admin.

The new MYBIBLI_SEED_DEV_USERS variable

Starting with 1.1.0, the seed migration is gated by `MYBIBLI_SEED_DEV_USERS`. The default is 0 (no seed). Only set it to 1 for local-development workflows or automated tests — never in production. The first-launch wizard becomes the canonical onboarding path.

8.3 Reading the release notes

Every mybibli release publishes a GitHub Release page with:

- a one-paragraph summary of the change set;
- the migration impact (none / additive / breaking);
- a link to the diff against the previous tag;
- attached PDFs of this manual (English + French) once the release-time PDF build is wired into CI.

Read the notes before upgrading. Breaking changes are tagged with a prominent *Breaking* label.

8.4 What's new in 1.7.1

Version 1.7.1 is a patch release that closes the v1.7.0 #301 gap and bundles four production-surfaced bugfixes from the user's v1.7.0 NAS upgrade. *No schema migration affecting catalog*

`data`; routine `docker compose pull && docker compose up -d`. One new settings row is seeded (`log_level`) by an idempotent migration.

Six fixes:

- **Runtime log-level admin UI (#308).** Finishes the v1.7.0 #301 feature properly. `/admin > System` gains a **Logging** section with a free-form `EnvFilter` directive field. Click Save and the next log line uses the new level — no `docker compose up -d`, no restart. Persisted to the `settings` table so the choice survives container restarts. The v1.7.0 release notes had advertised this surface but only the env-var bootstrap + rotation actually shipped; this release ships the missing piece.
- **API-key permanent-delete 422 (#309).** The `DELETE /admin/api-keys/{id}` handler was returning 422 Unprocessable Entity because it required a request-body form, but HTMX `hx-delete` puts form fields in the query string. Drop the body extractor; the CSRF middleware already validates the X-CSRF-Token header. Latent regression since v1.5.1 #284.
- **Provider-health probe timeout configurable, default bumped 3 s → 10 s (#310).** The hardcoded 3 s probe timeout was too tight for typical home-NAS DNS + TLS handshake; every probe was timing out, making `/admin > Health` permanently red even when providers were reachable for the actual metadata-fetch path. New env var `MYBIBLI_PROVIDER_HEALTH_TIMEOUT` (default 10) lets ops bump to 20–30 s on fragile uplinks or down to 4–5 s on fast infra.
- **Bulk cover-refetch now actually re-queries providers (#311).** The admin "Re-fetch missing covers" action was silently no-op'ing on cached titles because the metadata chain short-circuited on the cache hit set when the title was originally cataloged. New `force_refresh` path invalidates the cache before re-running the chain, so Google Books / BD Gest / etc. get re-asked for a cover URL. List-mode after HTMX search swap also restored (#315).
- **DB pool size 10 → 20 (#312).** A transient pool timed out while waiting for an open connection was observed on the user's prod NAS during a 64-title bulk-refetch. The `sqlx` default 10 was too tight for HTTP handlers + 3 always-on background tasks + concurrent admin HTMX polling + the per-title DB hits of the bulk worker. 20 gives comfortable headroom for single-household / small-association load.
- **Home filter chips: drop duplicate "Non classé" chip (#314).** The home page was rendering both a "Non classé" chip (auto-generated from the placeholder genre row) and a "Sans genre" chip (the dedicated #205 affordance) — both returning the same titles. Strip the placeholder from the chip list; the dedicated chip remains.

8.5 What's new in 1.7.0

Version 1.7.0 is the fifth themed minor of the post-v1 era — **Reach more users, debug more easily**. No schema migration, no breaking changes — routine `docker compose pull && docker compose up -d`. Operators who want to keep their logs across container restarts should bind-mount a host directory to `/var/log/mybibli` (see chapter 12); without the mount, logs still work but are ephemeral, exactly as before.

Three CRs, plus a non-user-visible foundation slice:

- **German UI translation (#275).** ~900 keys translated from `en.yml` to a new `de.yml`. Sie-form throughout (formal "you" — appropriate for library / small-association deployments).

Reference data (genre names, location-type names, condition states) stays in whatever language operators entered — only UI chrome is localized.

- **Italian UI translation (#276).** ~900 keys translated to a new `it.yml`. Tu-form informal throughout — appropriate for the home-library / small-association deployment shape. Italian guillemets «...» (U+00AB / U+00BB) used for inline quotes.
- **Persistent log directory + rotation + admin-controlled level (#301).** Logs now write to `/var/log/mybibli/` (bind-mountable to a host directory) with daily rotation and configurable retention. Admins can change the log level (`trace` / `debug` / `info` / `warn` / `error`) from the System settings tab without a redeploy. New chapter 12 (“Operations & debugging”) documents the full setup.

Foundation slice (prerequisite for the two UI translations, shipped under PR #302): `src/i18n/resolve.rs`, `src/middleware/locale.rs`, the admin → System → Default language selector, and the nav-bar language toggle were generalized to handle three or more languages. A new locale parity test (`tests/locale_parity.rs`) ensures all four locale files stay key-aligned on every CI run.

Out of scope (deferred to a future minor): German and Italian translations of this manual. The UI is in four languages; the manual remains in English + French.

8.6 What’s new in 1.6.2

Version 1.6.2 is a two-part hygiene patch on top of v1.6.1. *No code changes.* Routine `docker compose pull && docker compose up -d`.

Part 1 — dependency bump. `cargo` update applied 74 caret-compatible bumps across the dependency tree:

- `axum` 0.8.8 → 0.8.9
- `tokio` 1.50.0 → 1.52.3
- `openssl` 0.10.76 → 0.10.80 (security-relevant)
- `rustls` 0.23.37 → 0.23.40
- `hyper` 1.8.1 → 1.9.0
- `tower-http` 0.6.8 → 0.6.11
- `wasm-bindgen` 0.2.115 → 0.2.121
- the `icu_*` / `rustls-*` / `zerocopy` stacks
- (full list in the `Cargo.lock` diff against v1.6.1)

`cargo audit` reports no new CVEs. Two pre-existing upstream-blocked advisories remain: `rsa` 0.9.10 (RUSTSEC-2023-0071, medium, no fixed upgrade — comes via `sqlx-mysql`) and `time` 0.2.27 (RUSTSEC-2020-0056, comes via `genpdf` → `printpdf` → `lopdf`).

Part 2 — reproducible builds. `Cargo.lock` removed from `.gitignore` and committed for the first time. `mybibli` ships as a binary application, not a library, so the standard Rust guidance is to commit the lockfile. Pinning gives reproducible builds across (a) the Dockerfile multi-stage build, (b) the GitHub Actions release workflow, and (c) `git checkout vX.Y.Z` — a tag’s build deterministically matches the shipped Docker image. Previously the lockfile was git-ignored and the Docker build re-resolved versions at build time — two `:1.6.x` images built a week apart

could silently pull different upstream patches. One-time structural fix; future bumps will surface as `Cargo.lock` diffs in PR reviews.

909 unit tests + integration suite + `sqlx-prepare --check` all green; `cargo clippy --all-targets -- -D warnings` clean.

8.7 What's new in 1.6.1

Version 1.6.1 is a one-fix hotfix shipped hours after the v1.6.0 NAS upgrade. Routine `docker compose pull && docker compose up -d`, no schema migration.

Empty `parent_id` on root-location edit no longer 400s (#296). Repro: tick the new "Emplacement organisationnel" checkbox on a *root* location, submit → 400 Failed to deserialize form body: `parent_id: cannot parse integer from empty string`. Worked fine on non-root locations.

Root cause is pre-existing — latent since the location-edit feature shipped. The form's `<select name="parent_id"><option value="">None</option>...</select>` submits `parent_id=` (empty string) when the user keeps the root option. `#[serde(default)] Option<u64>` only short-circuits *absent* fields; an empty string still hits the u64 parser, which bails. CR #280 just gave users a new reason to edit root locations, exposing the latent bug.

Fix mirrors the existing pattern for `deserialize_optional_i32` / `deserialize_optional_f64` in `routes::series`:

- New `deserialize_optional_u64` helper treats "" / whitespace / absent as `None`, parses non-empty values as `u64`.
- Wired to `CreateLocationForm.parent_id`, `UpdateLocationForm.parent_id`, AND `VolumeEditForm.condition` (same `<option value=""></option>` shape, same latent bug — picking "no condition" on a volume would have 400'd the same way).

12 new unit + form round-trip tests; new E2E spec walks through create root → edit → tick organisationnel → submit → `/locations`. Earlier client-side workaround that stripped the `name="parent_id"` attribute in the location-edit form to dodge the 422 was removed.

8.8 What's new in 1.6.0

Version 1.6.0 is the fourth themed minor of the post-v1 era — **Tighten the catalog**. One additive schema migration (`volumes.under_audit_since` `DATETIME NULL`); no breaking changes — routine `docker compose pull && docker compose up -d`.

Four CRs aimed at catalog hygiene:

- **Organizational storage locations (#280).** A storage location can be marked "organizational" to refuse volume assignments — folders, not shelves. Checkbox on the create/edit form, flip-with-volumes guard (cannot mark organizational a location that already holds volumes), volume-edit picker rejects organizational targets at submit time with a localized feedback entry.
- **Shelf-audit workflow (#237).** Mark a volume "À contrôler" — single or bulk for an entire shelf. An amber home-dashboard indicator counts pending audits; a dedicated `/audit` list (sorted by location → V-code) shows them all. New `volumes.under_audit_since` column captures the mark timestamp; `admin_audit` rows are written for every mark and

every clear. Surplus-detection (locations with volumes the operator didn't expect) deferred to a follow-up.

- **"Titles without volumes" filter chip (#279).** New chip on the home dashboard surfaces titles that have no active volume row — either all soft-deleted away or catalogued but never shelved. Librarian+ only. SQL guard is `NOT EXISTS` against the `volumes` table with the `deleted_at IS NULL` predicate, so soft-deleted volumes don't keep their title from appearing.
- **Sortable home list view (#250).** The home page's list view becomes a sortable table — Title · Author · Genre · Dewey · Volumes. Replaces the previous cards-stacked-vertically layout. Column headers issue `HTMX hx-get + hx-push-url=true` mirroring the `/location/:id` pattern. Grid view unchanged; the sort dropdown is hidden in list mode (column clicks own sorting).

8.9 What's new in 1.5.2

Version 1.5.2 is a one-fix hotfix on top of v1.5.1: the v1.5.0 #243 valuation workflow was partially broken in production. Saving a *Purchase price* or *Estimated current value* on a volume landed the `UPDATE` (the totals on `/stats/value` and the home indicator both reflected the new amount) but returned a 500 to the user, and every subsequent view of the affected volume crashed.

Root cause: `SQLx 0.8` can't decode a raw MariaDB `DECIMAL(10,2)` into `Option<f64>` without the `bigdecimal / rust_decimal` feature flag. The write side worked because the driver coerces `f64` into `DECIMAL`; the read side didn't. The aggregation queries (`/stats/value` totals) already used `CAST(SUM(...)) AS DOUBLE` and were unaffected.

Fix: `VolumeModel::find_by_id` and `find_by_label` now wrap the two `DECIMAL` columns with `CAST(... AS DOUBLE)`. Two new `sqlx::test` regression guards `INSERT` a volume with non-NULL valuation columns and assert round-trip so the bug can't silently regress on a future refactor.

No schema migration, no breaking changes — routine `docker compose pull && docker compose up -d`.

8.10 What's new in 1.5.1

Version 1.5.1 is a five-fix production-hygiene patch shipped within hours of the v1.5.0 NAS upgrade. No schema migrations, no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Wish list PDF crash (#281).** The v1.5.0 Dockerfile never received `COPY static/fonts/` when CR #266 landed the vendored DejaVu Sans TTFs, so `/wishlist/export.pdf` returned 500 in production. Fix: one line added to the Dockerfile.
- **Trash permanent-delete crashed on titles (#282).** The 4 child FKs that reference titles (`volumes`, `title_contributors`, `title_series`, `pending_metadata_updates`) default to `RESTRICT` and MariaDB rejected the parent delete as soon as any child row referenced it — even soft-deleted children, because FK constraints don't consider `deleted_at`. Fix: `TrashService::permanent_delete` now opens a transaction and cascades through the FK-ordered children before deleting the parent.
- **Library valuation toggle was missing from /admin?tab=system (#283).** The v1.5.0 #243 added `default_currency + show_value_indicators` in the settings table and the

AppSettings cache but no UI was wired to flip them. Fix: a new "Library valuation" section on the System tab with a currency dropdown + the opt-in indicator checkbox.

- **No path to delete a revoked API key (#284).** The v1.4.0 #241 implementation deliberately preserved revoked rows for audit, but the list grew monotonically. Fix: a Delete button on revoked rows (UX-DR8 modal) calling a new DELETE /admin/api-keys/{id} handler. Active keys still show only Revoke. The audit references to deleted keys survive (the link is logical, not a schema FK).
- **Admin → Health tab marked every provider red (#285).** probe_once was sending HEAD without a User-Agent and treating 4xx responses as "down" — most providers' WAFs answer 4xx to anonymous HEADs even when the API works. Fix: the probe now sends User-Agent: mybibli/1.x.y matching the actual metadata-fetch path and treats any HTTP response as Reachable (only timeouts / DNS / refused connections count as Unreachable).

8.11 What's new in 1.5.0

Version 1.5.0 is the fourth themed minor of the post-v1 era — **Know what you have, and what it's worth.** One additive schema migration (volumes table gains 5 nullable valuation columns + two new settings rows), no breaking changes — routine docker compose pull && docker compose up -d.

- **Library of Congress metadata provider (#263).** Authoritative bibliographic source for English-language titles, slotted in the chain between Google Books and Open Library: BnF → Google Books → Library of Congress → Open Library. Fills the older / academic / US-government EN-book gap where Google Books returns nothing.
- **Donut chart for stats-by-genre on home (#265).** Replaces the per-row <progress> bars with a Chart.js v4 doughnut. Genres at ≥ 5% of the active catalog get their own slice; genres below the threshold roll up into a single "Other" slice with a <details> block listing the rolled-up names. Click on a slice or legend item navigates to /?filter=genre:N. Chart.js is vendored under static/js/vendor/ — no CDN, CSP stays strict.
- **Delete a title with no volumes (#271).** A new Delete button on /title/:id (Librarian+, visible only when the title has zero active volumes) opens a UX-DR8 destructive-action modal. On confirmation the title is soft-deleted and the user is redirected home; the row is recoverable from the Trash panel for 30 days.
- **Edit ISBN / ISSN / UPC in the metadata form (#272).** The previously read-only identifier lines in the title-edit form are now editable inputs. ISBN-13 checksum is validated server-side; a collision with another title's existing ISBN returns 409 + a friendly feedback entry. Companion to #271 — both fix the wrong-scan-stuck-title scenario.
- **Real PDF export for the wish list (#266).** Replaces the v1.3.1 browser-print HTML round-trip with a server-rendered PDF via genpdf (pure Rust, no subprocess). The Print button on /wishlist now downloads wishlist-YYYY-MM-DD.pdf directly — text is selectable, searchable, consistent across all browsers. DejaVu Sans vendored for unicode support.
- **Collection valuation (#243).** Two new optional fields per volume — *Purchase price* and *Estimated current value* — each with its own ISO 4217 currency. Aggregations on a new /stats/value page (Librarian-gated): total per currency, value by genre, value by series. Mixed-currency catalogs show one row per currency (FX-conversion deferred to a future CR). Opt-in home dashboard indicator — default OFF, admin toggle under

`/admin?tab=system`. The default currency is seeded to CHF and is admin-overridable. All value-related views are owner-facing (Librarian+); Anonymous visitors never see money.

8.12 What's new in 1.4.0

Version 1.4.0 is the third themed minor of the post-v1 era — **Bring your own agent**. One additive schema migration (`api_keys` table), no breaking changes — routine `docker compose pull && docker compose up -d`.

- **JSON HTTP API at `/api/v1/*` (#241).** A new authenticated API exposes the catalog to external scripts and LLM-powered tools. Read endpoints cover titles (list, detail, by-Dewey bucket), genres, locations, and series. A focused `PATCH /api/v1/titles/{id}` write endpoint updates four whitelisted fields (subtitle, description, Dewey code, genre) under optimistic locking, with a full audit row per change. See the new chapter 11 *API & integrations* for endpoint contracts, authentication, and a Python / curl / LLM-loop example.
- **API keys admin tab.** New `/admin?tab=api_keys` surface to mint, list, and revoke keys. Each new key's plaintext is shown *exactly once* in a UX-DR8 modal — argon2 hash + 12-character prefix persisted afterwards. Two scopes: *read only* (GET endpoints) and *read + write* (adds PATCH). The list shows label, scope chip, prefix, created timestamp, last-used timestamp, and status (active / revoked). Revoke is reversible-ish in the sense that you can always re-create a key with the same label, but a revoked key can never authenticate again.
- **No CSRF on `/api/*`.** Bearer-token auth doesn't ride on cookies, so the CSRF middleware short-circuits the entire `/api/*` prefix. Form endpoints under the human UI remain CSRF-protected.
- **Manual chapter 11 *API & integrations* (EN + FR).** Walks through key creation, the authentication shape, every endpoint, and three end-to-end examples — a curl crib sheet, a Python pagination loop, and an LLM classifier loop in pseudocode.

8.13 What's new in 1.3.1

Version 1.3.1 is a four-papercut polish patch surfaced by Guy within hours of v1.3.0's production deployment. No schema migrations, no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Loading spinner on the wish list "Rechercher" / "Look up" button (#258).** The ISBN lookup runs the full provider chain (BnF / Open Library / Google Books / ...) which takes 2 to 5 seconds. v1.3.0 shipped without a visual indicator, leaving the user to wonder whether their click had registered. The button now shows a spinner during the in-flight request and is disabled to prevent double-clicks.
- **Help-icon tooltip on the Omnibus checkbox (#259).** The Omnibus option in the series-assignment form on the title-detail page used to be opaque — a freshly-deployed user had to ask out loud what it does. A new help-icon tooltip now explains: *check this if the book bundles multiple series volumes; the form reveals an End-position field so the title covers a range of positions instead of just one*.
- **Polished wish list print page (#260).** The Print button now opens `/wishlist/print` in a *new tab* so closing the print dialog doesn't lose the wish list view. The print page

itself was rebuilt from a dense table to a card-based layout with cover thumbnails, a tighter typography hierarchy (title bold, author muted), a footer with the page counter, and CSS guards so a single entry never splits across two pages. Take the PDF to the bookstore and enjoy.

- **Folded "Recent additions" on the home page (#261).** As the catalog grows, the home dashboard becomes tall. The "Recent additions" section is now wrapped in a native `<details>` block, closed by default, with the open / closed state persisted across reloads via `localStorage`. Click the heading to expand. The dashboard indicator filters (overdue, unshelved, gaps, recent-cataloged, recent-returns) remain expanded since they're the result of an explicit filter click.

8.14 What's new in 1.3.0

Version 1.3.0 is the second themed minor of the post-v1 era — **Plan your next purchase**. One additive schema migration (`wishlist_items` table), no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Per-volume table on the title-detail page (#209).** Open any title and scroll past the contributors block : you now see an inline list of every active volume for that title, with the V-code (clickable, links to `/volume/:id`), the resolved location name and L-code, the condition state, and — for Librarians and Admins — Edit + Delete affordances. The Delete button opens a UX-DR8 confirmation modal ; success redirects back to the title-detail page with the volumes list re-rendered without the deleted row. On a mobile viewport, the table collapses into the established card-stack pattern. Closes GH issue #209.
- **Wish list as a first-class surface (#242).** A new `/wishlist` surface, available from the navigation bar (Librarian+ for mutations, browseable to anyone). Two add modes : scan or type a 13-digit ISBN to let mybibli pull the metadata through the existing provider chain (preview card with cover, then Confirm) — or pick *Free-form* and just type a title (everything else optional). Each entry can be marked as *Bought* from the list, the detail page, OR automatically when the same ISBN is later cataloged through `/catalog` — mybibli surfaces an emerald banner on the scan feedback offering *Remove from wish list* in one click (no silent deletions). The `/wishlist/print` page renders a clean print-friendly layout ; use the browser's *Print to PDF* to take the list along to the bookstore. A new home-dashboard chip shows the wish-list count (hidden when empty). Full audit-trail entries on add and on mark-as-bought. See [section 3.8](#) for the full walk-through. Closes GH issue #242.

The 1.2.x patch line stays in maintenance for production bug fixes that don't warrant a feature bump.

8.15 What's new in 1.2.2

Version 1.2.2 is a three-fix *defense-in-depth* patch on top of 1.2.1, focused on the soft-deleted-genre semantics. No schema migrations, no breaking changes — routine `docker compose pull && docker compose up -d`.

Story 8-4's deletion guard prevents the orphan-FK state via the supported delete path: you can't actually soft-delete a genre that still has titles assigned. These fixes harden the read-side contract so a future migration, a manual SQL `UPDATE`, or removal of the guard wouldn't silently mislead the librarian.

- **Stale genre-filter links no longer return empty lists silently.** Open the home page in two browser tabs, soft-delete a genre from the admin reference-data tab in one of them, then click the genre’s filter chip in the other. Before: the catalog SQL filtered by the now soft-deleted genre id, the LIST came back empty, and you had no explanation of why your filter showed zero results. Now: the home handler validates the filter’s target genre, silently clears the chip if it points at a deleted genre, and surfaces a localized info notice (*This genre no longer exists — showing all titles.*) above the results so you understand what happened. Closes GH issue #112.
- **Orphan-genre titles no longer disappear from the catalog.** The four catalog SQL queries (search results, recent-additions, recent-cataloged) now use a left outer-join on the genre table with a COALESCE on the genre name. A title whose genre was soft-deleted (orphan FK) used to silently vanish from every catalog listing; it now appears with a localized placeholder (*Deleted genre*) / (*Genre supprimé*) so you can find it, re-assign a real genre, and move on. Closes GH issue #107.
- **Stats-by-genre percentages now reflect the true catalog share.** The home dashboard’s per-genre percentage panel used to compute its denominator from the sum of displayed rows. With orphan-FK active titles, this inflated each displayed row’s percentage (12 Roman + 8 BD + 5 orphan-Z would display Roman at 60% instead of the truthful 48%). The denominator now comes from the full active-title count emitted by the glance card. When orphan-FK titles exist, the displayed sum is <100%; the gap is the orphan delta and signals that some titles point at a soft-deleted genre. Closes GH issue #111.

8.16 What’s new in 1.2.1

Version 1.2.1 is a four-fix patch on top of 1.2.0. No schema migrations, no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Deleting a series with assigned titles now shows the meaningful error message.** Trying to delete a series that still has titles attached used to surface a generic *An internal error occurred* feedback strip; the actual cause (*Cannot delete X: N title(s) assigned. Remove all title assignments first.*) only reached the server log. The handler now also recognises the `Conflict` variant returned by the title-assignment guard, so the user sees the actionable message and knows to unassign the titles first. Closes GH issue #139.
- **Help-icon button no longer sits inside its associated `<label>` element.** Six forms (setup wizard admin step, setup wizard preferences step, borrowers create/edit email & phone, series-type, volume condition) rendered the help-icon `<button>` nested inside the form-control `<label>`. HTML5 forbids interactive content inside `<label>`; some screen readers also propagated label clicks to the help button, producing stray focus jumps. The button is now a sibling of the label in a small flex container, with the tooltip popping up below the trigger so it never overlaps the form control. Closes GH issue #154.
- **Return-loan modal now bounces back to the original page after login.** An anonymous user landing on the `/borrower/:id` page and clicking the *Return* button on an active loan used to be sent to `/login` and then dropped on the generic `/loans` list page, losing the borrower context. The handler now reads the `HX-Current-URL` header sent by HTMX, validates the path through the existing same-origin guard, and uses it as the post-login return target. Closes GH issue #133.
- **403 Forbidden body now renders in the user’s language.** A French-locale user clicking an admin-only action as a Librarian (or any 403-producing flow) saw the English

error text instead of the French translation. The locale now travels through the permission-check helper into the error response, so every Forbidden body is rendered in the locale the user already sees everywhere else. Closes GH issue #219.

8.17 What's new in 1.2.0

Version 1.2.0 is the first minor release of the post-v1 era — six change-requests bundled around a single theme: **see what you own, find it faster**. No schema migrations beyond a small additive column the admin bulk-fetch needs to track, no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Dewey code right next to the genre.** The title- detail page now shows the Dewey code as a chip in the media-type / genre / volume-count row at the top of the page (when set). The old standalone Dewey paragraph further down is gone. Closes GH issue #236.
- **Sort series by Dewey or title.** The series-detail page gains a small sort-link row above the gap grid. Three keys: *position* (the original behaviour, kept as default), *Dewey code* (assigned titles with a Dewey first, then assigned without, then gaps — gaps never lead the list, even in descending direction), and *title* (alphabetic, case-insensitive). The location-detail page already had the same Dewey sort column from Epic 5; v1.2 finishes the symmetry. Closes GH issue #235.
- **Fold / unfold the Locations tree.** Every parent node in the `/locations` tree now has a small `/` toggle in front of its row. Clicking it hides the node's subtree; the collapsed-node set is persisted in the browser's `localStorage`, so the layout survives a reload. Useful once your hierarchy grows beyond a screen-full. Closes GH issue #200.
- **One-click "Uncategorized" filter.** The home page gains a dedicated filter chip next to the named-genre chips. Clicking it surfaces titles whose genre is the default `Non classé` bucket — the ones the metadata-fetch chain couldn't classify and you haven't re-classified yet. Closes GH issue #205.
- **Spinner on "Apply selected changes".** The metadata-conflict confirmation form, the one that appears when re-fetching metadata on a title with manual edits, now shows a spinner inside the *Apply* button while the request is in flight. No more wondering whether the click registered. Closes GH issue #215.
- **Bulk re-fetch missing covers.** The `/admin?tab=health` panel gains a *Maintenance* section with a count of titles missing a cover and a one-click *Re-fetch missing covers* button. Clicking it spawns a background task that walks the eligible titles (those with at least one ISBN / ISSN / UPC code, ordered oldest first) and re-triggers the full metadata- fetch chain on each — including the Open Library Covers fallback that landed in v1.1.6/1.1.7. A 500 ms inter- title gap keeps the providers happy; a single-instance lock refuses a concurrent click with a friendly message. Closes GH issue #214.

8.18 What's new in 1.1.9

Version 1.1.9 is a single-fix release on top of 1.1.8. No schema migrations, no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Saving a title no longer fails silently when numeric fields are empty.** The edit form has four numeric inputs (page count, track count, total duration, issue number). For a title where the field is `NULL` (the common case on French books where BnF didn't

supply a page count), the input renders empty. Clicking *Save* used to return HTTP 422 *Unprocessable Entity* — the browser submitted `page_count=` (an empty string) and the form deserializer rejected it as not-a-number. The symptom was the worst kind: the page sat there, no feedback, no change persisted, only the browser console mentioned the 422. 1.1.9 routes the four numeric `Option` fields through the empty-string-tolerant deserializer that was already in place for the series form, so an empty input is correctly understood as *not provided* and the save proceeds. Closes GH issue #238.

8.19 What's new in 1.1.8

Version 1.1.8 closes the third bug in the cover-fallback saga (#225 → #228 → #232) and brings the first illustrations to the manual. One small schema migration ships with this release — review the section on routine updates if you have not seen a release with a migration in a while, but the change is data-only (no breaking schema diff).

- **Re-fetch finally downloads covers on titles whose only customization was the genre.** The previous releases wired the Open Library Covers fallback into both the background scan path (1.1.6) and the user-triggered Re-fetch metadata path (1.1.7), but the form coming back from Re-fetch still hid the cover behind an unchecked *accept cover* checkbox whenever `manually_edited_fields` was non-empty. Investigation showed that mybibli used to record every *genre* change in that list, even though no metadata provider in the chain ever returns the genre. Result: every title you had ever re-classified was permanently pinned to the conflict-confirmation flow, and a one-click Re-fetch silently dropped the cover. v1.1.8 stops tracking `genre_id` in `manually_edited_fields`, ships a one-shot data migration that strips the spurious entry from existing rows, and NULLs out arrays that become empty so the next Re-fetch on those titles takes the *direct-apply* branch — metadata refreshed and cover downloaded silently when Open Library has one. Closes GH issue #232.
- **Illustrations.** Chapter 3 (*Everyday use*) and Chapter 9 (*Best practices*) now carry annotated screenshots taken from a real production install — a Home dashboard with the search bar, genre chips, counters and recent additions, and a Locations tree with the household's room / bookcase / shelf hierarchy. The same images now lead the project README and the website landing page.

8.20 What's new in 1.1.7

Version 1.1.7 is a single-fix follow-up to 1.1.6 — same area (the Open Library Covers fallback introduced for issue #225), different code path. No schema migrations, no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Re-fetch metadata now applies the cover fallback.** v1.1.6 wired the Open Library Covers fallback into the *background* metadata-fetch path that fires on a fresh ISBN scan, but missed the user-triggered *Re-fetch metadata* button on a title's detail page. Symmetric bug: FR titles resolved by BnF still landed cover-less on re-fetch even though Open Library may have had the cover indexed by ISBN. v1.1.7 extracts a single helper (`CoverService::resolve_cover_url_with_fallback`) used by every code path that downloads a cover for a resolved metadata result — background scan, direct-apply re-fetch (no manual edits), and the conflict-confirmation form's *accept cover* checkbox. To recover a previously cover-less title, re-open its detail page and click *Re-fetch metadata*; the cover now arrives in the same request when Open Library has it. Closes GH issue #228.

8.21 What's new in 1.1.6

Version 1.1.6 is a single-fix release on top of 1.1.5. No schema migrations, no breaking changes — routine `docker compose pull && docker compose up -d`.

- **Cover images for titles resolved by BnF.** The metadata-provider chain is "first-match-wins for the whole result". When BnF (the chain's primary FR book provider) resolved a title it would set the metadata fields but never a cover URL — UNIMARC XML, the format BnF returns via its SRU API, simply doesn't carry image URLs. The same happened for Google Books results without an `imageLinks` block (common for self-published Amazon KDP titles). The chain didn't fall through to Open Library for the cover, so most catalog entries landed with no jacket image even when Open Library had one.

mybibli now performs a cover-only fallback: when the chain returns metadata but no `cover_url` and the lookup code is an ISBN, it HEAD-probes the *Open Library Covers API* (<https://covers.openlibrary.org/b/isbn/>). On a 200 the cover is downloaded, resized to 400 px wide, and saved as JPEG 80 % as if the original provider had supplied it. On 404 / network error / non-ISBN code the title stays cover-less, but quietly (no warn-log noise per missed cover).

To recover covers on titles cataloged before this release, open each affected title's detail page and click *Re-fetch metadata*. A bulk-fetch admin action is tracked at issue #214. Closes GH issue #225.

8.22 What's new in 1.1.5

Version 1.1.5 continues the production-feedback patch series on top of 1.1.4. No schema migrations and no breaking changes — routine `docker compose pull && docker compose up -d`. Three improvements:

- **Save title with no metadata.** When the provider chain returns nothing for a scanned ISBN (BnF / Google Books / Open Library all silent), the title was created with the `Non classé` default genre but a manual save from the edit form could trip an FK violation if the submitted `genre_id` was 0 — a corner case that hit users editing freshly-scanned titles whose dropdown had no pre-selection. The handler now falls back to `Non classé` for any `genre_id == 0` input, and the underlying lookup self-heals (reactivates a soft-deleted row, re-seeds when missing). Closes GH issue #203.
- **Modal-Confirm header hardening.** The `ModalConfirmRetargetGuard` now requires `X-Modal-Confirm: true` AND `HX-Request: true` before it will strip `HX-Retarget` / `HX-Reswap` from a response. A non-HTMX client (curl, wget, scripted non-browser) can no longer trip the strip just by sending the modal header. Defense-in-depth — the threat model keeps this at low severity for the single-tenant LAN shape. Closes GH issue #218.
- **Google Books configuration walk-through.** The metadata-providers chapter (Chapter 5) now carries a step-by-step Google Cloud Console procedure for the Books API key, plus a *Troubleshooting empty metadata* checklist for when a scan lands with empty fields. Docs slice of GH issue #202; the broader per-provider visibility ideas (metadata-source badge, per-provider retry, structured error surfacing) remain open for future stories.

8.23 What's new in 1.1.4

Version 1.1.4 is a single-fix release. No schema migrations and no breaking changes for the *application* side — but the upgrade path required a `docker-compose.yml` change, which is why the cover persistence story is documented prominently in the Installation chapter rather than buried here.

- **Covers persist across container upgrades.** The published `docker-compose.yml` template now declares a `mybibli_covers` named volume mapped to `/app/covers`, parallel to `mybibli_db_data`. Pre-1.1.4 the cover JPGs lived inside the container's writable layer and were destroyed by every `docker compose up -d` after a `pull`. Existing users upgrading from a pre-1.1.4 install need to re-fetch covers from each affected title's detail page — a bulk-fetch admin action is tracked at GH issue #214. Closes GH issue #213.

8.24 What's new in 1.1.3

Version 1.1.3 is the first production-feedback-driven patch on top of 1.1.2. No schema migrations and no breaking changes — routine `docker compose pull && docker compose up -d`. Two user-visible improvements:

- **Home page reorder.** The filter results / search results area (the list that updates when you click a genre chip or type in the search box) now sits directly under the search box, ABOVE the indicator widgets (recent additions, recent returns, stats by genre, etc.). Pre-1.1.3 the results landed at the very bottom of the page — clicking a chip triggered the swap, but the swap happened five or more scroll-screens below the user's viewport, so it looked like nothing happened. Closes GH issue #204.
- **Locations tree rows are now clickable.** On the `/locations` page, clicking a location's name (or its icon, L-code, type label, or volume count) navigates to `/location/:id` — the per-location volume list with sortable columns (title, author, genre, Dewey). The edit/delete/add-child buttons stay individually clickable OUTSIDE the link. Closes GH issue #208.

8.25 What's new in 1.1.2

Version 1.1.2 closes Epic 10 — mobile UX and security closeout. There are no schema migrations and no breaking changes. The upgrade is a routine `docker compose pull && docker compose up -d`. The user-visible improvements are:

- **Mobile data tables become cards.** The catalog, loans, borrowers, and series tables now collapse into a compact card list on narrow viewports (anything below the md Tailwind breakpoint). The desktop table layout is unchanged. The same rows are rendered in both shapes — the responsive switch is pure CSS, no JavaScript involved.
- **The admin tabs become a dropdown on mobile.** On narrow viewports, the five-tab admin nav (Health, Users, Reference data, Trash, System) renders as a `<select>` dropdown instead of overflowing horizontally. Deep links to `/admin?tab=...` keep working.
- **“Session expired” feedback is now explicit.** When the per-session CSRF token drifts (e.g. because the browser was idle past the inactivity timeout and the session rotated), the next state-changing action returns a clear server-rendered banner instead of a silent failure. The same flow handles session-timeout 403s.

- **Full accessibility coverage extended to entity pages.** The WCAG 2.2 AA gate in CI now covers every `/title/:id`, `/borrower/:id`, `/contributor/:id`, `/series/:id`, and the first-launch `/setup` wizard, not just the top-level navigation routes. Two contrast violations on the title detail page and the home recent-additions block were fixed inline.
- **Auth posture documented end-to-end.** `docs/auth-threat-model.md` now formalises the single-tenant LAN/NAS deployment shape: what CSRF protects, what session cookies do, why specific design choices were made, and what would need revisiting if multi-tenancy became a goal. Referenced from `README.md` and `CLAUDE.md`.

8.26 Rollback

Rollback is *not* a first-class flow. Schema migrations move forward only — there is no `down.sql`. The supported recovery strategy is:

1. Restore the pre-upgrade database backup (chapter 7).
2. Pin the image tag back to the prior version.
3. Start the stack.

This works because the backup is consistent with the prior schema. Trying to point an older image at a newly-migrated database will fail to boot.

Chapter 9

Best practices

Concentrated lessons from real-world deployments. Every section is optional — pick what fits your scale.

9.1 Storage hierarchy




Figure 9.1: The Locations page on a live install: a household hierarchy with a study room, an in-room bookcase, ten shelves, and two flat-level rooms — each node carrying its L-code and a live volume count.

- **Three levels are usually enough.** Room → shelf → row covers most households. A fourth level (e.g. “box inside row”) becomes paperwork you have to maintain.
- **Name locations the way you talk about them.** “Living-room shelf, top” beats “LR-S01-R01”. The names are visible on barcode labels; treat them as user-facing.
- **Print L-codes once.** Re-printing because you rebalanced shelves is wasted effort — L-codes are stable identifiers; the human-readable name on the label can drift and that is fine.

9.2 V-code rolls

- Print V-codes in rolls of 50 to 100. Pre-printed stickers in a holder beside the scanner station make cataloging *fast*.
- **Stick the V-code before scanning.** It removes a future failure mode where you mislabel a book and have to scan-find it again.
- **Place the sticker consistently.** Inside back cover, upper-right corner, parallel to the spine — pick a rule and stick to it. Future-you doing an audit will thank present-you.

9.3 Genre versus Dewey

mybibli supports both broad genres and Dewey codes. Pick one as your primary navigation aid:

- **Use genres** when the audience browses by mood (“I want a thriller”) and the collection is small enough that “Thriller” is granular enough.
- **Use Dewey** when the audience browses by topic (“books on 18th-century French history”) or when the collection has reached the size where “Non-fiction” is no longer a useful filter.
- Using both is fine; mybibli does not force a choice.

9.4 Audit cadence

- **Quarterly** for collections under a few hundred volumes.
- **Monthly** for collections above that, especially if more than one person catalogs.
- Always run an audit *after* a major shelf rebalance — moving books physically without telling mybibli is the single biggest source of “where is this volume?” bugs.

9.5 User accounts

- **One human, one account.** Sharing a single admin account between two people loses the audit trail and makes the last-active-admin guard less useful.
- **Most editors should be Librarians, not Admins.** The admin panel handles installation-time concerns and rare cleanups; you do not need admin rights to catalog or loan.
- **Rotate the seeded admin password on first login.** On 1.1.0+ the setup wizard forces this; on legacy 1.0.x installs do it manually before exposing the instance to a network.

9.6 Loans

- **Keep borrowers identifiable.** “Marie L.” beats “Marie” the moment the second Marie shows up. Email or phone in the notes field helps for chase-ups.
- **Run the overdue list weekly.** The dashboard indicator nags you for a reason — overdue loans compounded over a year become books you no longer own.
- **Mark damage at return time.** Switching a returned volume to *Worn* or *Damaged* captures the information while the borrower is in front of you; chasing the cost later from memory rarely works.

9.7 Metadata providers

- **Keep keys minimal.** You do not need all three keyed providers from day one — BnF + Open Library cover books for free. Add OMDb + TMDb the day you catalog your first DVD.
- **Treat keys as secrets.** `.env` is in `.gitignore`; keep backups of `.env` encrypted (chapter 7).

9.8 Upgrades

- **Read release notes before pulling :latest.** See chapter 8.
- **Back up immediately before any upgrade.** Even “additive” migrations have non-zero failure surface; a same-minute backup turns a worry into a chore.
- **Pin a version tag for production.** Auto-updates for a household instance are fine; for a community library, an unattended major bump can interrupt service.

Chapter 10

Troubleshooting

When things go wrong, start with the logs and the `/admin > Health` dashboard. This chapter catalogues the problems users actually hit, in rough order of frequency.

10.1 Reading the logs

```
docker compose logs --since 10m mybibli
docker compose logs -f mybibli
```

Logs are emitted as structured JSON. The most useful fields are `level`, `target` (which module logged this), `message`, and any contextual key=value pairs that the log line carries (`user_id`, `volume_id`, ...).

To filter:

```
docker compose logs --since 10m mybibli | grep -i error
```

10.2 App does not start

Symptom: container exits immediately

Run `docker compose logs mybibli` and look for one of:

- missing required environment variable: `DATABASE_URL` — the `.env` file is missing or the variable is unset. See section [1.5](#).
- Failed to create database pool — the database hostname / port / credentials in `DATABASE_URL` are wrong, or the DB is not reachable. From the host, try `docker compose exec mybibli sh -c 'apk add curl; curl <db-host>:<db-port>'` to confirm reachability.
- Failed to run database migrations — usually a permissions issue. Confirm the DB user has `ALL PRIVILEGES` on the database (see section [1.4](#)).
- Address already in use — the published host port is taken. Change `HOST_PORT` in `.env`.

Symptom: every page returns 404

The container is up but routes are missing. Most common cause: you pulled an image tag that does not exist (typo) and Docker fell back to a stale older tag. `docker compose pull` and restart.

10.3 Cannot log in

Symptom: “Invalid credentials” on a fresh 1.0.x install

The seed credentials are `admin/admin`. If they do not work either:

- the database migration that creates them failed silently — look for migration errors in the logs;
- the previous admin already rotated the password (check with whoever installed mybibli);
- you are looking at a 1.1.0+ install, where the seed credentials are planted by the migration but immediately soft-deleted by the issue-#173 seed gate at startup — use the wizard at `/setup` instead.

10.4 Stray users in the users table after the wizard

Symptom: four users in the database, only one was created in the wizard

If you inspect the `users` table directly (via phpMyAdmin or the MariaDB CLI) right after completing `/setup`, you will see *four* rows even though you only entered one username and password. This is by design. Three of the rows cannot be used to log in.

username	role	active	deleted_at
admin	admin	1	set, $\approx$ first-boot time
librarian	librarian	1	set, $\approx$ first-boot time
SYSTEM	system	0	NULL
<your-username>	admin	1	NULL

- `admin` and `librarian` are planted by the seed migrations (which run on every fresh install for the dev / E2E workflow) and immediately *soft-deleted* by the issue-#173 seed gate at startup. Their `deleted_at` is set, so the login query (which filters `deleted_at IS NULL`) refuses them.
- `SYSTEM` is a non-loginable utility account used to attribute the audit-log entries produced by the 30-day auto-purge background task. Its role is `system` (outside the login enum `admin/librarian`), its `password_hash` is the literal string `NO_LOGIN_SYSTEM_USER_HASH_NOT_A_VALID_ARGON2` and its `active` flag is 0.

Tip

To confirm the soft-deleted rows are inert, try logging in as `admin/admin` — it returns an authentication failure. The `/admin > Users` panel filters out both the `SYSTEM` row and any `deleted_at`-set row, so the admin UI shows the single account you created in the wizard.

The two soft-deleted rows are physically removed from the table by the 30-day auto-purge scheduler — expect them to disappear roughly one month after the boot date stamped in `deleted_at`.

Symptom: login redirects back to `/login`

The session cookie is being rejected by the browser. Causes:

- `MYBIBLI_COOKIE_SECURE=true` on a plain-HTTP deployment — the browser refuses to accept a Secure cookie over HTTP. Set the variable to `false` or front the instance with HTTPS.
- Cookies blocked by browser policy (private window with strict third-party blocking, etc.). Test in a default profile.

Symptom: locked out — last admin deactivated

The last-active-admin guard should prevent this, but if a previous admin was deactivated then permanently deleted, recovery requires direct SQL access:

```
docker compose exec db mariadb -u root -p"$MYSQL_ROOT_PASSWORD" mybibli \  
-e "UPDATE users SET deleted_at=NULL, active=1 WHERE username='your-admin';"
```

Then log in normally.

10.5 Metadata never resolves

Symptom: skeleton title stays unresolved

Open `/admin > Health`. Look at each provider’s status. A red cross usually means one of:

- the provider is genuinely down (rare for BnF / Open Library; more common for the smaller services);
- the API key is missing or invalid — open `/admin > System > Metadata Providers` and re-enter the key;
- the host cannot reach the upstream. Confirm with `docker compose exec mybibli sh -c "wget -q https://www.googleapis.com/books/v1 -O -"` or equivalent.

Symptom: only French ISBNs resolve

You disabled (or never set) the Google Books / Open Library keys. BnF is excellent for French-published titles but limited for foreign-language titles. Add at least Open Library (no key required).

10.6 Search returns too little

Symptom: known titles do not surface

- Confirm the title is not soft-deleted (check `/admin > Trash`). Search hides soft-deleted items.
- Confirm the title’s contributors are spelled the way you searched. mybibli’s search is fuzzy but not magical; “Camus” will find “Albert Camus” but “Camu” might not.

Symptom: search is slow

For collections above 10,000 titles, queries that combine contributor + free-text can take a second or two. This is a known limitation; planned indexes will close the gap.

10.7 Cover images are missing

Symptom: every title shows the placeholder

The cover directory is empty or unreadable. Confirm:

- `docker-compose.yml` declares the `mybibli_covers` named volume (or a bind mount on `COVERS_HOST_PATH`) mapped to `/app/covers`. Pre-1.1.4 installs did NOT declare this volume — see the next subsection;
- the metadata fetch actually returned a cover URL (check the title page — if the URL is missing, the provider did not supply one).

Symptom: covers vanished after upgrading the image

You upgraded from a pre-1.1.4 install and all previously cataloged covers are gone. This is issue #213 — the older compose template did not persist `/app/covers`, so the writable layer holding those JPGs was destroyed when the new container replaced the old one.

Fix forward: pull the 1.1.4 `docker-compose.yml`, run `docker compose down && docker compose up -d`. Subsequent covers are persisted in the `mybibli_covers` volume.

Recovery of lost covers: open each affected title's detail page and click *Re-fetch metadata* to repopulate from the provider chain. A bulk action is tracked at issue #214.

10.8 Where to file bugs

Issues you cannot resolve, or behaviour that contradicts this manual: open a ticket at <https://github.com/guycorbaz/mybibli/issues>. The issue template asks for the version (visible in the page footer), the install method, and the steps to reproduce — please fill them in.

Chapter 11

API & integrations

mybibli ships a small JSON HTTP API at `/api/v1/*` that lets external scripts and AI tools read your catalog and update the classification of titles. The API runs on the same web server as the human-facing UI — there is no separate daemon, no extra port to open.

This chapter walks through:

- creating an API key in the admin panel,
- the two scopes (*read only* vs. *read + write*),
- the endpoints exposed under `/api/v1/*`,
- three end-to-end examples (`curl`, Python, an AI classifier loop).

11.1 Why an HTTP API?

The everyday use case is offloading classification work to an external tool — for instance, asking an LLM to suggest a Dewey code, a genre, or a clean subtitle for every newly-added English book. The LLM walks the catalog over `GET /api/v1/titles`, fetches the detail of each via `GET /api/v1/titles/{id}`, computes a suggestion, and writes it back with `PATCH /api/v1/titles/{id}`. Every `PATCH` is audited so you can roll back changes if a model produced garbage.

The API is *not* meant for high-throughput batch ingest — for that, use the SQL database directly during a maintenance window. It *is* meant for the kind of slow, careful, opinionated write that fits inside an LLM agent loop.

11.2 Creating an API key

API keys are managed under **Administration** → **API keys**. Only Admin users see this tab.

1. Click **Administration** in the navigation bar.
2. Open the **API keys** tab.
3. Fill in a **label** — a name only you see, e.g. “*AI classifier*” or “*Backup script*”. The label is for the audit log; the key plaintext itself is random and unrelated to it.
4. Choose a **scope**:

- **Read only** — GET endpoints only. Safe for read-only dashboards, mirroring tools, or curiosity-driven scripts.
 - **Read + write** — additionally allows PATCH `/api/v1/titles/{id}` for a narrow allow-list of fields (genre, dewey, description, subtitle). Use this for LLM agents or curation tools.
5. Click **Create key**. A modal appears showing the full plaintext exactly once — copy it now. Once you close the modal there is no recovery path; if you lose the key, revoke it and create a new one.

The plaintext shape is `mybibli_<scope>_<random40>`, where `<scope>` is `ro` or `rw`. Only the `argon2` hash and the first 12 characters of the plaintext (the prefix) are persisted; the prefix is what `mybibli` uses to narrow candidates during authentication. The hash itself cannot be reversed.

11.3 Authenticating

Every `/api/v1/*` request must carry one of:

- Authorization: Bearer `<plaintext>` — preferred.
- X-API-Key: `<plaintext>` — convenience fallback.

A missing or invalid key returns 401 `Unauthorized` with a JSON envelope:

```
{ "error": "unauthorized",
  "reason": "missing_api_key" }
```

A read-only key on a write endpoint returns 403 `Forbidden` with the missing scope named:

```
{ "error": "forbidden",
  "reason": "insufficient_scope",
  "scope_required": "write" }
```

CSRF protection is intentionally bypassed for `/api/*` because bearer-token authentication does not ride on cookies and the cross-site replay risk that CSRF defends against does not apply.

11.4 Endpoints

11.4.1 GET `/api/v1/titles`

Paginated list. Supports `?page=N`, `?q=<search>`, `?genre_id=N`, and `?dewey_prefix=8`.

Returns:

```
{ "items": [ { "id": 42, "title": "...",
               "subtitle": null,
               "language": "en",
               "media_type": "book",
               "isbn": "9780...",
               "publisher": "...",
               "publication_year": 2020,
               "genre_id": 7,
               "genre_name": "Fiction",
               "dewey_code": "813.54" },
  ... ],
```

```
"page": 1, "total_pages": 12,
"total_items": 287 }
```

11.4.2 GET /api/v1/titles/{id}

Full detail, including contributors, volumes (with their location), and series assignments. The `version` field is the optimistic-locking handle — keep it next to the data your script will edit.

11.4.3 GET /api/v1/genres, /locations, /series

Reference data. Use these when an LLM needs to know which genres exist before it suggests a `genre_id`.

11.4.4 GET /api/v1/dewey/{prefix}

Bucket lookup. Pass a partial Dewey prefix (e.g. 8 for literature) and get back every title whose `dewey_code` starts with it. Useful for “find all books that already have a Dewey code in the 800s so I can compare”.

11.4.5 PATCH /api/v1/titles/{id}

Write-scope only. Updates the four whitelisted fields:

- `subtitle` — nullable.
- `description` — nullable.
- `dewey_code` — nullable; digits, dots, slashes only.
- `genre_id` — must reference an existing, non-deleted genre.

The JSON body uses partial-update semantics:

- Field **omitted** → no change.
- Field set to `null` → clear to `NULL` (for the three nullable fields above; `null` on `genre_id` returns 400).
- Field set to a value → update.

`version` is required. On mismatch the response is 409 `Conflict` carrying both the expected and the supplied version — refetch and retry.

A successful PATCH writes one row into `admin_audit` with action `api_patch_title`, attribution via the issuing admin (`user_id`) and the API key (`details.via_api_key_id + details.via_api_key_label`), and a per-field `old/new` diff. A no-op PATCH (no field actually changed) returns 200 but does NOT bump `version` and does NOT write an audit row.

11.5 Examples

11.5.1 curl

```
KEY="mybibli_ro..."
```

```
# List the first page of titles.
curl -H "Authorization: Bearer $KEY" \
```

```

https://library.example.org/api/v1/titles

# Show one title in full.
curl -H "Authorization: Bearer $KEY" \
      https://library.example.org/api/v1/titles/42

# PATCH a Dewey code (needs a write-scope key).
KEY_WRITE="mybibli_rw_..."
curl -X PATCH \
      -H "Authorization: Bearer $KEY_WRITE" \
      -H "Content-Type: application/json" \
      -d '{"version": 3, "dewey_code": "813.54"}' \
      https://library.example.org/api/v1/titles/42

```

11.5.2 Python

```

import os, requests

BASE = "https://library.example.org/api/v1"
KEY  = os.environ["MYBIBLI_API_KEY"]
H    = {"Authorization": f"Bearer {KEY}"}

# Walk every page.
page = 1
while True:
    r = requests.get(f"{BASE}/titles?page={page}",
                     headers=H, timeout=30)
    r.raise_for_status()
    data = r.json()
    for t in data["items"]:
        print(t["id"], t["title"], t.get("dewey_code"))
    if page >= data["total_pages"]:
        break
    page += 1

```

11.5.3 An LLM classifier loop

The intended write-flow looks like this — pseudocode:

```

for t in walk_titles(KEY_READ):
    if t["dewey_code"]:
        continue                # already classified
    detail = get(f"/titles/{t['id']}")
    suggestion = ask_llm(detail) # returns {"dewey_code": ...}
    if not suggestion.get("dewey_code"):
        continue
    patch(
        f"/titles/{t['id']}",
        json = {
            "version": detail["version"],
            "dewey_code": suggestion["dewey_code"],
        },
        key = KEY_WRITE,

```

)

If two classifiers race on the same title, the loser gets a 409 and re-fetches; no silent clobber. The audit log records who (key) suggested what — handy for rolling back a model that has a bad day.

11.6 Revoking a key

Revocation is one click in the admin panel. A revoked key *cannot* be restored; the audit history of past usage is preserved alongside the now-defunct row. If you suspect a key has leaked (committed to a public repo, e-mailed by mistake, found in a log file), revoke immediately and create a new one.

11.7 What the API does *not* expose

By design, the v1 API is read-mostly. It does not:

- create or delete titles, volumes, contributors, borrowers, loans, or series;
- update any field outside the four-field PATCH allow-list;
- return passwords, password hashes, session tokens, or any other credentials;
- accept multipart uploads (no cover-image PUT).

These are intentional v1 limits — the use case that justified the API is classification, and that fits inside the narrow allow-list. Wider write-surface would change the security posture (CSRF considerations, mass-mutation rate limits) and is deferred until a concrete need surfaces.

Chapter 12

Operations & debugging

This chapter documents where logs live on disk, how to control their verbosity, and the commands you typically need when a user reports odd behavior and you want to see what the app actually did.

12.1 Where the logs live

Starting with **v1.7.0**, mybibli writes its logs to two places simultaneously:

- **stdout** (unchanged) — visible via `docker compose logs mybibli`. Best for live-tailing during an active session.
- **A daily-rotating file on disk**, by default under `/var/log/mybibli/` inside the container, mounted on a named volume (`mybibli_logs`) on the host so files survive `docker compose pull && up -d`. Best for *post-mortem* debugging.

Files are named `mybibli.log.YYYY-MM-DD` (one per UTC day). A background task purges files older than 30 days automatically.

Changing the in-container log path

The default in-container log directory `/var/log/mybibli/` is controlled by the `MYBIBLI_LOG_DIR` environment variable. Change it only if you remount the volume elsewhere — the shipped `docker-compose.yml` maps the `mybibli_logs` named volume (or a bind mount via `LOGS_HOST_PATH`) onto `/var/log/mybibli`, so changing one without the other lands logs in the writable layer where they get wiped on every `docker compose up -d` after a pull.

See section [1.3](#) for the bind-mount setup (host-visible logs for DSM Log Center / journald shipping).

12.2 Tailing logs from the host

```
# Tail the current day's file in real time
docker compose exec mybibli tail -f /var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d)

# Last 100 lines of the current day's file
docker compose exec mybibli tail -n 100 /var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d)

# Show all available days
docker compose exec mybibli ls /var/log/mybibli/
```

```
# Search across all retained days for a keyword
docker compose exec mybibli sh -c "grep ERROR /var/log/mybibli/mybibli.log.*"
```

If you mounted the volume as a bind mount (`LOGS_HOST_PATH` in `.env`, see section 1.3), the files are visible directly on the host without needing `docker compose exec`.

12.3 Log levels

Two environment variables control verbosity:

- `MYBIBLI_LOG_LEVEL` (preferred, v1.7.0+) — accepts either a plain level (`trace`, `debug`, `info`, `warn`, `error`) or a tracing-subscriber `EnvFilter` directive list, e.g. `mybibli=debug,sqlx::query=warn`.
- `RUST_LOG` — legacy fallback, same format. Honored when `MYBIBLI_LOG_LEVEL` is unset.

Default is `info` — production-safe (`info` + `warn` + `error`, no per-query SQL noise, no per-request trace floods).

Recommended levels

info Production default. Catches authentication events, metadata fetches, scheduled-task runs, user-facing errors. Roughly 10–100 lines per day on a single-household install.

mybibli=debug Active investigation of an mybibli feature. Adds request-routing, cache hits/misses, scan dispatch decisions. Roughly 10× the `info` volume.

mybibli=trace Deepest one-shot investigation. Includes per-call tracing within mybibli’s own code. Very noisy; use briefly and bump back to `info`.

mybibli=debug,sqlx::query=info Active investigation AND see every SQL statement (useful when a query is suspect).

How to change the level

Two paths, depending on whether you want the change to survive container restarts or just for one investigation.

From the admin UI (v1.7.1+, recommended)

Log in as Admin, go to `/admin?tab=system`, scroll to the **Logging** section. Enter the new directive in the `Log level` field (anything that `tracing-subscriber::EnvFilter` accepts — plain level or full directive list). Click **Save**. The next log line uses the new filter — no container restart, no `docker compose up -d`. The value is persisted to the `settings` table, so it survives container restarts.

From `docker-compose.yml` / `.env`

Edit the file, set `MYBIBLI_LOG_LEVEL`, then:

```
docker compose up -d
```

The container restarts and the new level takes effect from the next log line. **Caveat:** the admin-saved value in the `settings` table wins on the next boot, so if you’ve ever saved a level via the

admin UI, that’s what will load — the env var only matters on a fresh install before the first save.

12.4 Reading the structured JSON

Logs are emitted as line-delimited JSON. Each line carries:

- `timestamp` — ISO 8601 UTC.
- `level` — TRACE / DEBUG / INFO / WARN / ERROR.
- `target` — the Rust module that emitted the line (e.g. `mybibli::routes::catalog`). Filter on this to narrow a search to a subsystem.
- `fields.message` — the human-readable summary.
- `additional fields.*` — contextual keys like `user_id`, `volume_id`, `isbn`, `title_id`.

To pretty-print:

```
docker compose exec mybibli tail -n 50 /var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d) \
| jq -r ' [.timestamp, .level, .fields.message] | @tsv'
```

12.5 Sharing a log excerpt for a bug report

When opening an issue at <https://github.com/guycorbaz/mybibli/issues>:

1. Note the wall-clock time when the bug occurred (e.g. “around 14:30 UTC today”).
2. Extract the relevant 5-minute window:

```
# Replace 14:25 / 14:35 with your window.
docker compose exec mybibli sh -c \
  "awk '/\"timestamp\": \"$(date -u +%Y-%m-%d)T14:2[5-9]/, /\"timestamp\": \"$(date -u +%Y-%m-%d)T14:3[0-5]/' \
    /var/log/mybibli/mybibli.log.$(date -u +%Y-%m-%d)"
```

3. Paste the output into the issue body, fenced as ````json`.

Redact anything you consider sensitive — log lines may contain borrower names, ISBNs of titles you own, or session-token prefixes (but never full session tokens or password hashes; those are kept out of log emission by design).

12.6 Retention & rotation

- Rotation happens **automatically** at UTC midnight, handled in-process by `tracing-appender::rolling`.
- Files older than **30 days** are purged once per day by a background task. The retention window is hardcoded in v1.7.0; a future release will expose it as an admin setting.
- The current day’s file is never deleted (the cutoff is strict).

If you want to keep logs longer than 30 days, copy the rotated files out of the volume to long-term storage:

```
docker compose cp mybibli:/var/log/mybibli/. ./logs-archive/
```

Chapter 13

Glossary

Admin A user account with role `admin`. The only role that can open the `/admin` panel and edit users, reference data, trash or system settings.

Anonymous A visitor without a session cookie. Read-only access to the public catalog. No cataloging, no loans, no edits.

Audit The routine “walk the shelves and confirm everything is where the database says it is” workflow described in section 3.6.

BDGest A French-language comic / BD catalogue used as a metadata source for that media type.

BnF *Bibliothèque nationale de France*. The French national library catalogue, used as the primary metadata source for French-language books.

Borrower A person who loans volumes from the library. Modeled as its own entity (name, optional contact information, loan history).

Cataloging The act of registering a new title and / or volume in the database. Typically barcode-driven: scan, confirm, stick label, repeat.

Cover image The visual front of a book / disc / record, downloaded from a metadata provider and cached under `COVERS_DIR` on disk.

CSP Content Security Policy. A browser-enforced instructions header that restricts which scripts and styles the page may execute. mybibli ships a strict policy that forbids inline scripts and styles.

CSRF token A per-session secret embedded into every form. mybibli rejects POST / PUT / DELETE requests whose token does not match — a defense against cross-site request forgery.

Dewey code The Dewey Decimal Classification number, used for topic-based catalog navigation. Optional per title.

EAN-13 The 13-digit barcode standard. ISBNs are encoded as EAN-13 since 2007; CDs, DVDs and other media carry EAN-13 codes too.

Exemplaire Synonym for *volume* (the physical copy). Used by some French libraries; mybibli prefers *volume*.

Genre A free-form classification of a title's subject (Fiction, Non-fiction, Mystery, ...). Defined in the genres reference table.

Hard-delete Deleting a row outright. mybibli only hard-deletes from the `/admin > Trash` panel and from the auto-purge background task; everything else is a soft-delete.

HTMX The JavaScript library mybibli uses to swap fragments of the page in place without full reloads. Operators do not need to understand HTMX to use mybibli.

ISBN International Standard Book Number. Identifier printed on the back cover of books, encoded as an EAN-13 barcode.

Last-active-admin guard The rule that prevents deactivating or permanently deleting the last user with the `admin` role. See section 2.6.

L-code A storage-location barcode label. Format: L followed by a zero-padded number.

Librarian A user account with role `librarian`. Can catalog, loan, return and edit most data but cannot administer users or settings.

Loan A record that a specific volume is currently out with a specific borrower. Tracks the loan date, due date, return date, and the volume's pre-loan location.

Location The physical place where a volume lives. Modeled as a tree (room → shelf → row), identified by an L-code.

MariaDB The database engine mybibli stores its data in. A fork of MySQL.

Media type The kind of object a title represents: Book, Comic / BD, Audio, Film / series. Drives provider selection and field display.

Multi-position volume A single physical volume that contains several positions in a series — typically an omnibus comic album.

MusicBrainz A volunteer-maintained music metadata catalogue, used as the primary source for audio media.

Node type A label applied to storage locations to indicate what kind of container they are (room, shelf, row, box). Pure label — no semantic enforcement.

OMDb Open Movie Database. Metadata source for films and TV series.

Open Library An open volunteer-driven catalogue of books, used as a metadata source.

Optimistic locking The technique mybibli uses to detect concurrent edits. Two users editing the same title get a clear error on the second save instead of silently overwriting each other.

Reactivation Restoring a soft-deleted row. For users, clears `deleted_at` and lets them log in again. For reference data, the term also covers re-creating a row with the same name as a soft-deleted one, which transparently undoes the delete.

Reference data The four taxonomy tables: genres, contributor roles, volume states, location node types. Editable from `/admin > Reference data`.

Series A named ordering of related titles (e.g. a trilogy or an open-ended comic series). Each title's position in the series is a small integer or a range (for omnibuses).

Setup wizard The four-step first-launch flow at `/setup` that an empty instance presents until an admin account exists.

Soft-delete Setting a row's `deleted_at` column to “now” instead of physically removing the row. Soft-deleted rows are invisible to normal queries but surface in `/admin > Trash` for 30 days, after which the auto-purge hard-deletes them.

Title The bibliographic entity (a book, a film, an album). A title can have multiple volumes — for example, two copies of the same novel on different shelves.

TMDb The Movie Database. A community-curated metadata source for films and TV series.

Trash The `/admin > Trash` tab. Lists soft-deleted entities of every type and lets an admin restore them or permanently delete them.

V-code A per-volume barcode label. Format: `V` followed by a zero-padded number.

Volume The physical object on a shelf — one copy of a title. A title can have many volumes; each carries its own V-code, location and state.

Volume state A reference-data row describing the condition of a volume: New, Good, Worn, Damaged,

Index

1.1.0 upgrade, 31

admin panel, 10
admin role, 26, 62
anonymous, 26, 62
audit, 17, 62

backup
 .env, 28
 covers, 28
 database, 28
BDGest, 22, 62
BnF, 22, 62
borrower, 17, 62

Code-128, 12
comic / BD, 20
contributor roles, 13
covers, 62
 directory, 3
 persistence, 6
CSRF, 62

Dewey, 47, 62

EAN-13, 12
env vars
 .env.example, 7
 provider keys, 24

film, 20

genres, 13, 63
Google Books, 22

HID keyboard mode, 12

i18n, 13
inactivity timeout, 27
installation, 3

L-code, 11, 63
last-active-admin guard, 63

librarian, 26, 63
loan, 63
location, 16, 63
location tree, 10
logs
 directory, 4

media types, 20, 63
metadata providers, 3, 22
multi-position, 20, 63
MusicBrainz, 20, 22, 63
MYBIBLI_SEED_DEV_USERS, 32

node type, 11, 13, 63

OMDb, 21, 22, 63
omnibus, 20
Open Library, 22, 63

roles, 14, 26

scan field, 15
security
 1.1.0 floor, 3
series, 20, 63
session cookie, 26
setup wizard, 6, 9, 64
shelving, 18
soft-delete, 64
storage locations, 10
SYSTEM user, 50

TMDb, 21, 22, 64
troubleshooting
 DATABASE_URL, 49

V-code, 11, 64
volume, 10, 11, 64
volume state, 13, 16, 64

wish list, 18
/wishlist, 18